Why Are We Near Real-Time?

Kristo Raun

Introduction to Near Real-Time Data Analytics

August 2022

UNIVERSITY OF TARTU
1632

Everything is relative

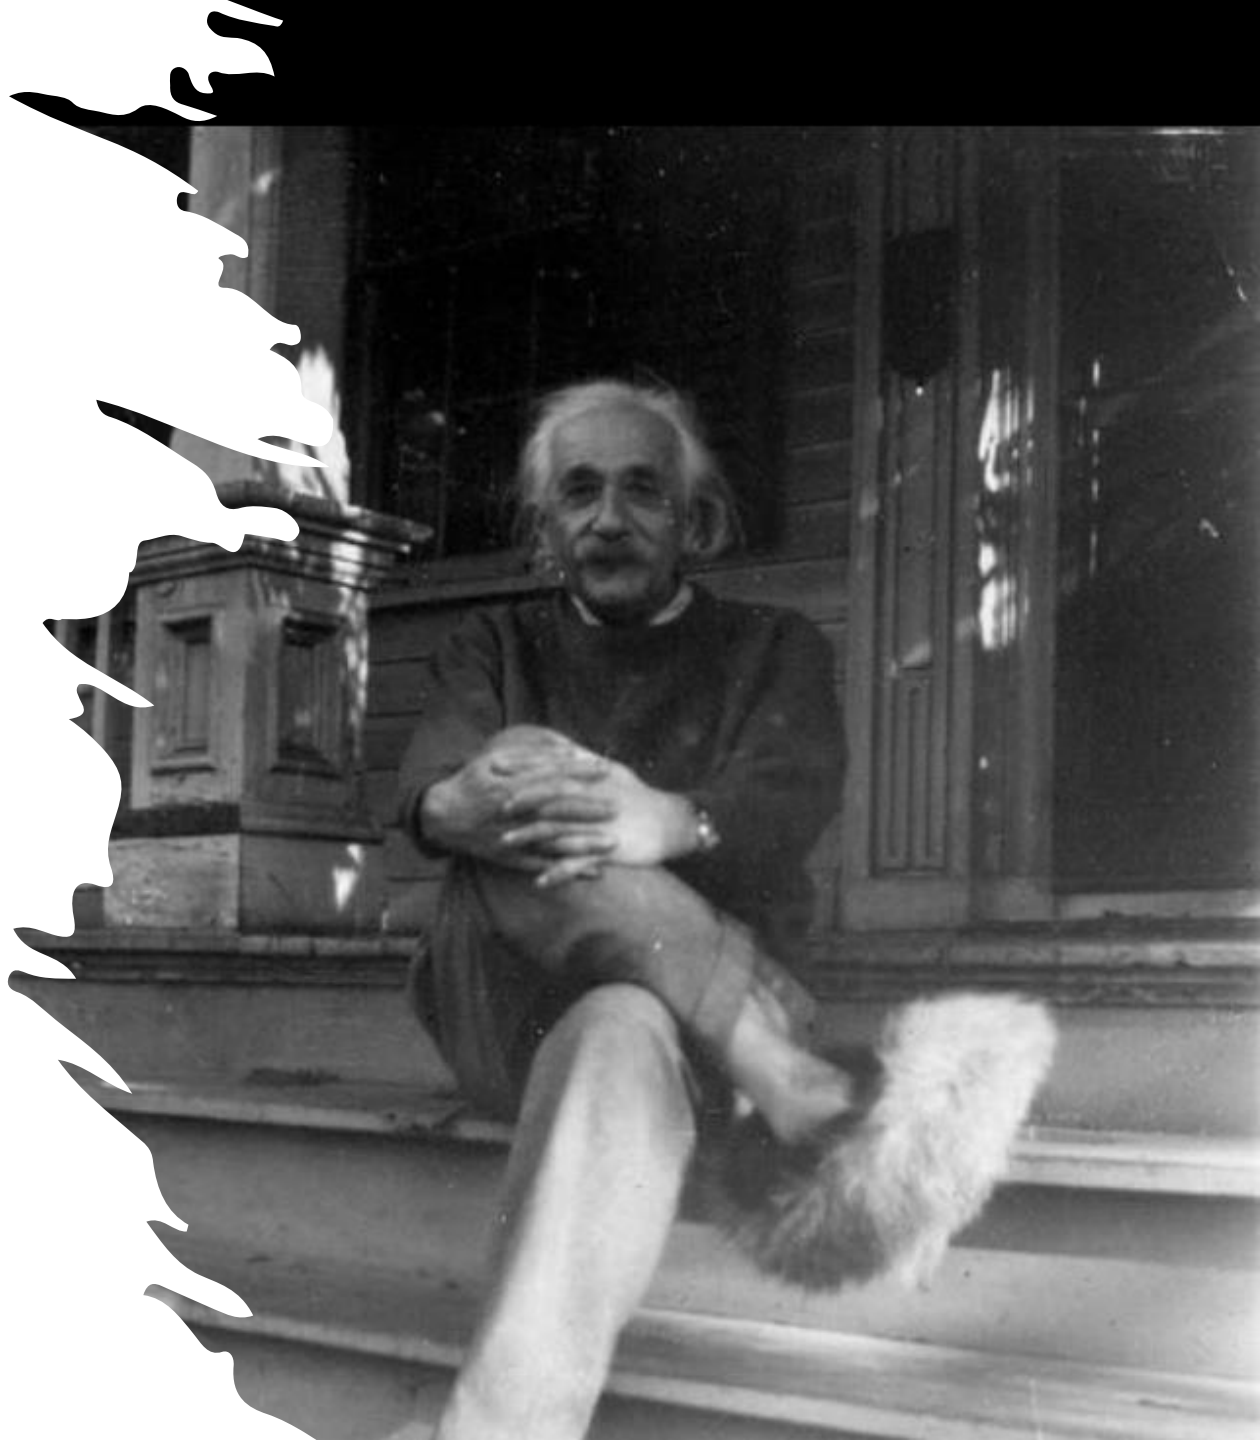

Everything is relative

When you sit with a nice girl for two hours you think it’s only a minute.

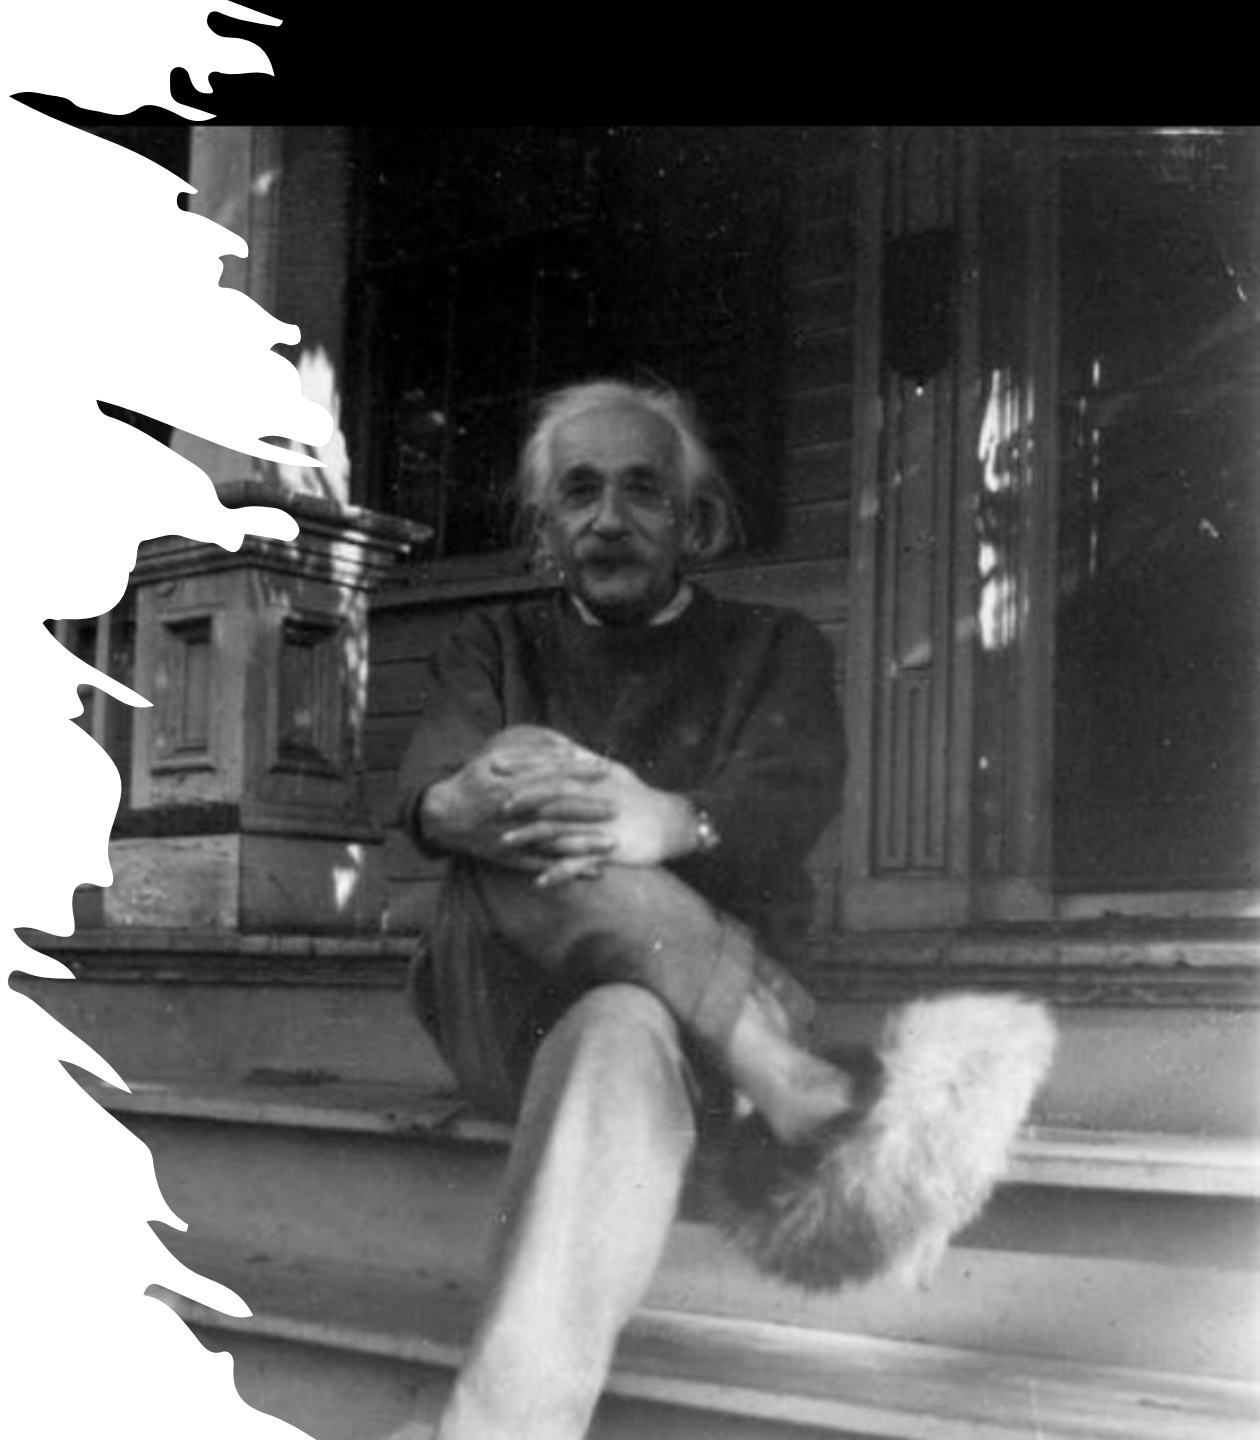

Everything is relative

When you sit with a nice girl for two hours
you think it’s only a minute.

When you receive streaming data for a minute
you think it’s two hours.

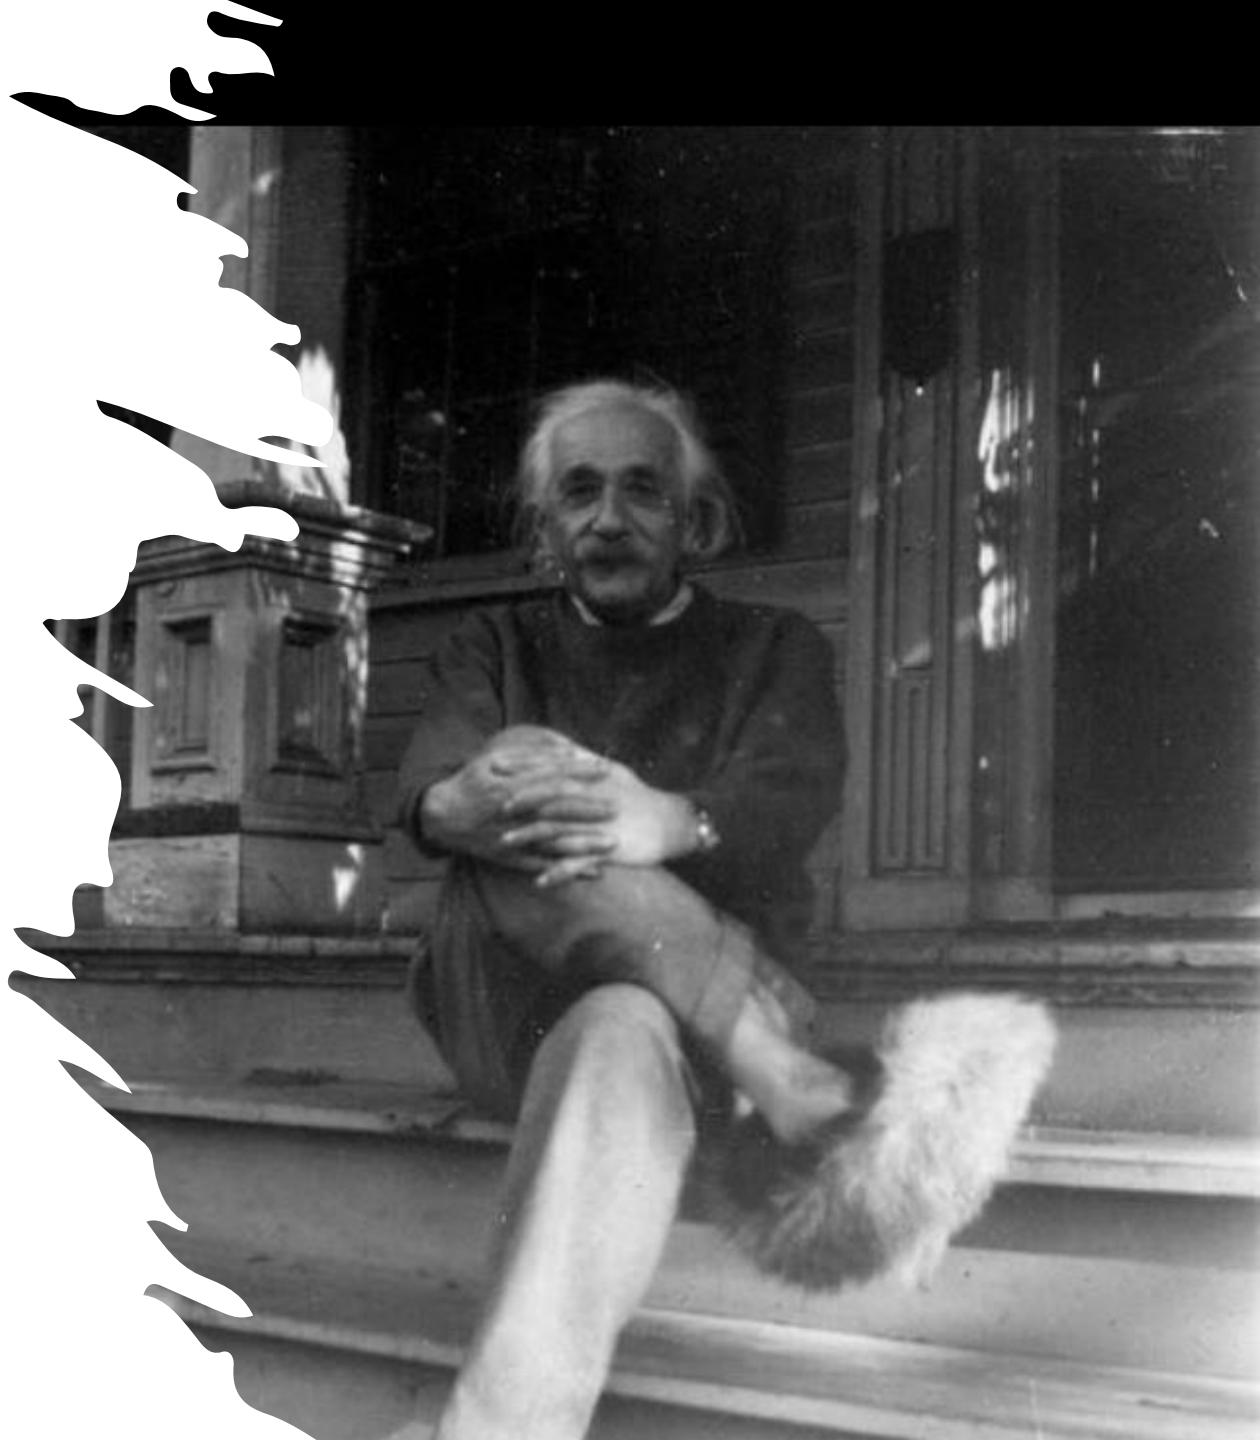

About me

About me

- 2008 – BA start

About me

- 2008 – BA start
- 2011 – working w data

About me

- 2008 – BA start
- 2011 – working w data
- 2015 – MSc start, 1st IT job

About me

- 2008 – BA start
- 2011 – working w data
- 2015 – MSc start, 1st IT job
- 2018-19
  - 1st child
  - live in farm
  - data engineer

About me

- 2008 – BA start
- 2011 – working w data
- 2015 – MSc start, 1st IT job
- 2018-19
  - 1st child
  - live in farm
  - data engineer
- 2021 – PhD start

Agenda for the week

Agenda for the week

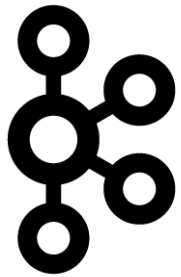

Agenda for the week

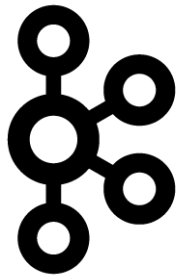

Agenda for the week

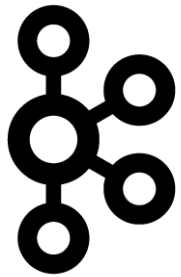

Agenda for the week

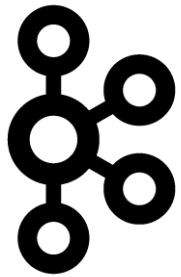

Agenda for today

Agenda for today

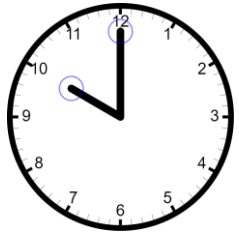

Why are we near real-time?

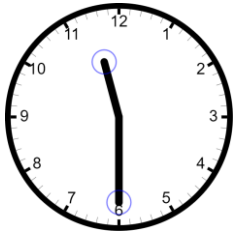

Agenda for today

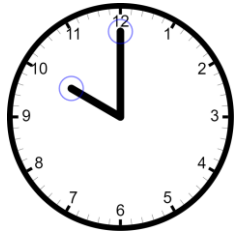

Why are we near real-time?

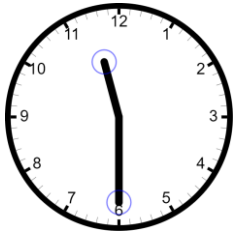

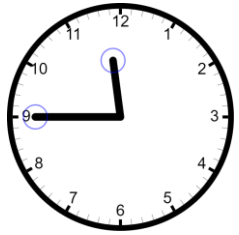

Apache Kafka setup

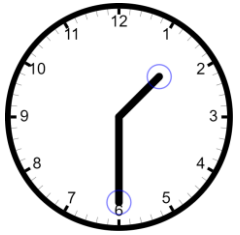

Agenda for today

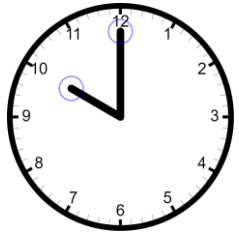

Why are we near real-time?

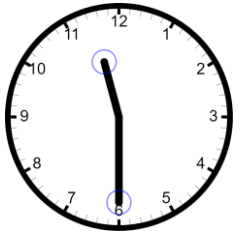

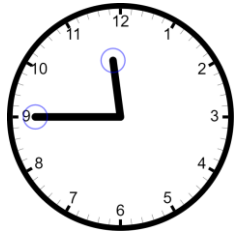

Apache Kafka setup

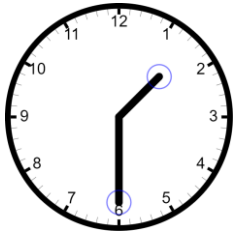

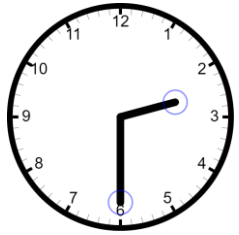

Apache Kafka practice

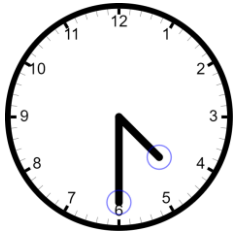

Why?

Why?

Why?

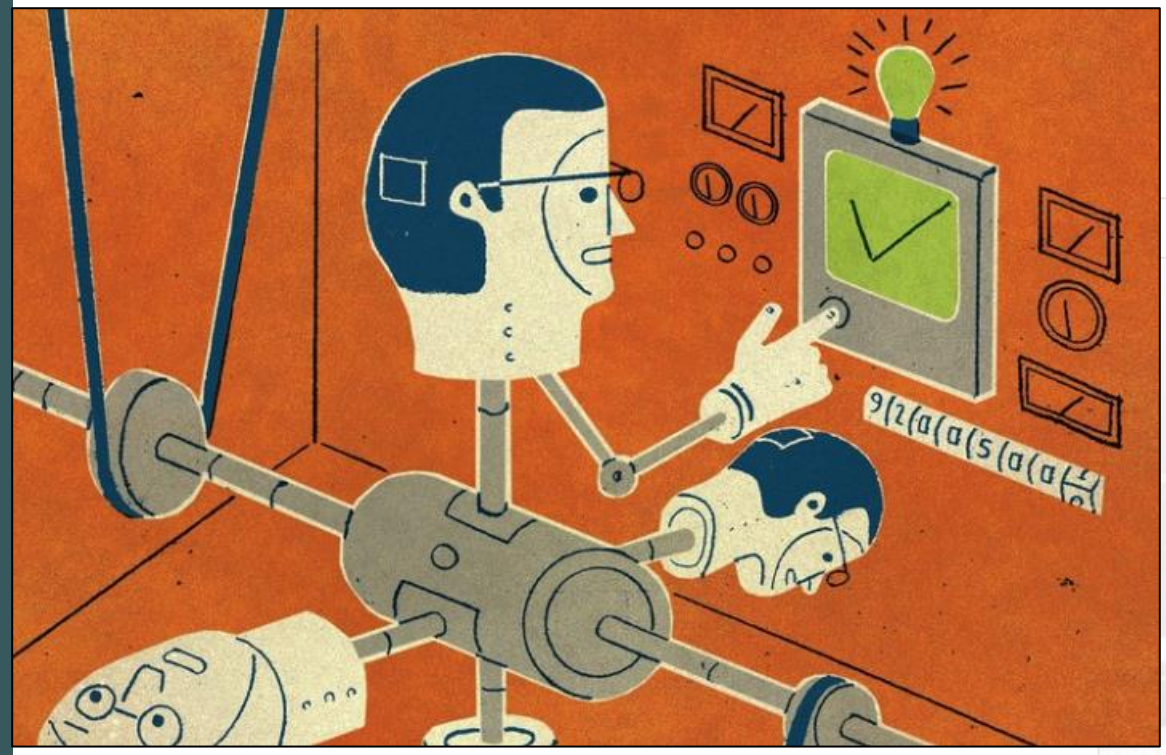

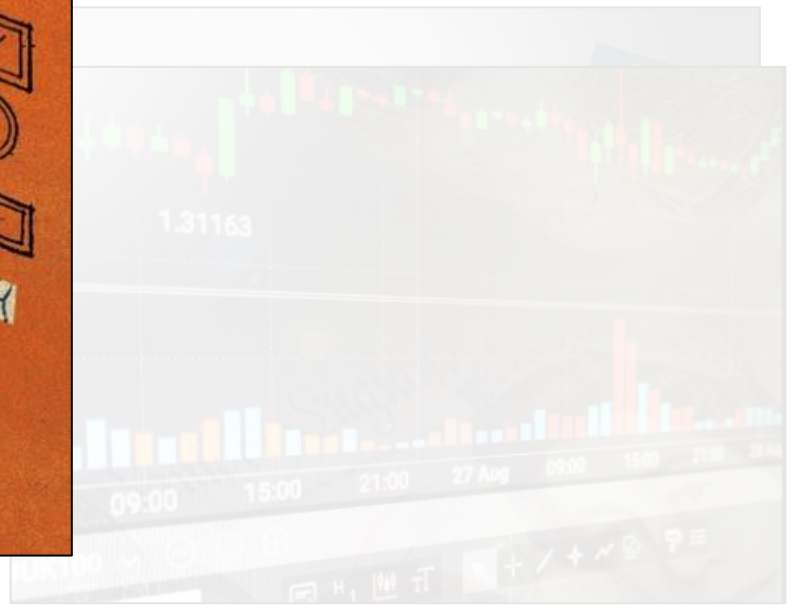

Why?

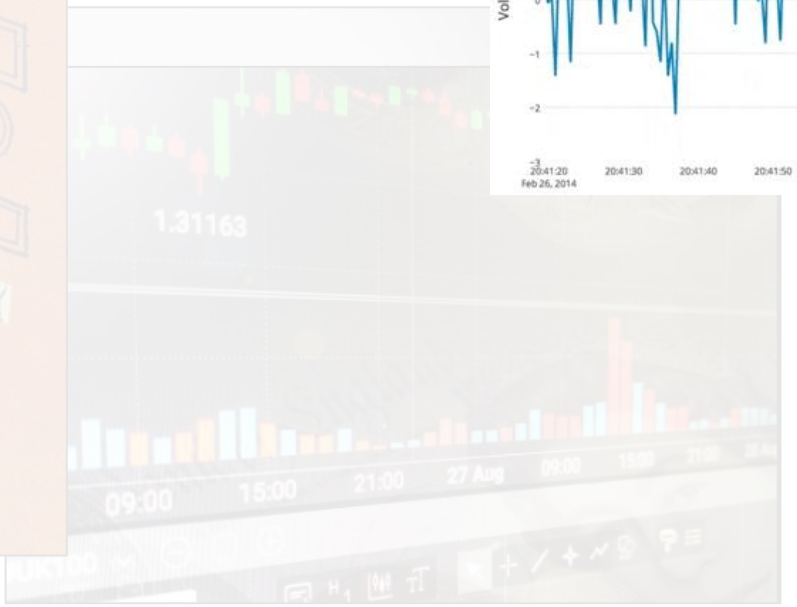
1.31163
09:00
15:00
21:00
27 Aug

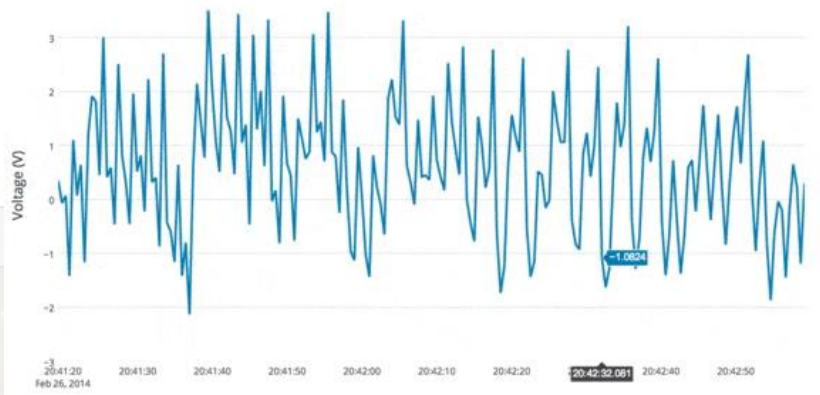
Voltage (V)
3
2
1
0
-1
-2
-3
-1.0624
20:41:20
Feb 26, 2014
20:41:30
20:41:40
20:41:50
20:42:00
20:42:10
20:42:20
20:42:32.081
20:42:40
20:42:50

Why?

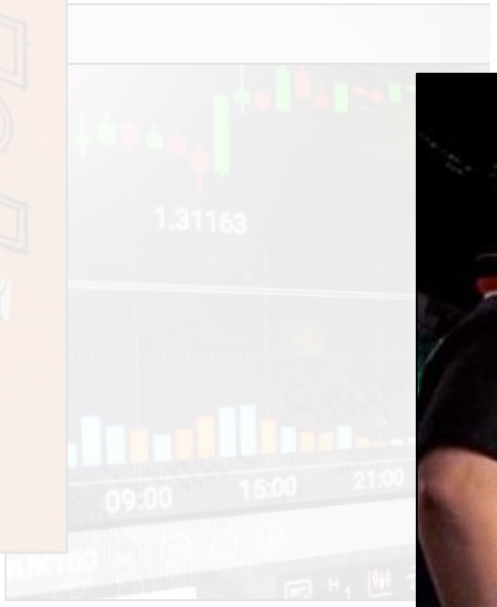

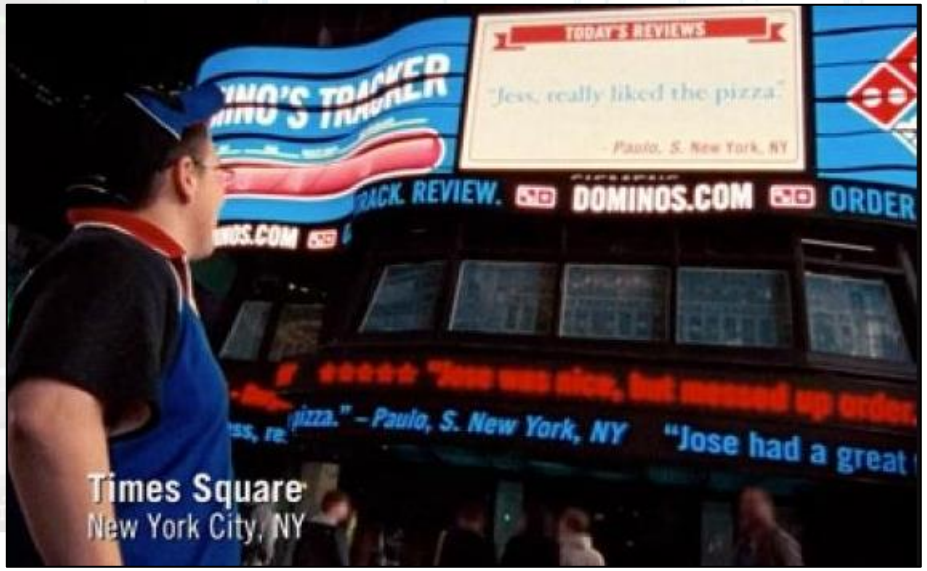
TODAY'S REVIEWS
"Jess, really liked the pizza."
Paulo, S. New York, NY
DOMINOS.COM
ORDER
"Jose had a great
Times Square
New York City, NY

Why?

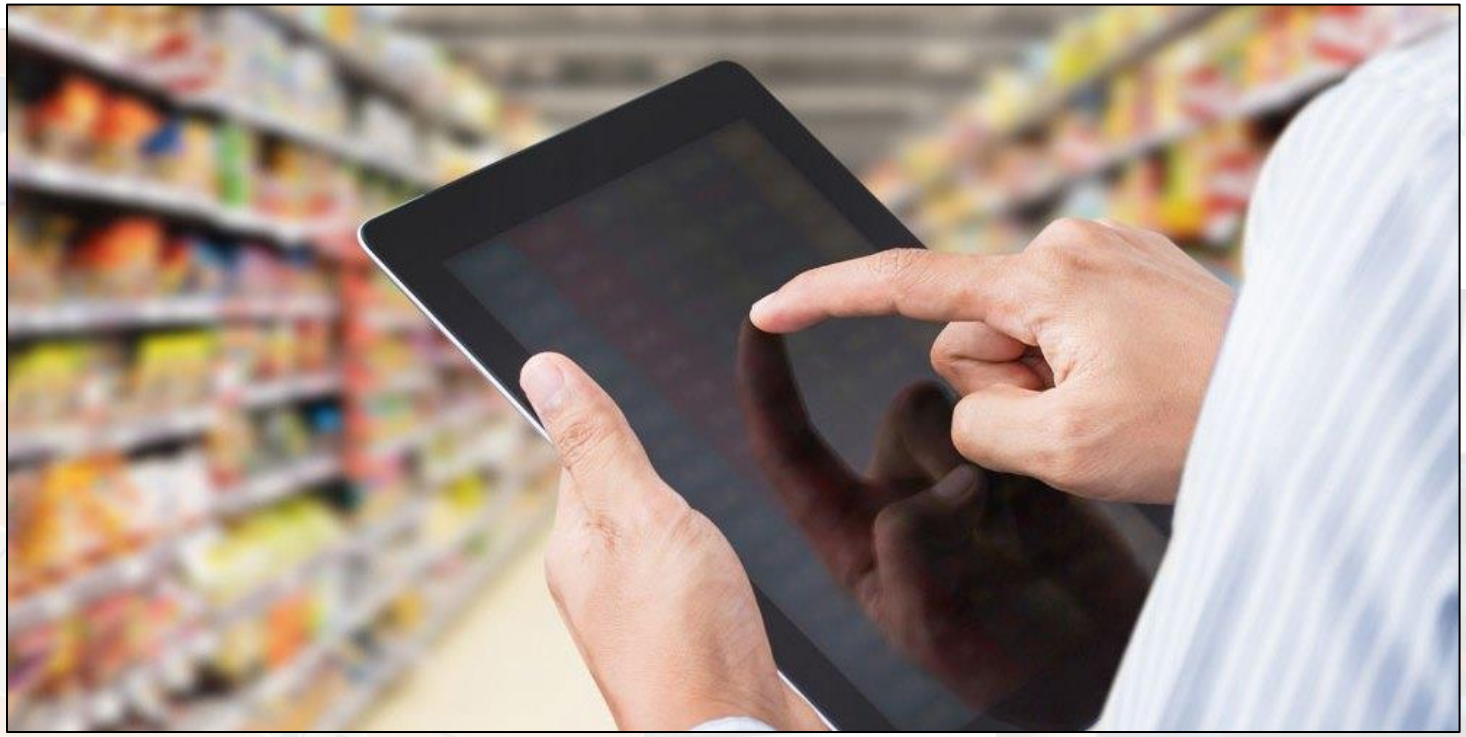

Why?

Timeline

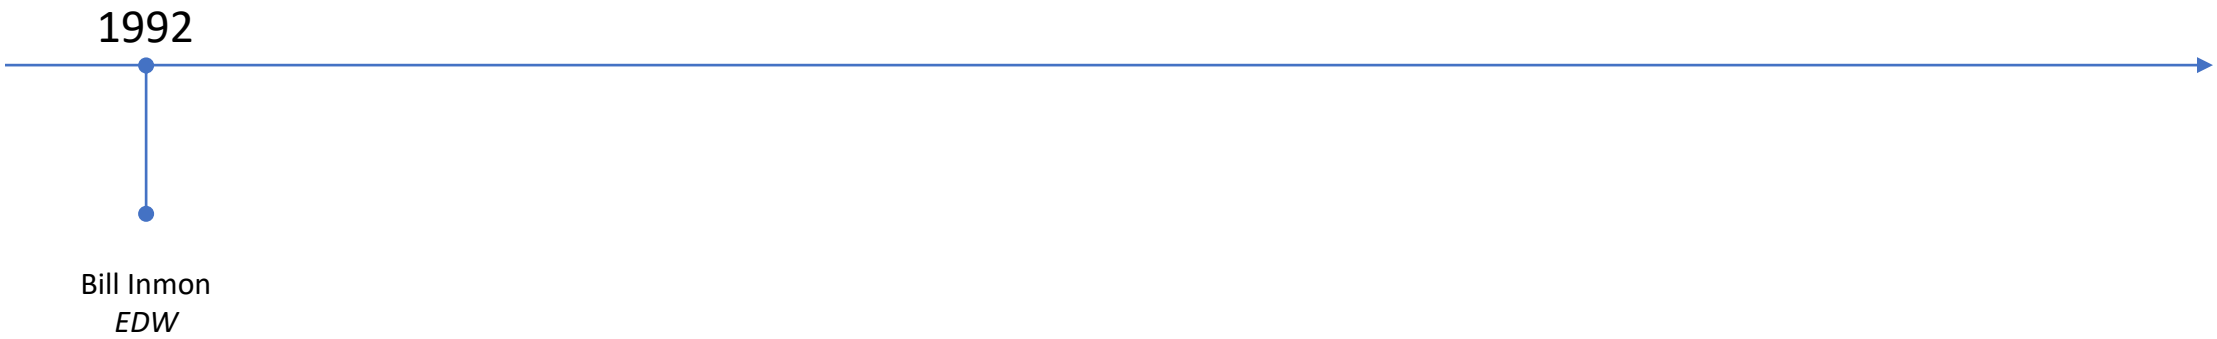

Batch processing, data warehouses

Timeline

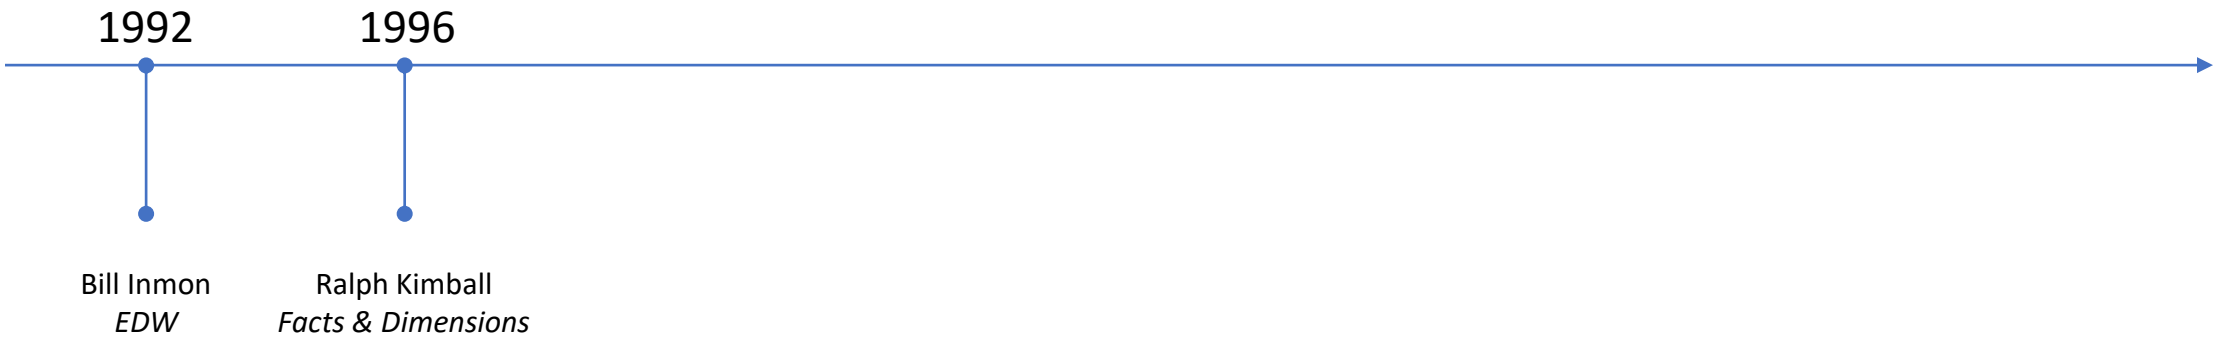

Batch processing, data warehouses

Timeline

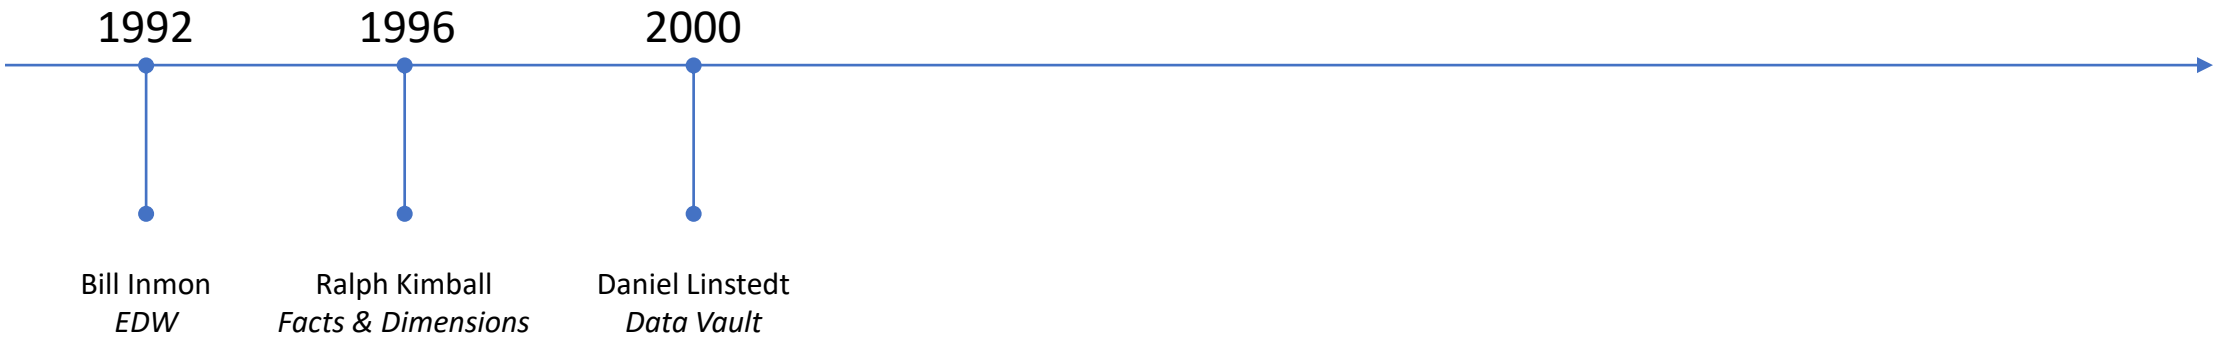

Batch processing, data warehouses

Timeline

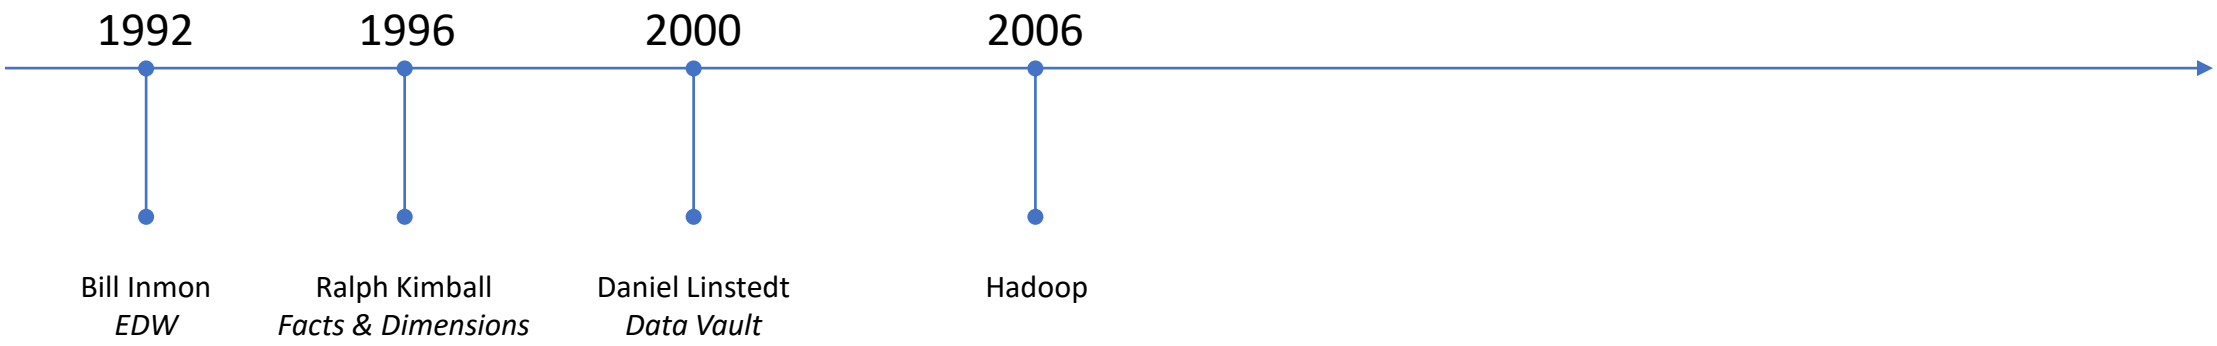

Batch processing, data warehouses …… data lakes

Timeline

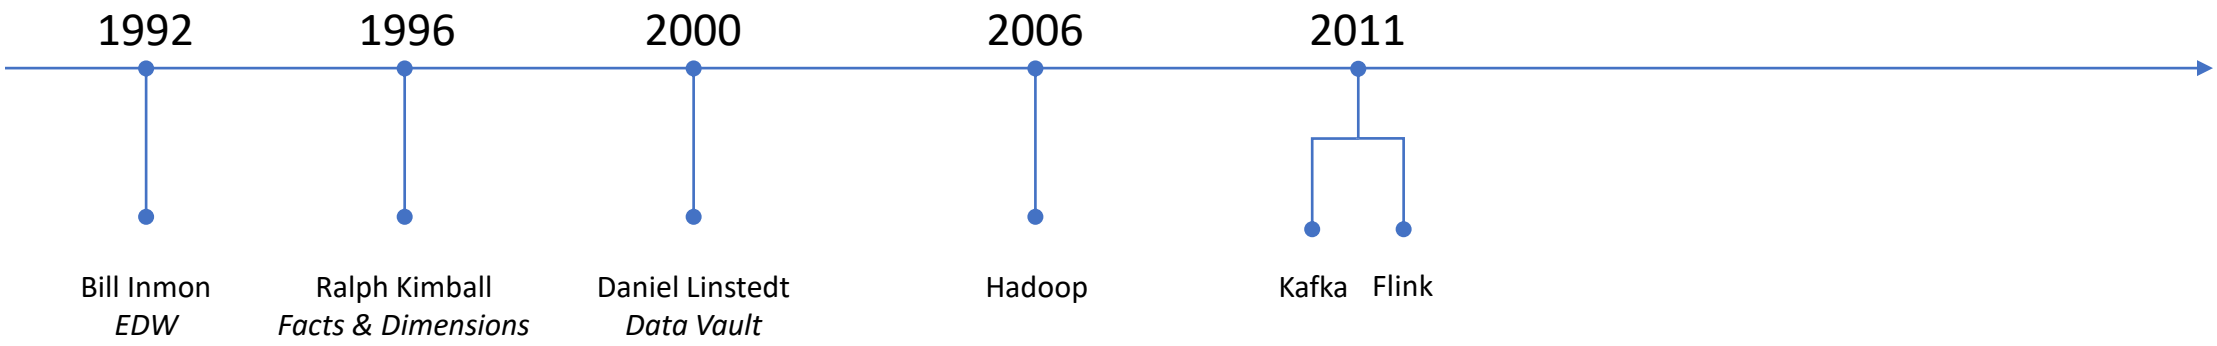

Batch processing, data warehouses …… data lakes …… stream processing

Timeline

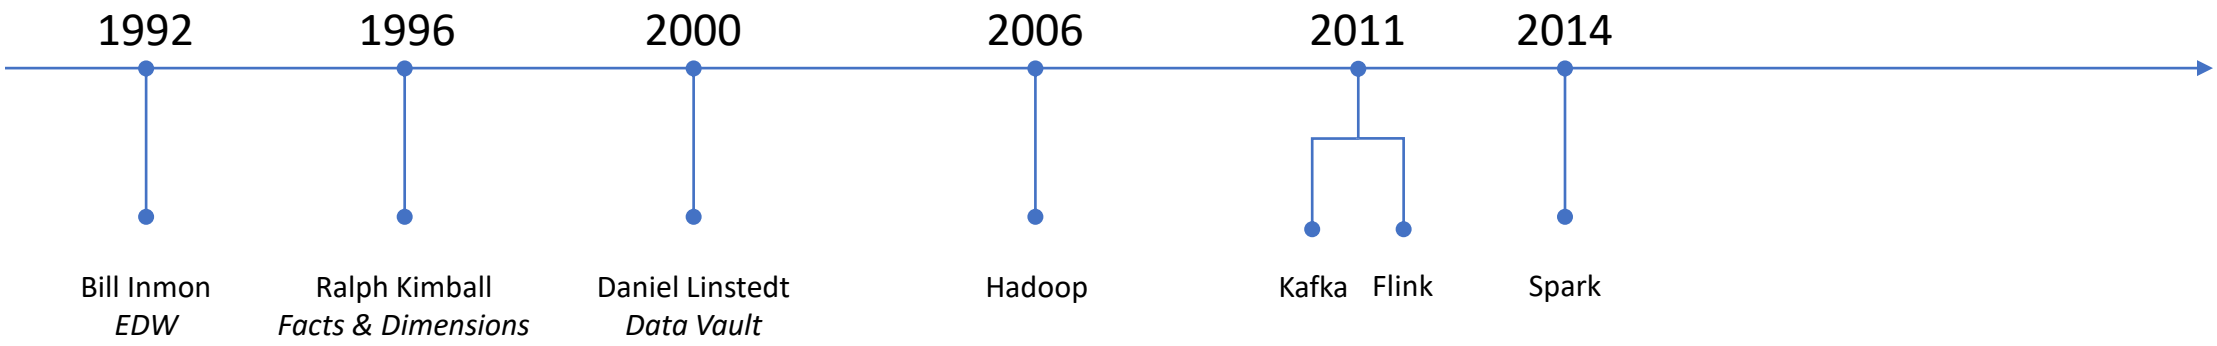

Batch processing, data warehouses …… data lakes …… stream processing

Timeline

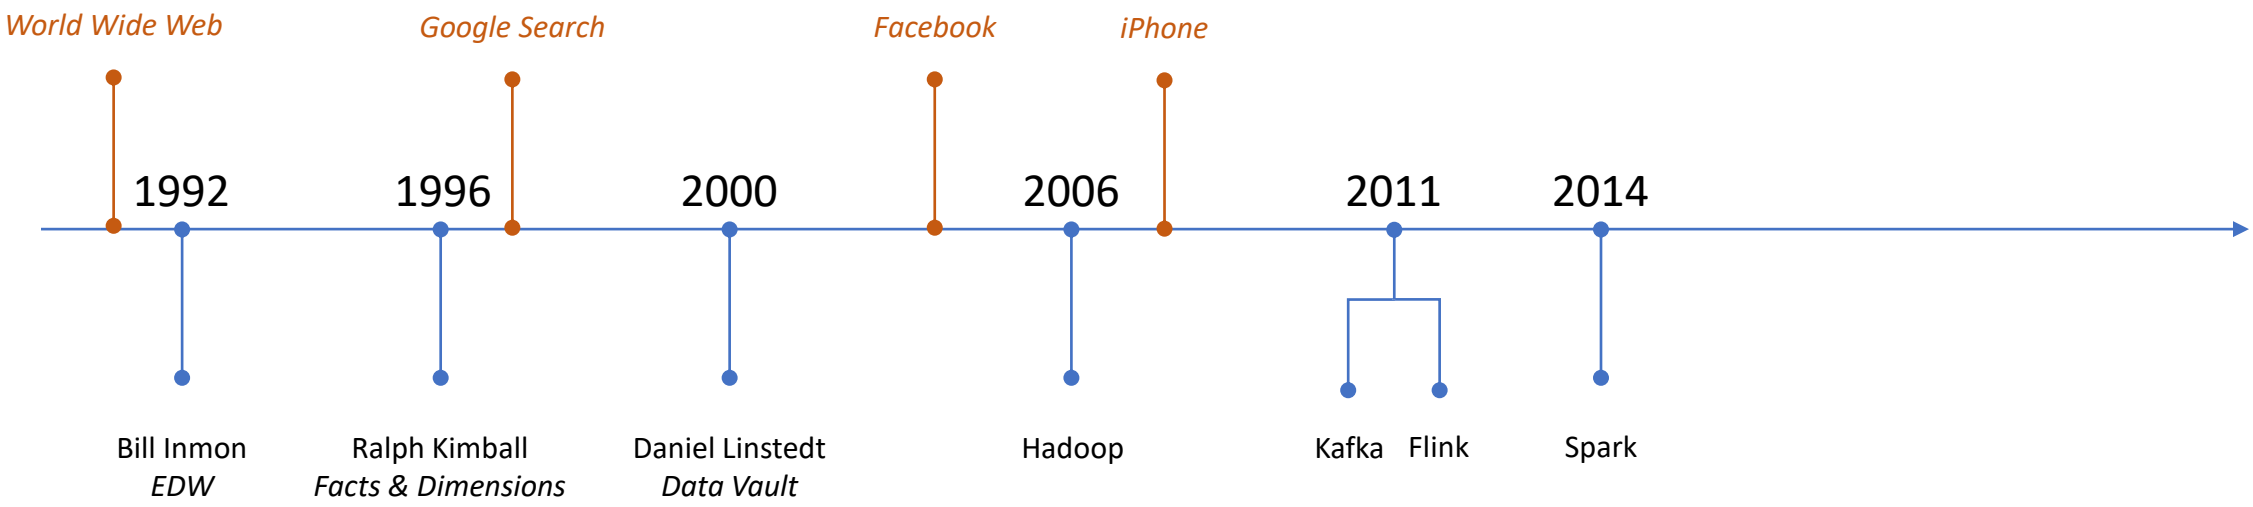

Batch processing, data warehouses …… data lakes …… stream processing

Role

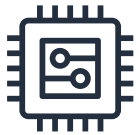

Role

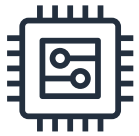

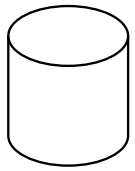

Role

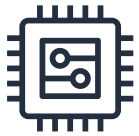

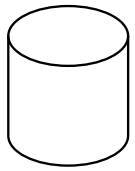

Role

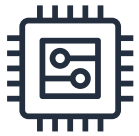

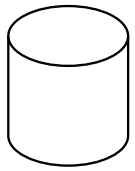

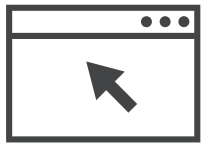

Role

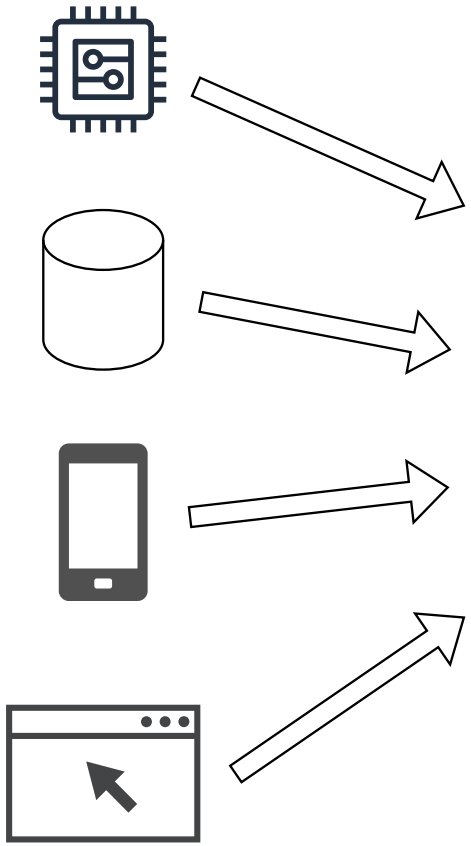

Role

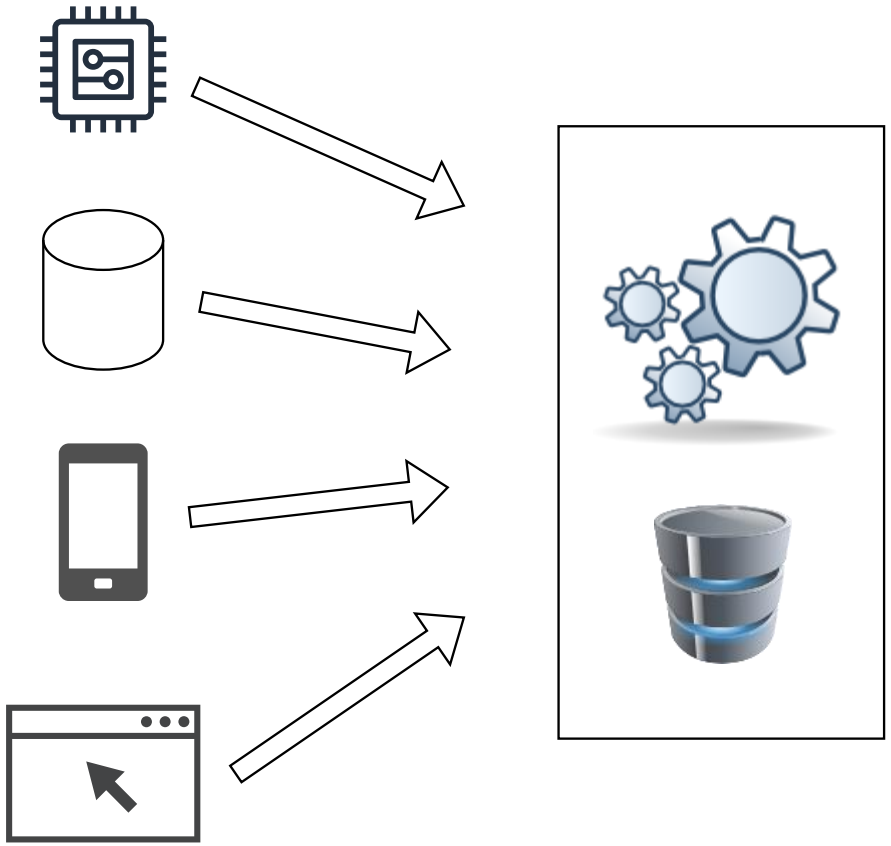

Role

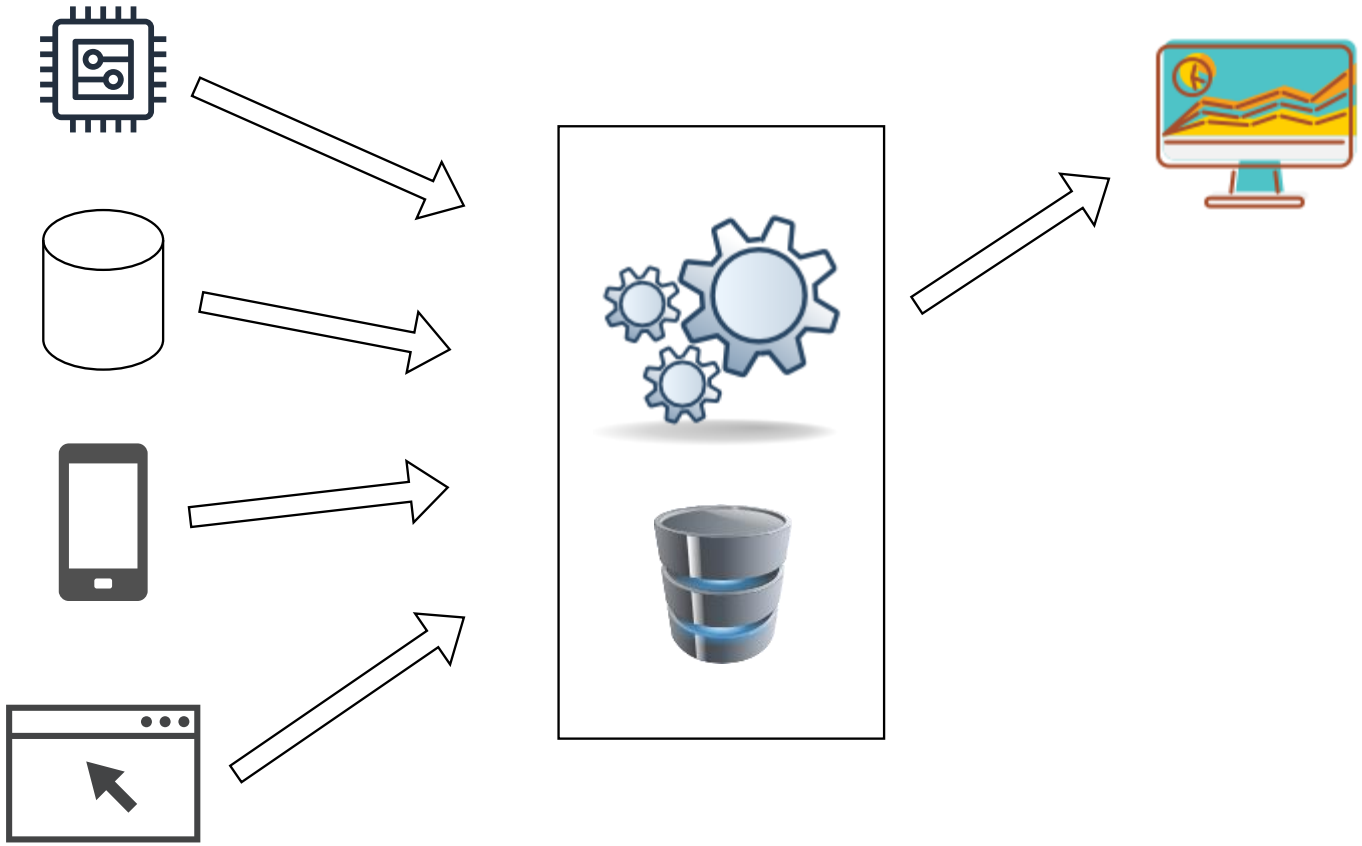

Role

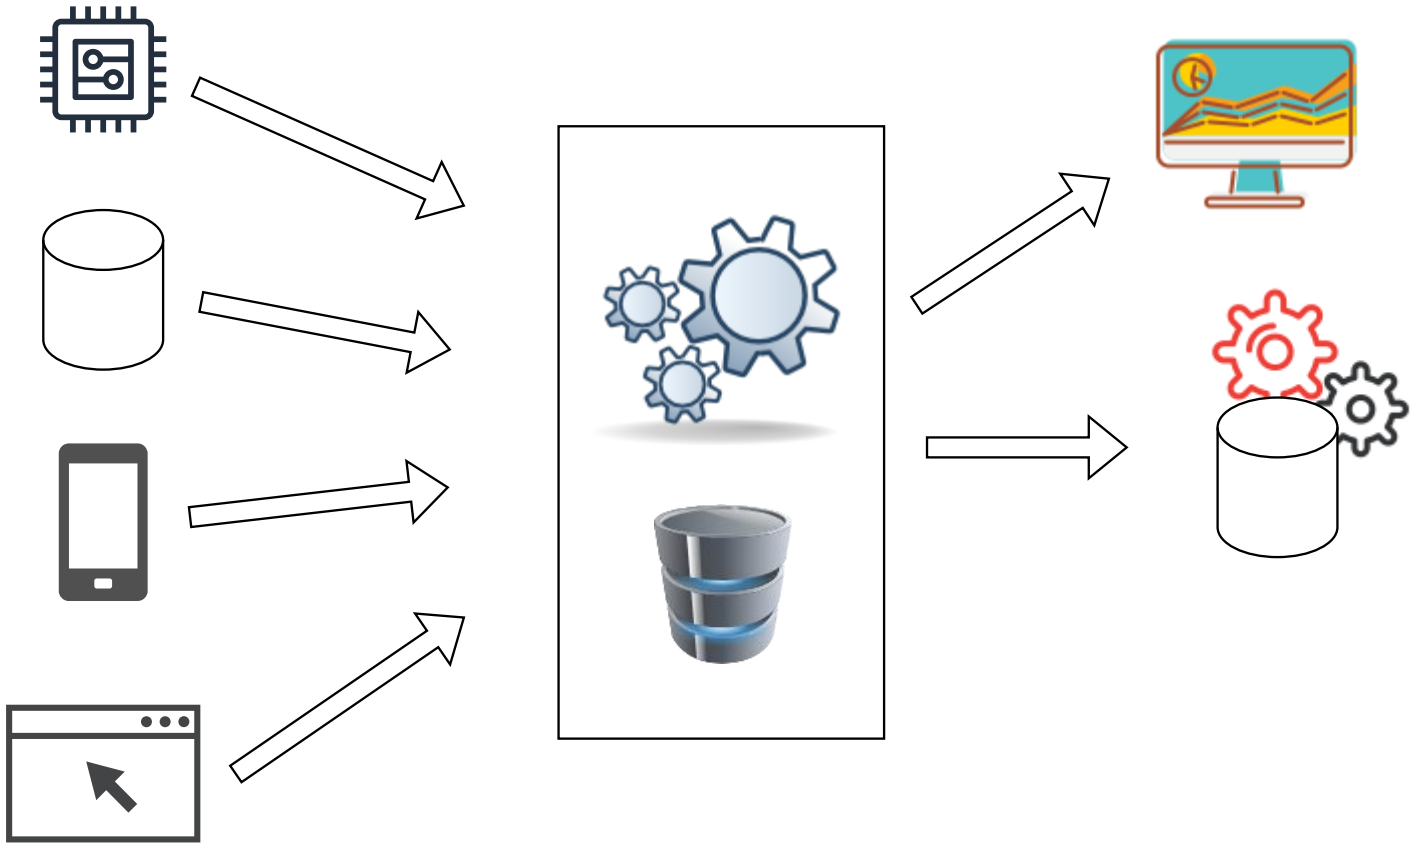

Role

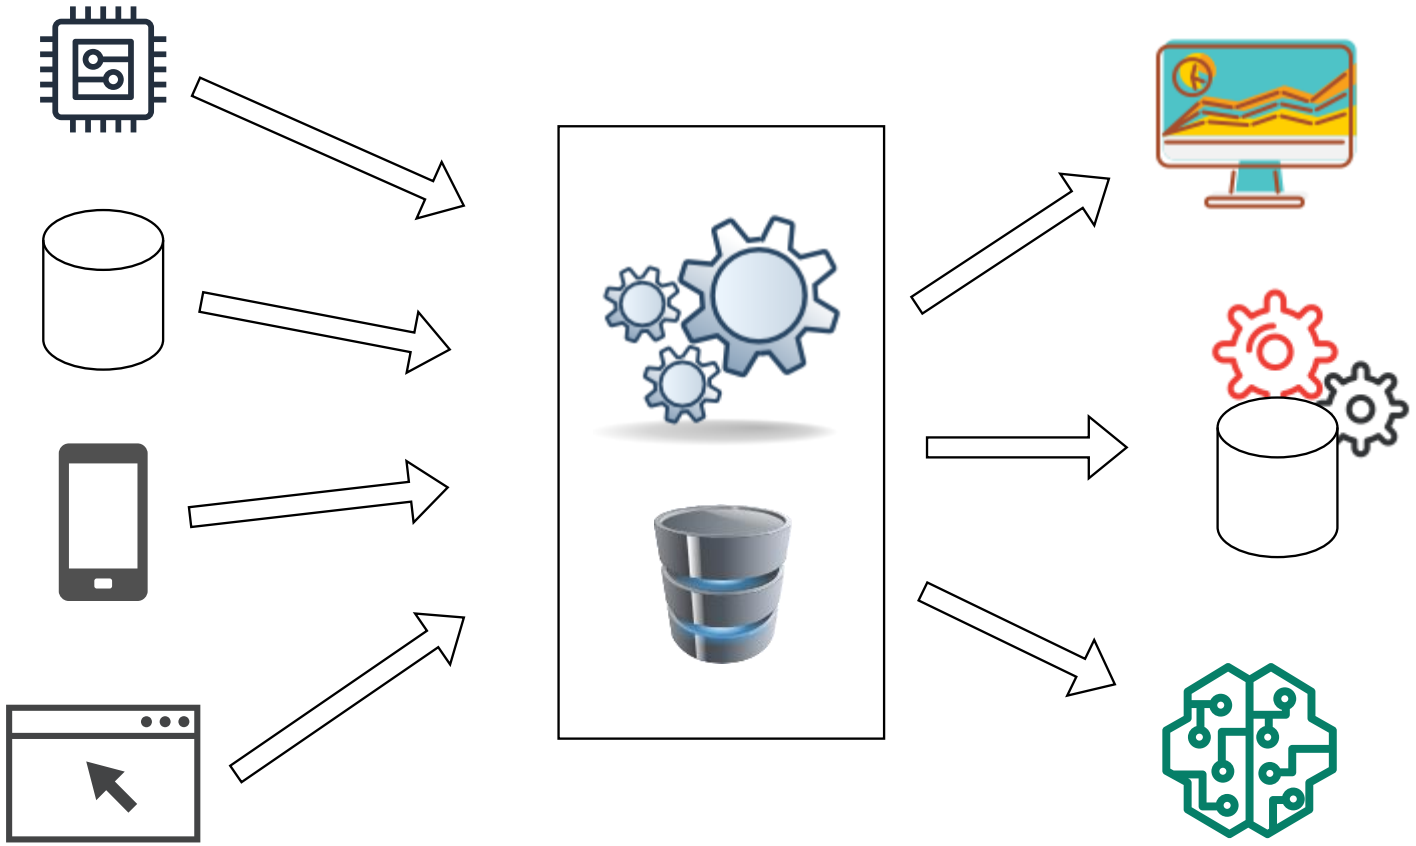

Data architecture

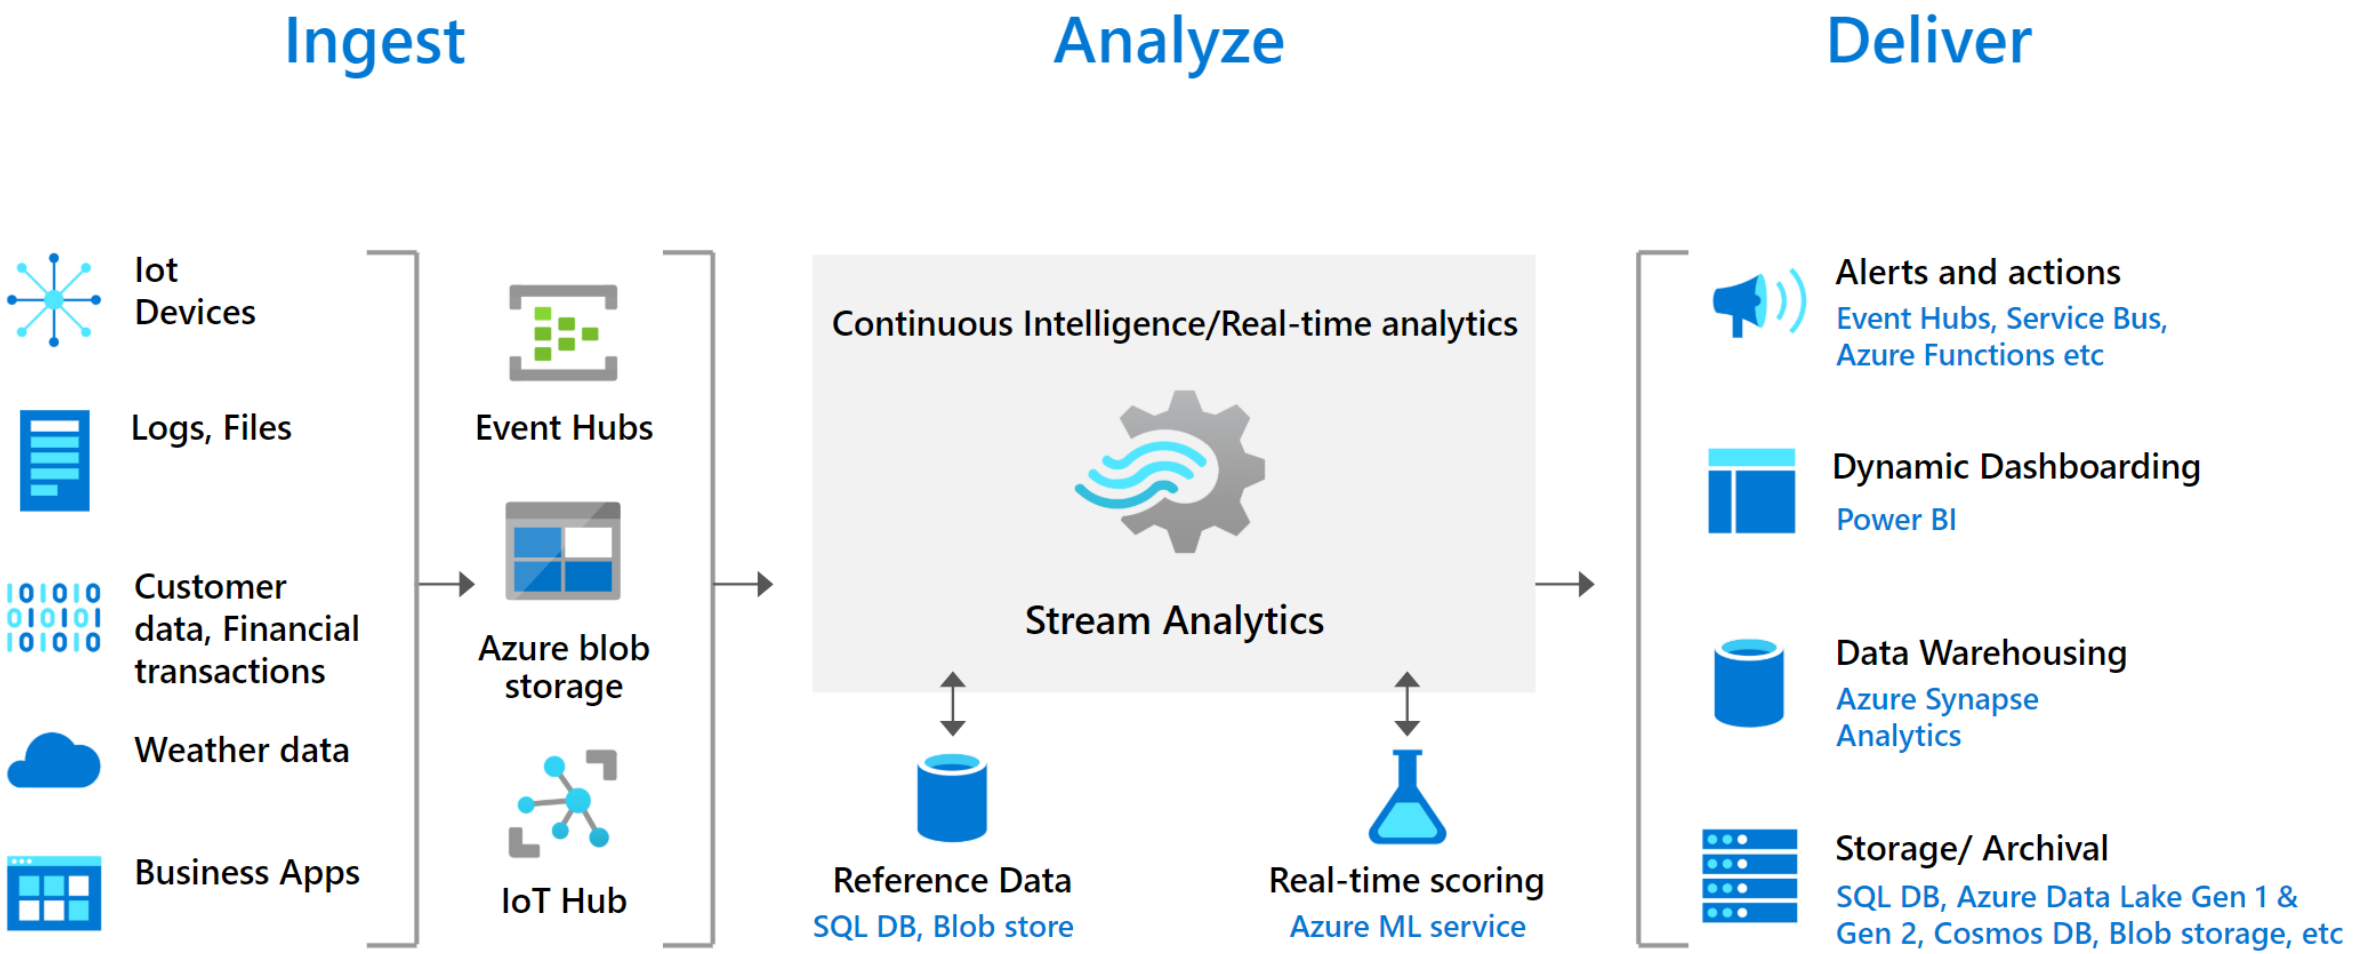

Source: https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-introduction

Accepted Latency

Time Dimension v/s SLA

SLA is order of Hours / Day

Batch

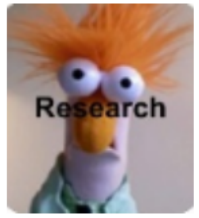

- Pre generated reports
- Cross grain resolution - trends

Source: Practical Real-time Data Processing and Analytics (2017)

Accepted Latency

Time Dimension v/s SLA

SLA is order of Hours / Day

Batch

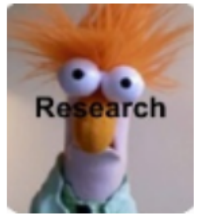

- Pre generated reports
- Cross grain resolution - trends

SLA is of order of Mins / Hour

Near Real Time

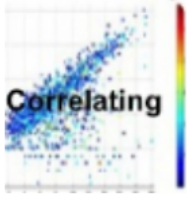

- Adhoc queries
- Mid resolution – aggregated counters

Source: Practical Real-time Data Processing and Analytics (2017)

Accepted Latency

Time Dimension v/s SLA

SLA is order of Hours / Day

Batch

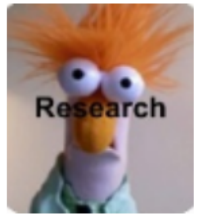

- Pre generated reports
- Cross grain resolution - trends

SLA is of order of Mins / Hour

Near Real Time

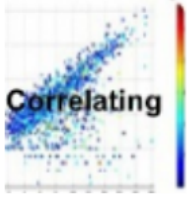

- Adhoc queries
- Mid resolution – aggregated counters

SLA is of order of Msec / Secs

Real Time

- Event Driven
- High resolution – each event counts

Source: Practical Real-time Data Processing and Analytics (2017)

Analytics

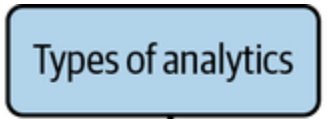

Source: Fundamentals of Data Engineering (2022)

Analytics

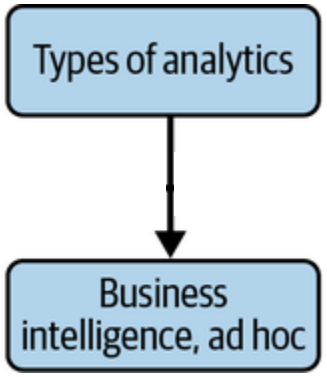

Source: Fundamentals of Data Engineering (2022)

Analytics

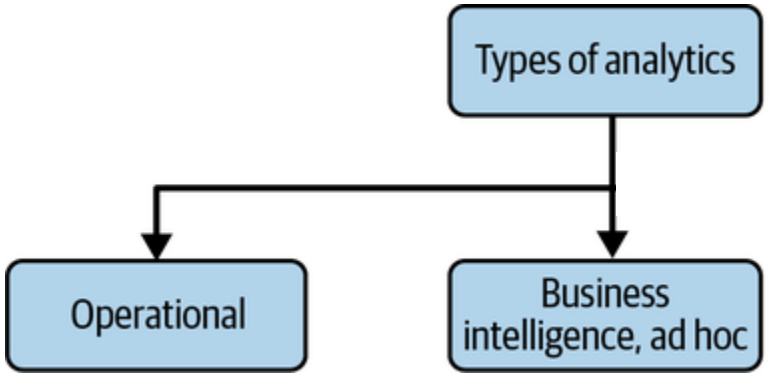

Source: Fundamentals of Data Engineering (2022)

Analytics

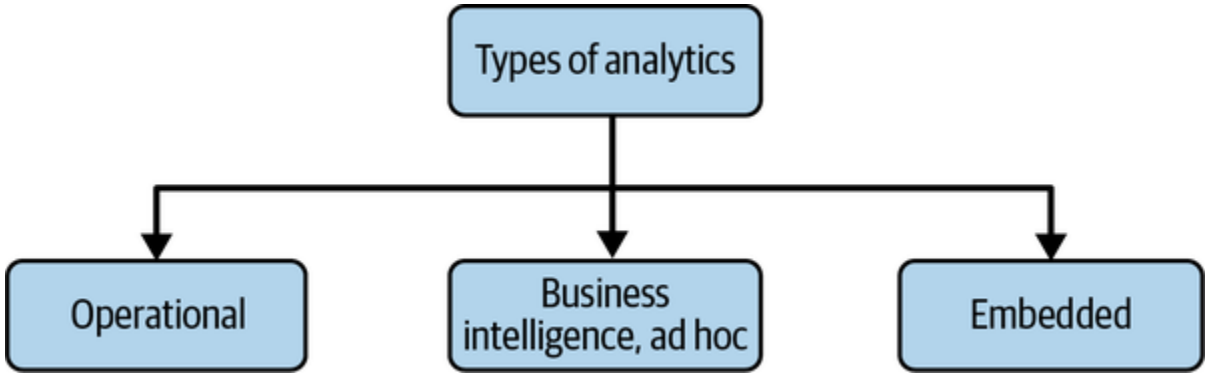

Source: Fundamentals of Data Engineering (2022)

Analytics

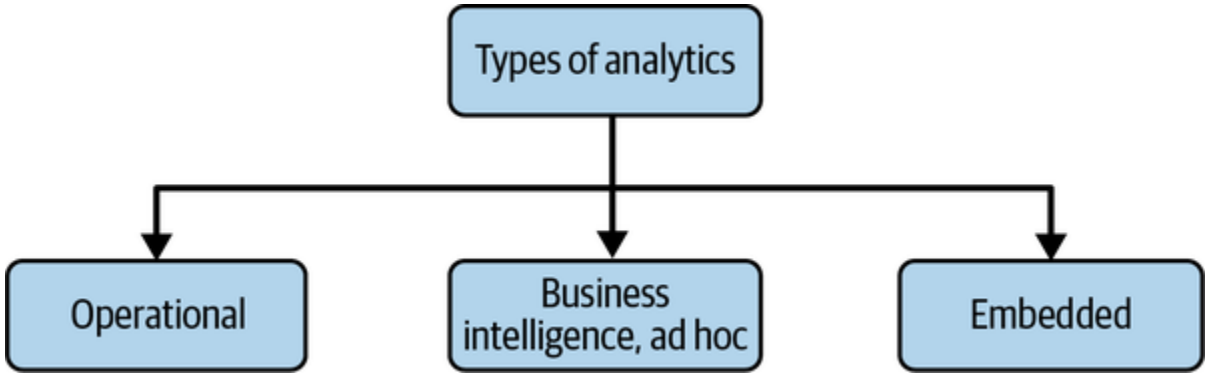

Live view

Past and present of business

Multi-tenant (outbound BI)

Source: Fundamentals of Data Engineering (2022)

Analytics

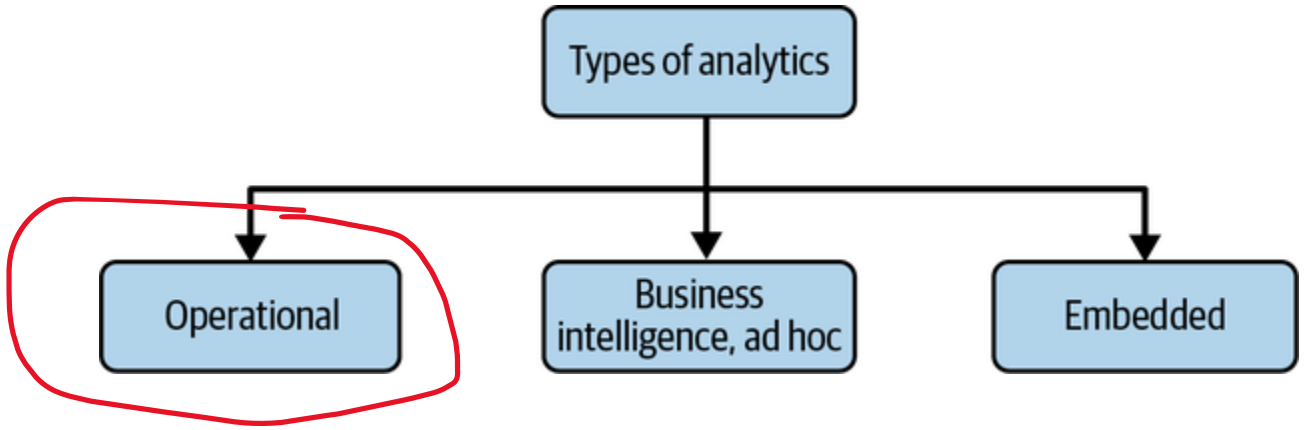

Live view

Past and present of business

Multi-tenant (outbound BI)

Source: Fundamentals of Data Engineering (2022)

Challenges

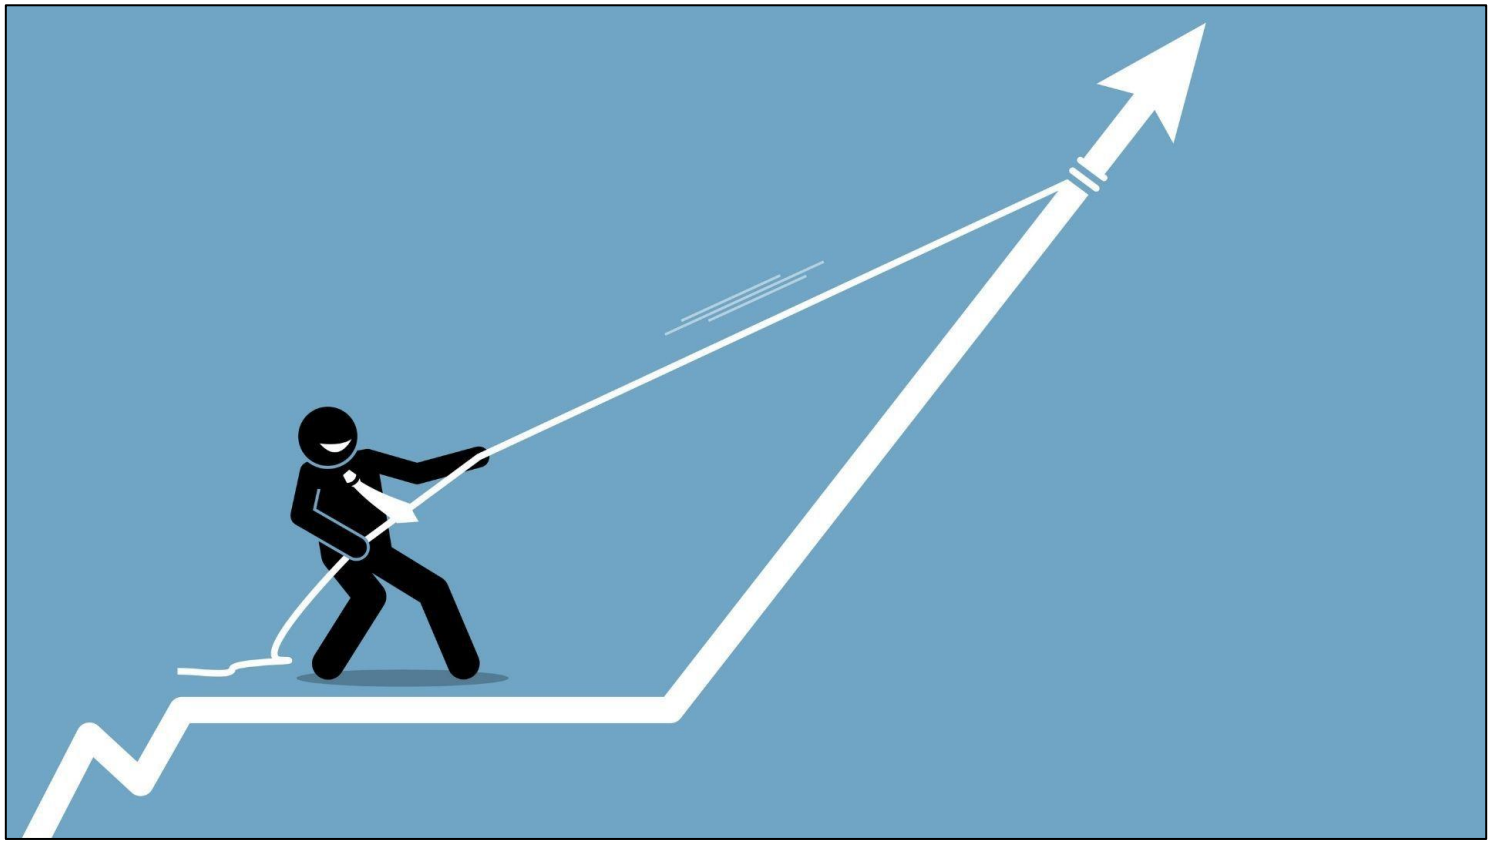

Scalability

Challenges

Ordering

Challenges

Consistency

Challenges

Fault Tolerance

Recap

- Data generation boom
- Data processing shift
  - From batch to streaming
- We need a live view
  - Operational analytics
  - Near real-time (seconds/minutes latency)
- Challenges
  - Scalability – Ordering – Consistency – Fault Tolerance

Not discussed but important

- Reverse ETL
- Machine Learning
- DataOps
- Security
- Data Management

Kafka

Messaging

Messaging

Messaging

- Hello, would you like to hear a TCP joke?

Messaging

- Hello, would you like to hear a TCP joke?
- Hello, yes, I’d like to hear a TCP joke.

Messaging

- Hello, would you like to hear a TCP joke?
- Hello, yes, I’d like to hear a TCP joke.
- OK, I’ll tell you a TCP joke.

Messaging

- Hello, would you like to hear a TCP joke?
- Hello, yes, I'd like to hear a TCP joke.
- OK, I'll tell you a TCP joke.
- OK, I'll hear a TCP joke.

Kafka motivation

Traditional
Message-Queue

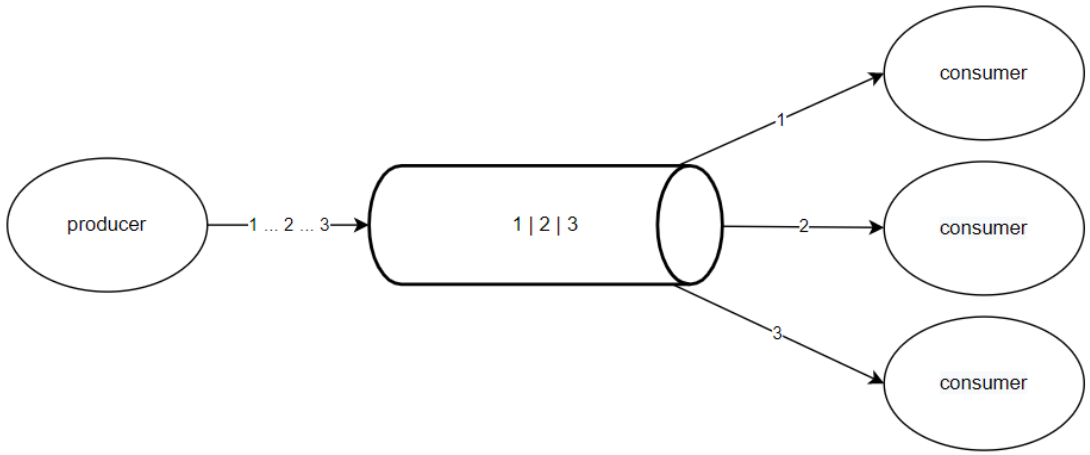

Kafka motivation

Traditional Message-Queue

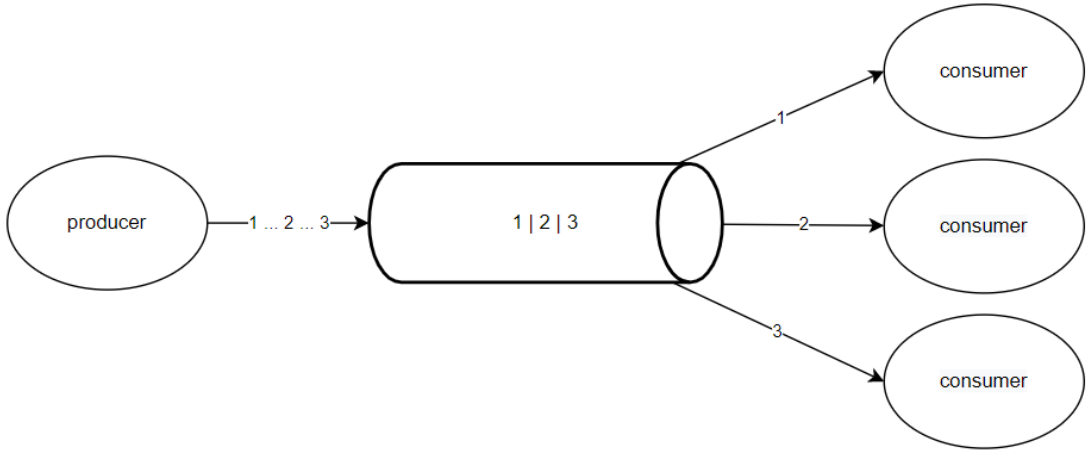

Kafka

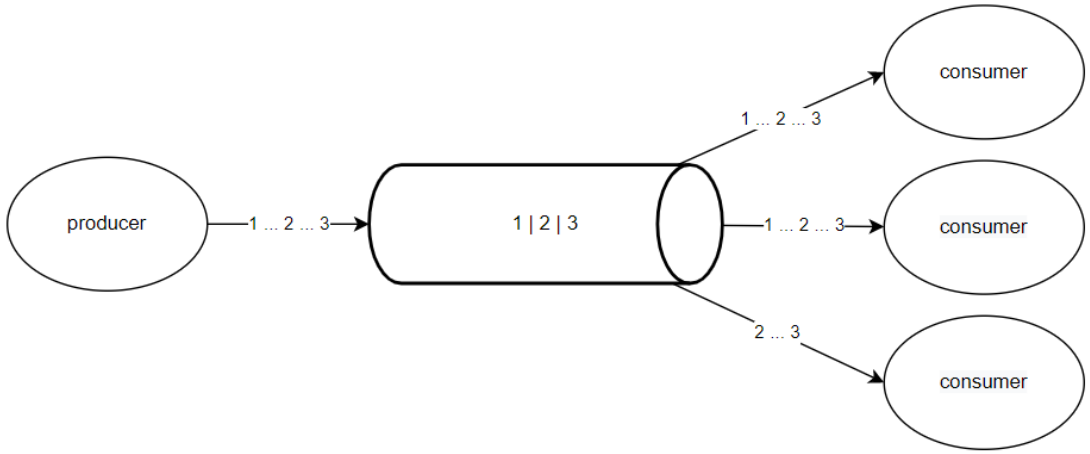

- Distributed streaming platform
  - Publish & subscribe
  - Store streams durably
  - Process streams as they occur

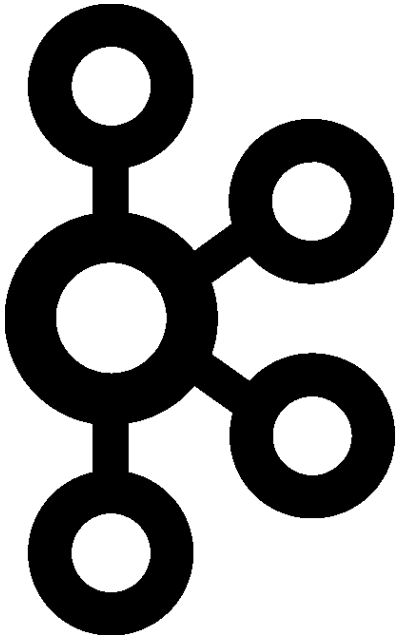

kafka

Kafka cluster

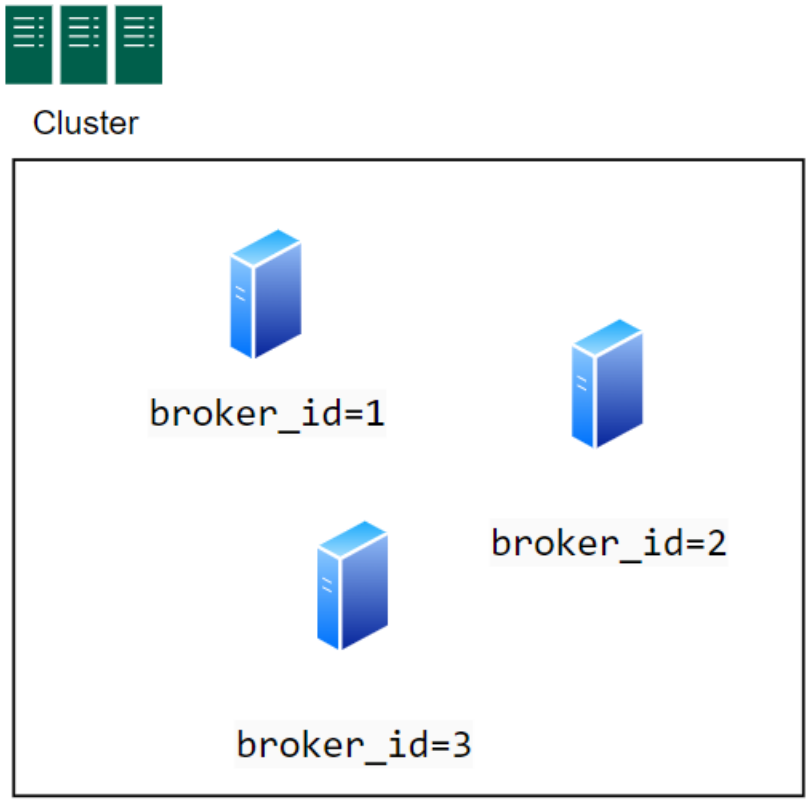

Kafka topic

The core abstraction Kafka provides for a stream of records is the **topic**.
A topic is a category or feed name to which records are published.

Kafka topic

producers

E-shop

Retail

Kafka topic

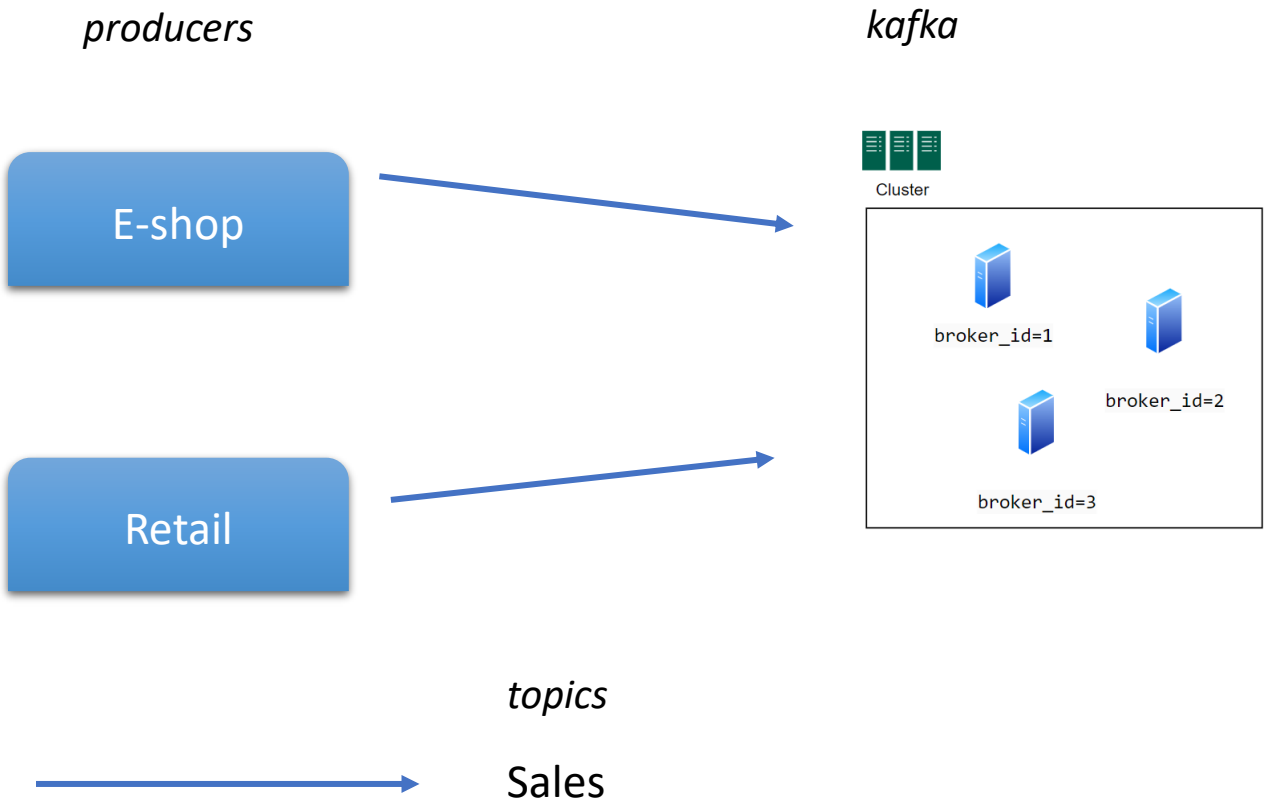
producers
kafka
E-shop
Retail
Cluster
broker_id=1
broker_id=2
broker_id=3
topics
Sales

Kafka topic

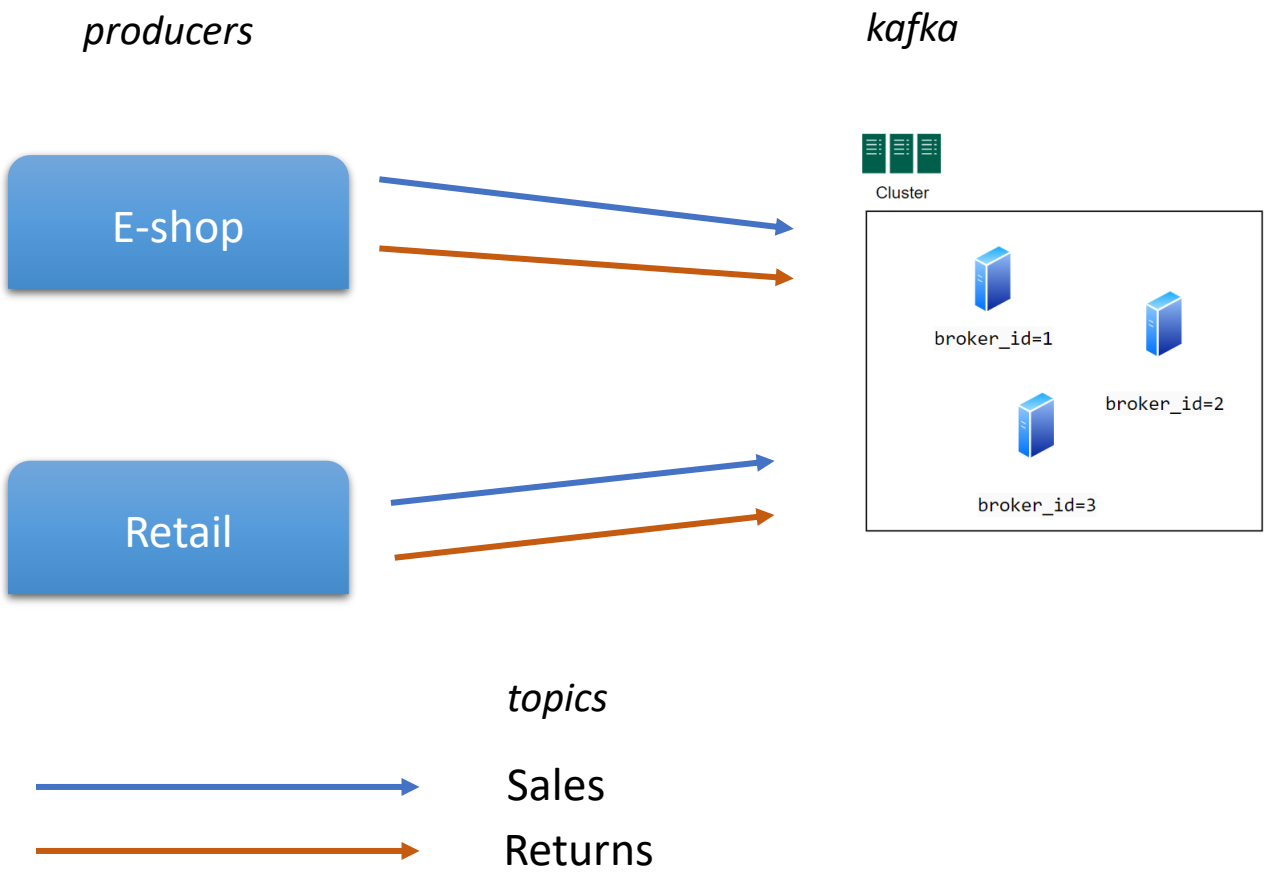
producers
kafka
E-shop
Retail
Cluster
broker_id=1
broker_id=2
broker_id=3
topics
Sales
Returns

Kafka topic

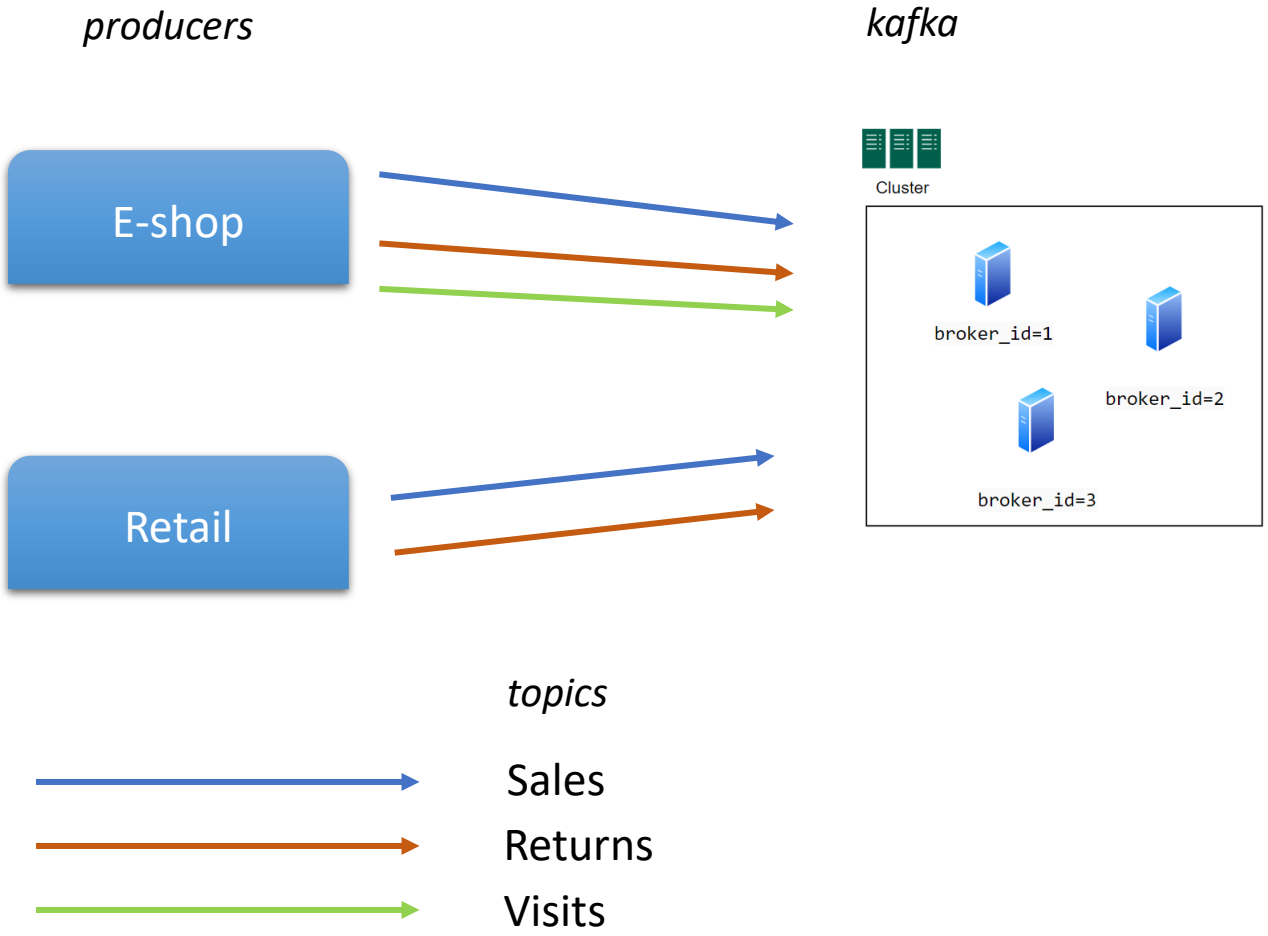
producers
kafka
E-shop
Retail
Cluster
broker_id=1
broker_id=2
broker_id=3
topics
Sales
Returns
Visits

Kafka topic

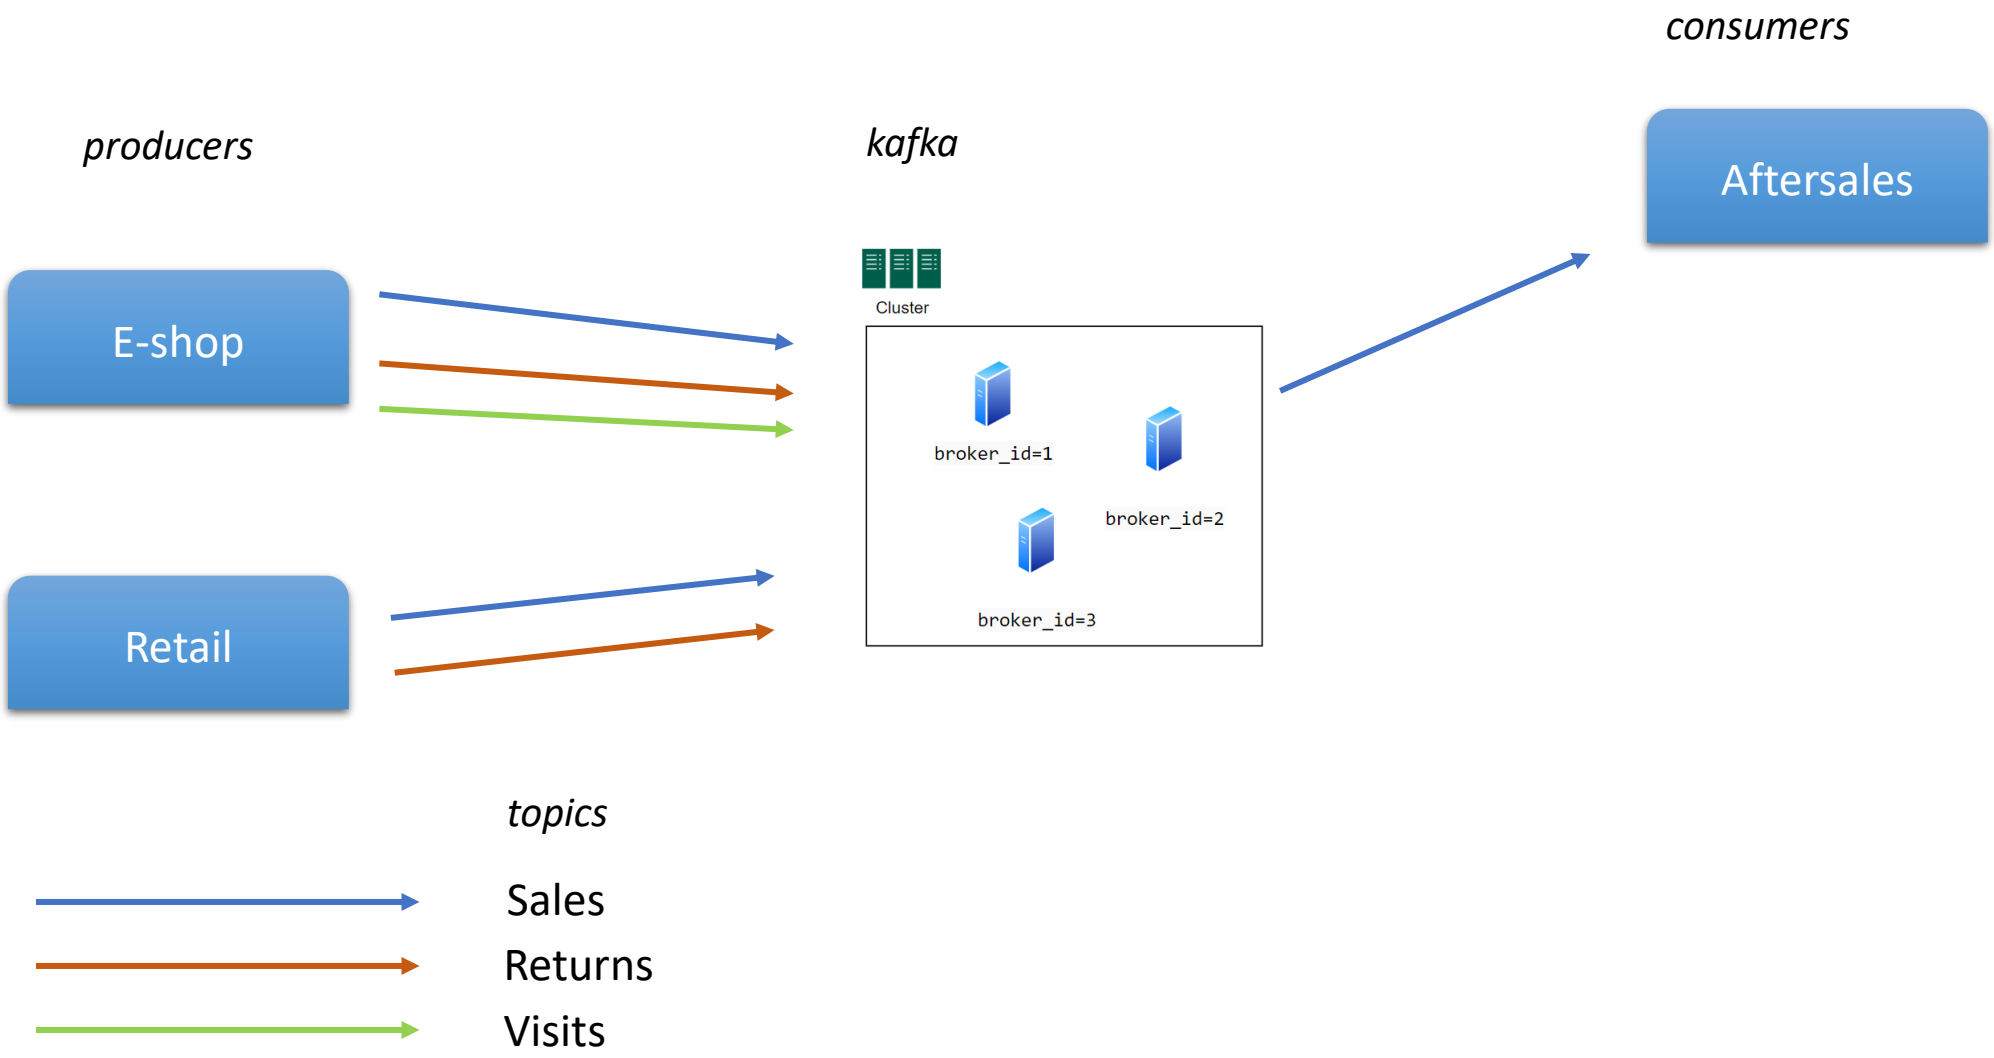
consumers
producers
kafka
Aftersales
E-shop
Cluster
broker_id=1
broker_id=2
broker_id=3
Retail
topics
Sales
Returns
Visits

Kafka topic

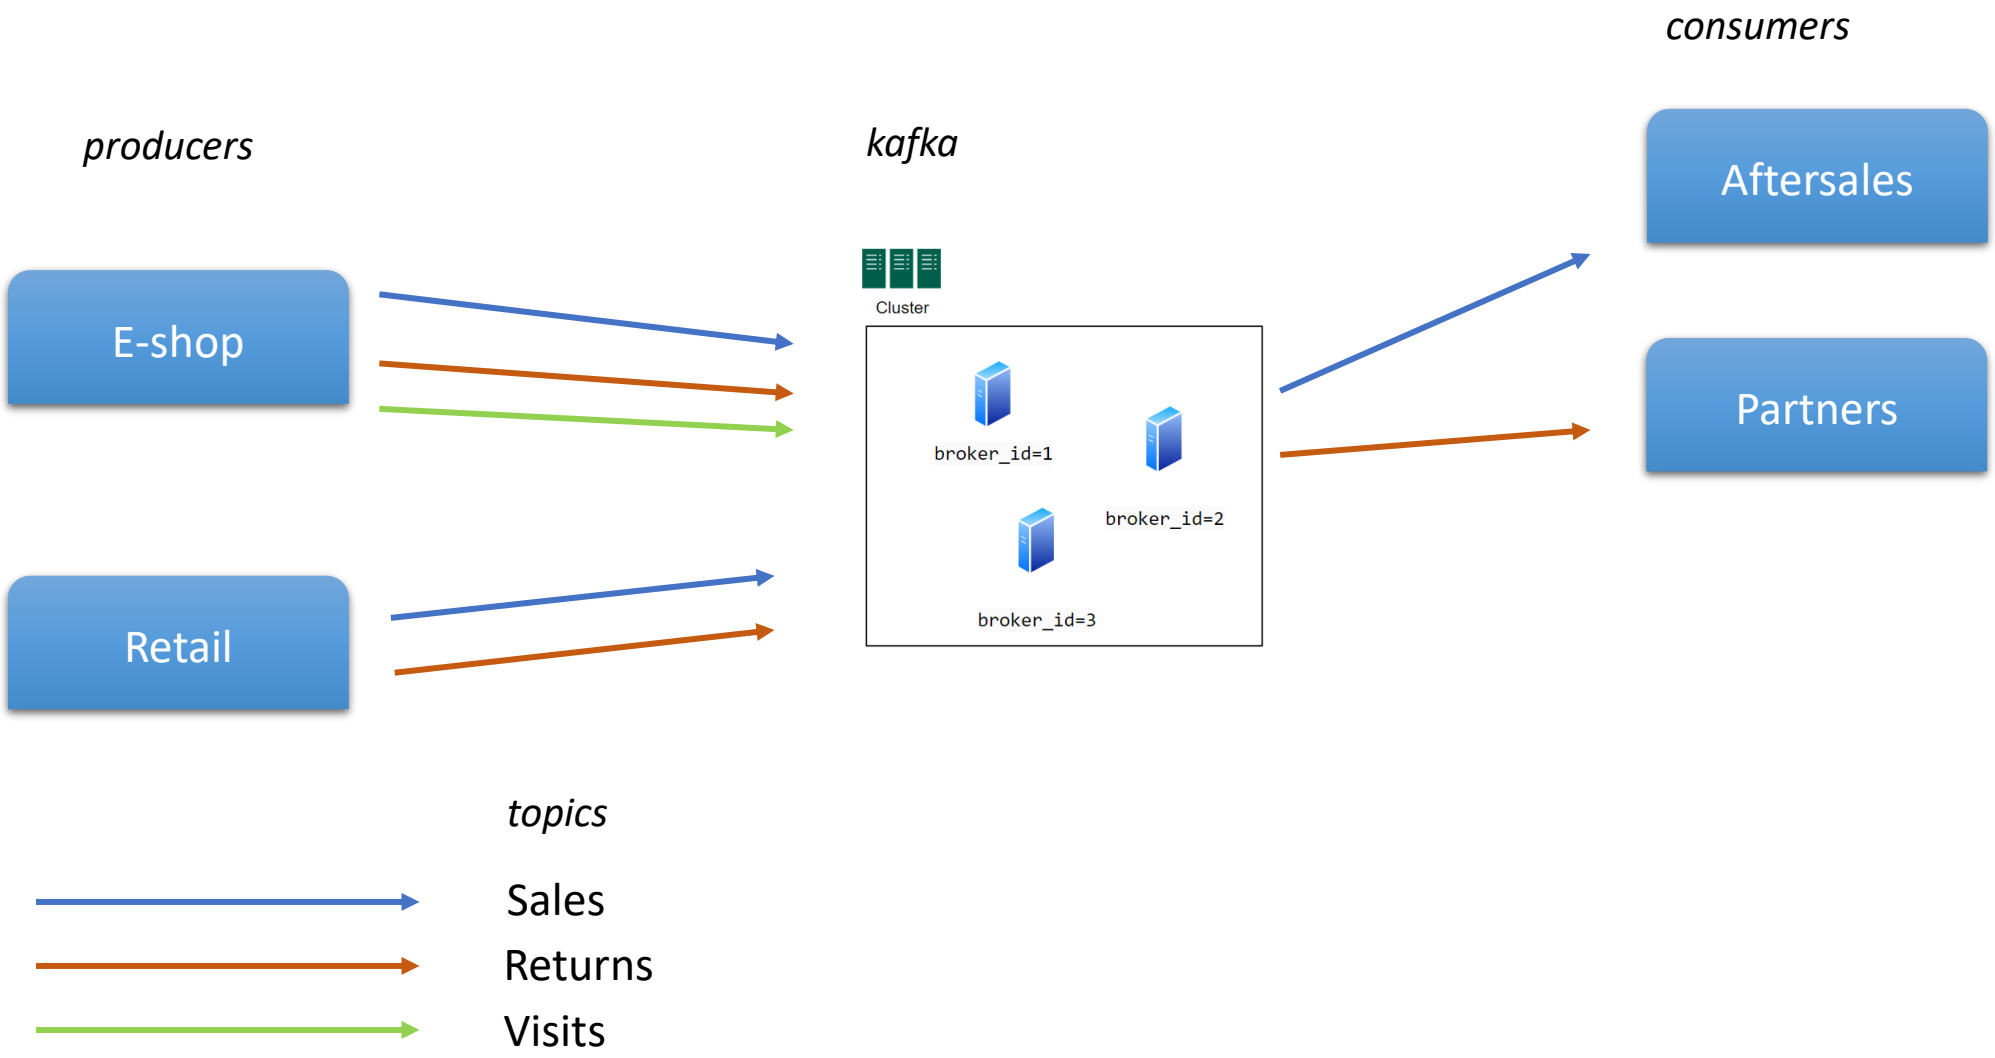
consumers
producers
kafka
Aftersales
E-shop
Cluster
broker_id=1
broker_id=2
broker_id=3
Partners
Retail
topics
Sales
Returns
Visits

Kafka topic

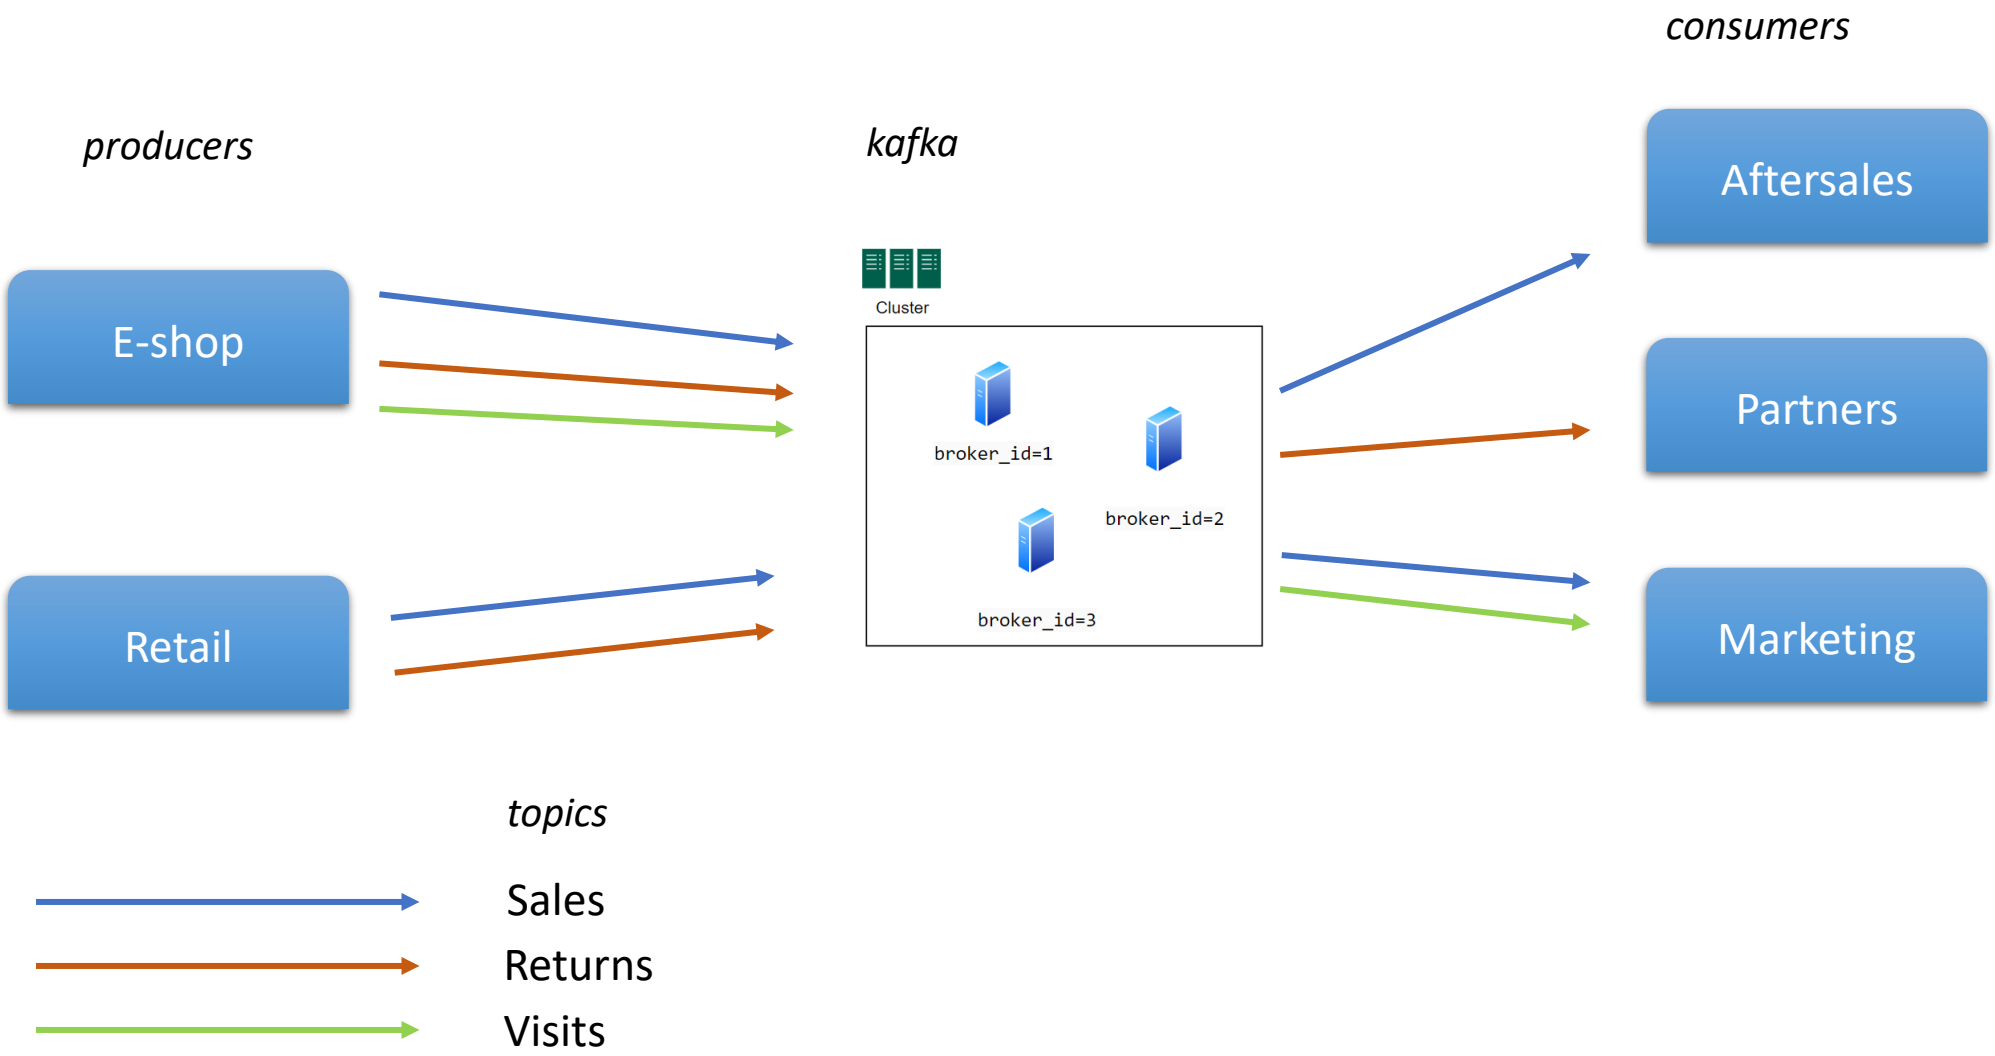
consumers
producers
kafka
Aftersales
E-shop
Cluster
Partners
broker_id=1
broker_id=2
Retail
broker_id=3
Marketing
topics
Sales
Returns
Visits

Kafka topic

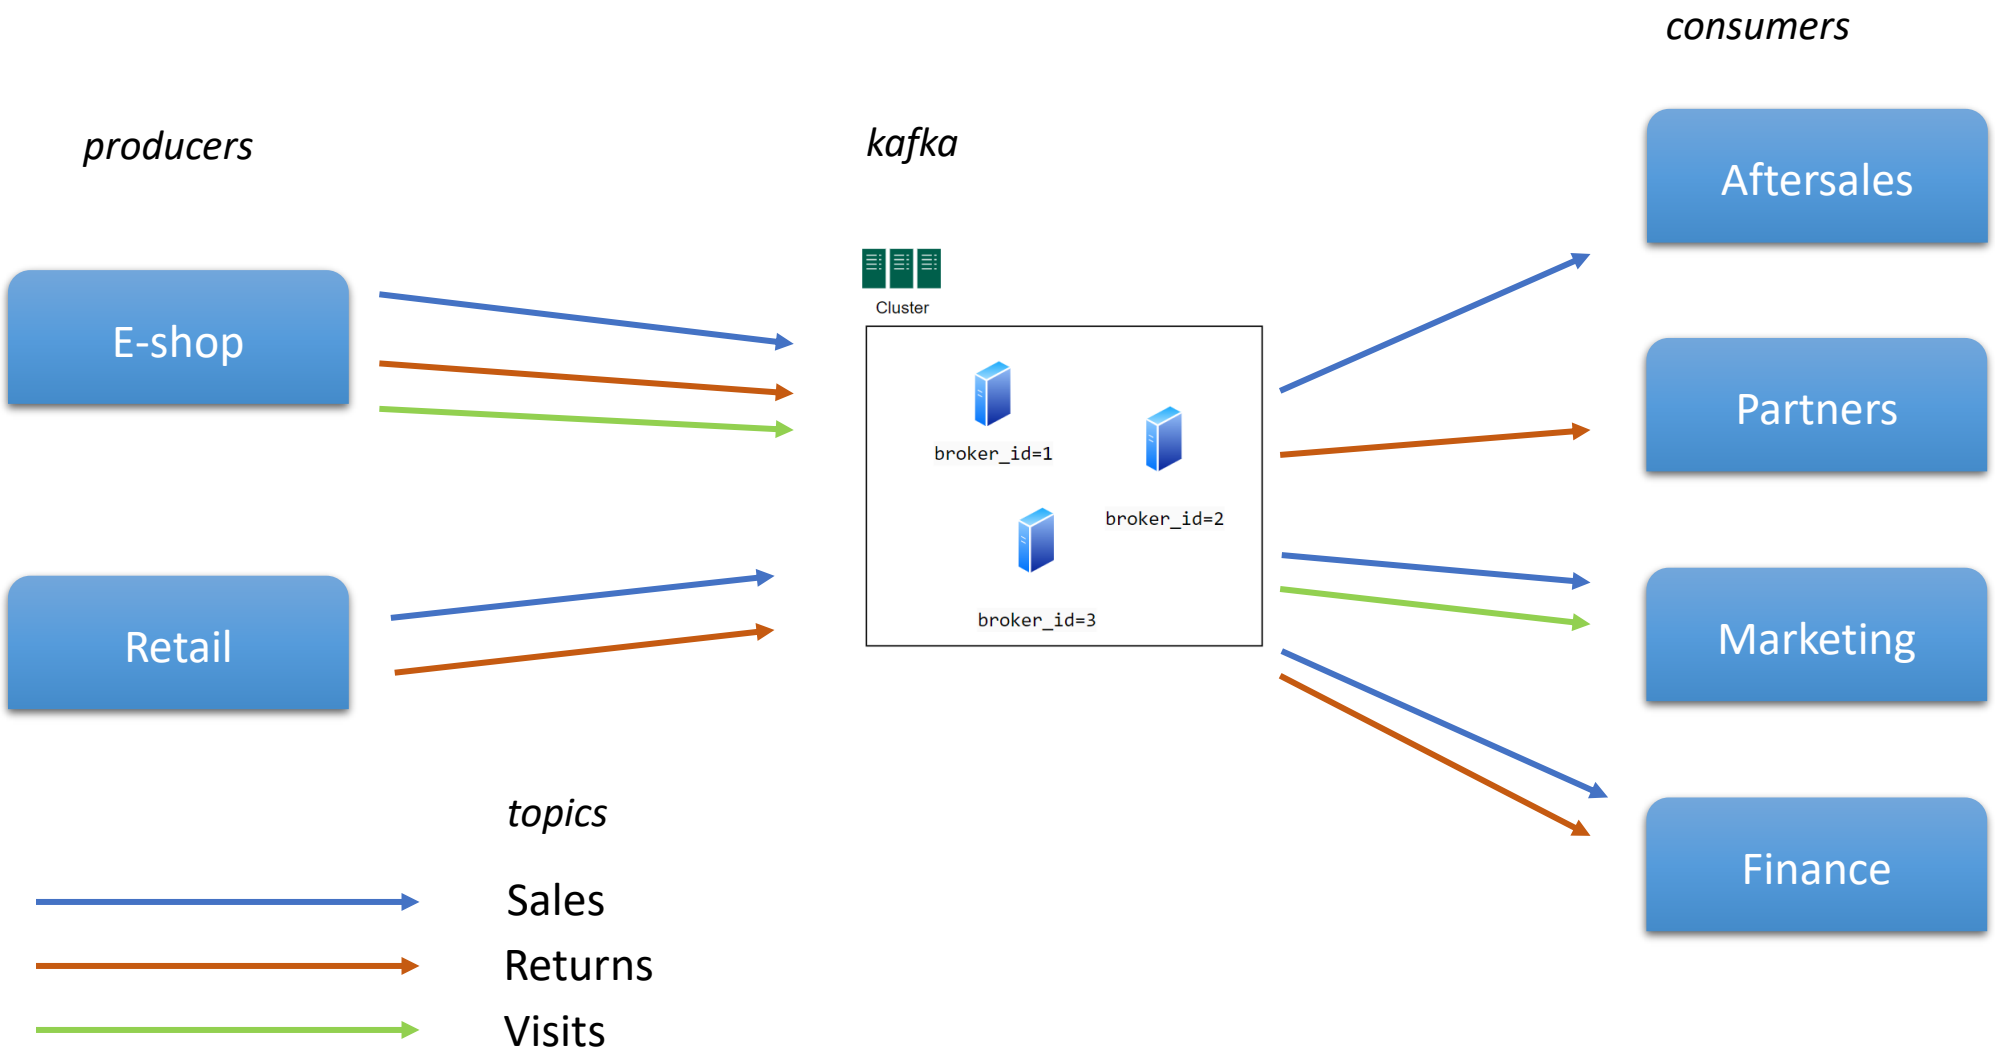
consumers
producers
kafka
Aftersales
E-shop
Cluster
Partners
broker_id=1
broker_id=2
Retail
Marketing
broker_id=3
Finance
topics
Sales
Returns
Visits

Kafka message?

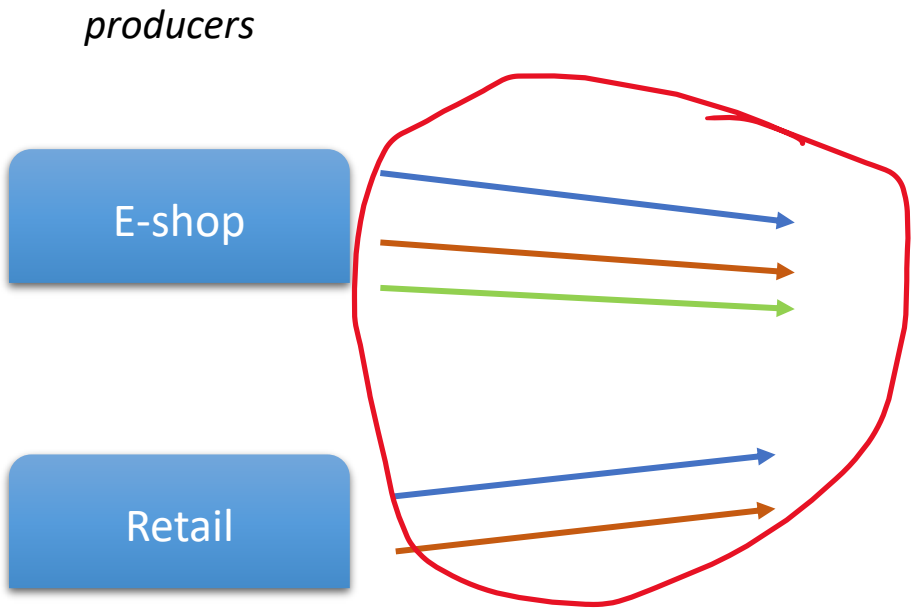

Kafka message?

- (Key)
- Value

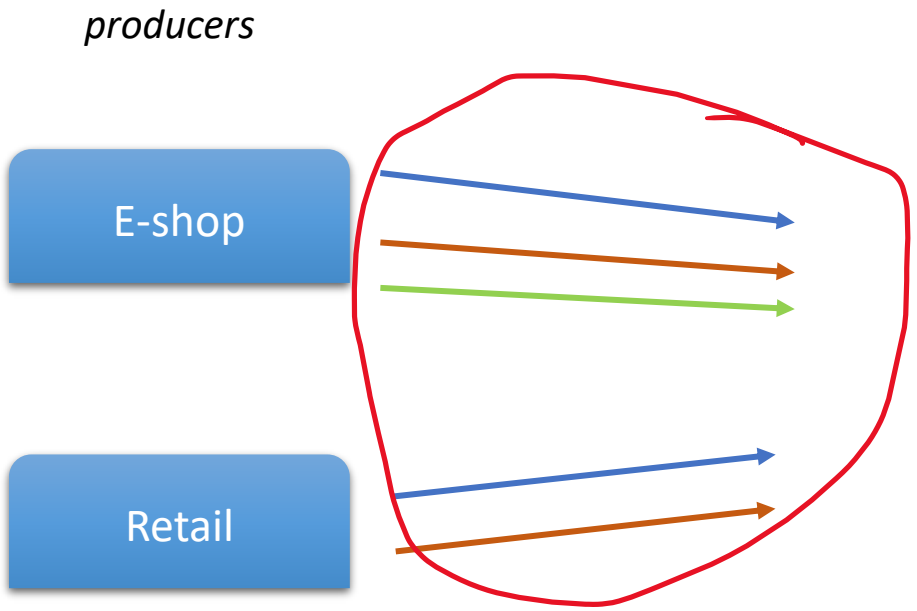

Kafka partitions

- Scalability

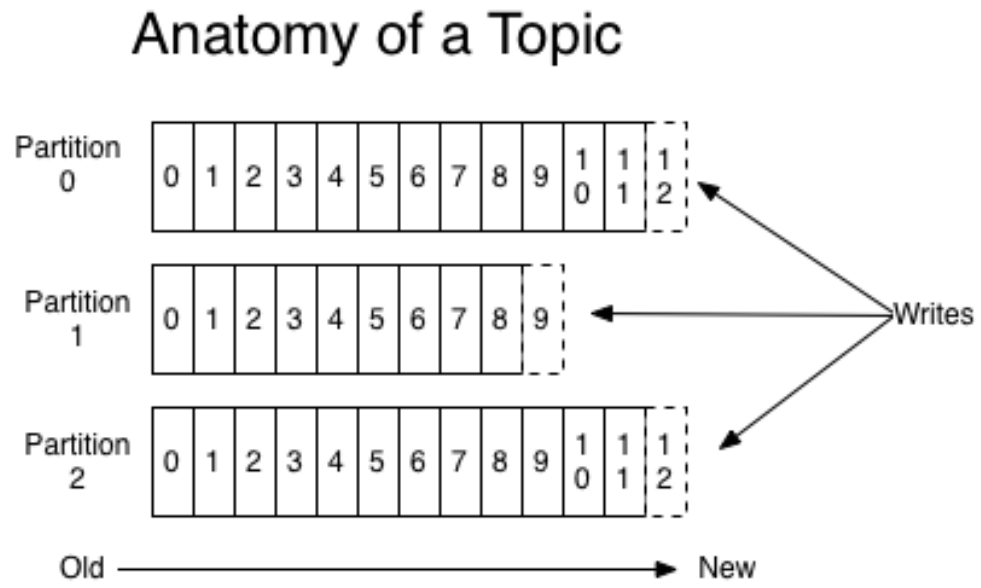

Kafka replication

- Fault-tolerance: if a broker is down, another broker can serve the data

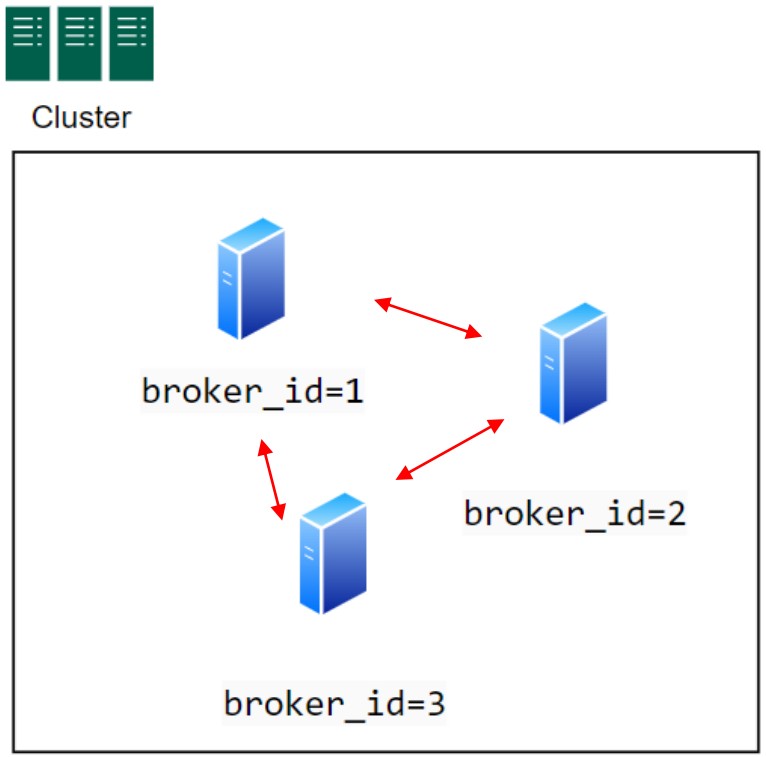

Kafka replication

- Fault-tolerance: if a broker is down, another broker can serve the data

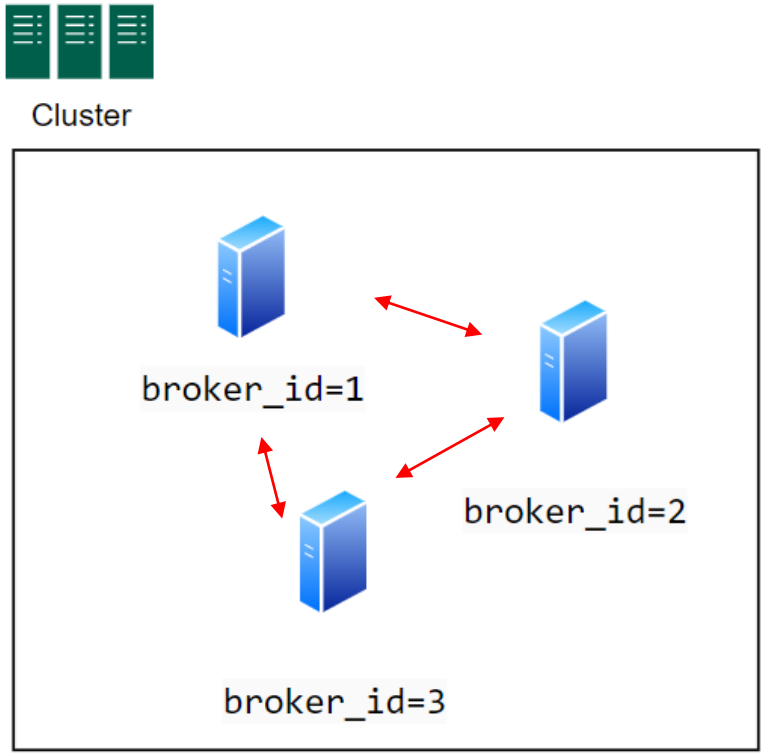

ISR = In-Sync Replica

Example

Kafka core

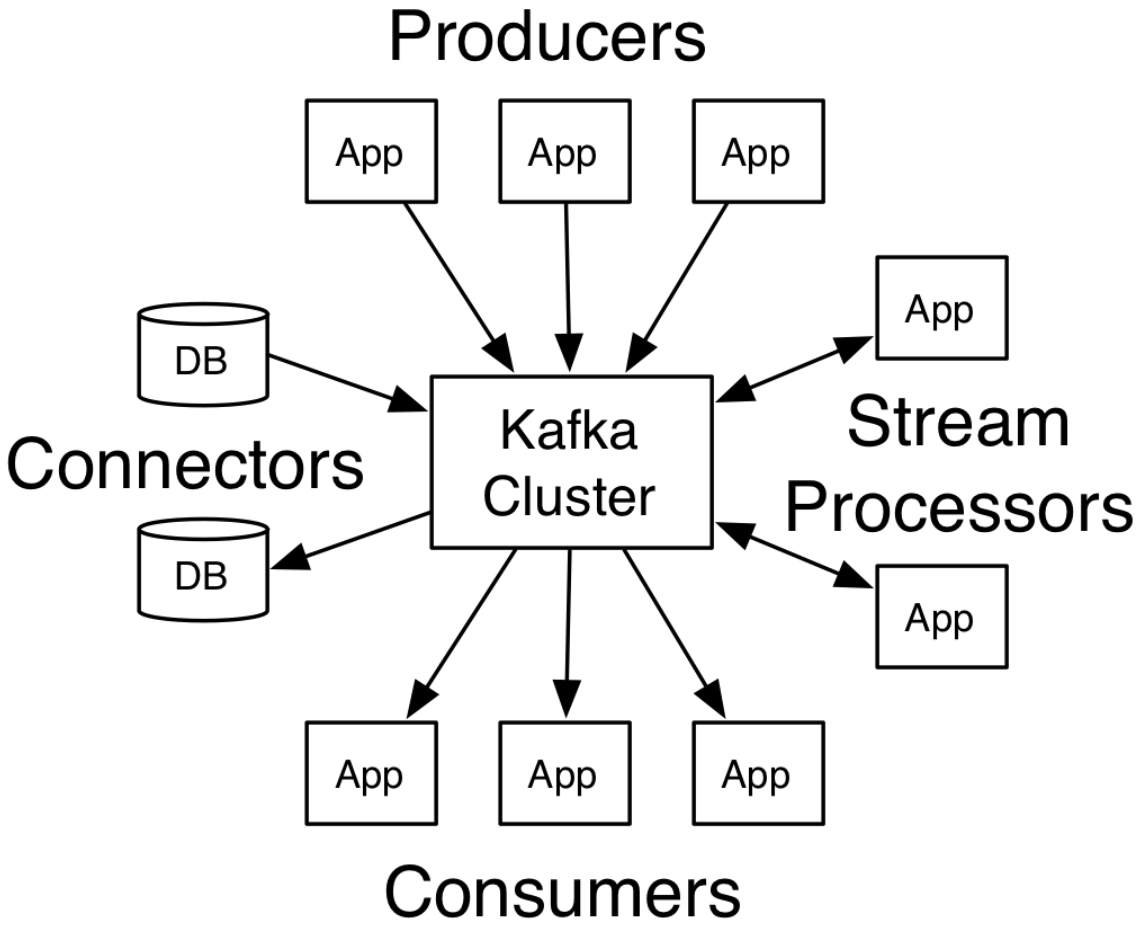
Producers
App
App
App
DB
Connectors
DB
Kafka
Cluster
App
Stream
Processors
App
App
App
App
Consumers

Kafka extended

- Schema registry
- Kafka Connect
  - CDC (Change Data Capture)
- Zookeeper, Kafka Raft (Kraft)
- ksqlDB

Let's Roll