Data Streaming Concepts and Tools

Kristo Raun

Introduction to Near Real-Time Data Analytics

August 2022

UNIVERSITY OF TARTU
1632

Agenda for the week

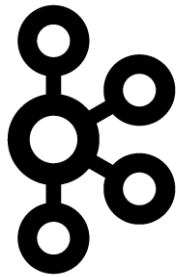

Agenda for the week

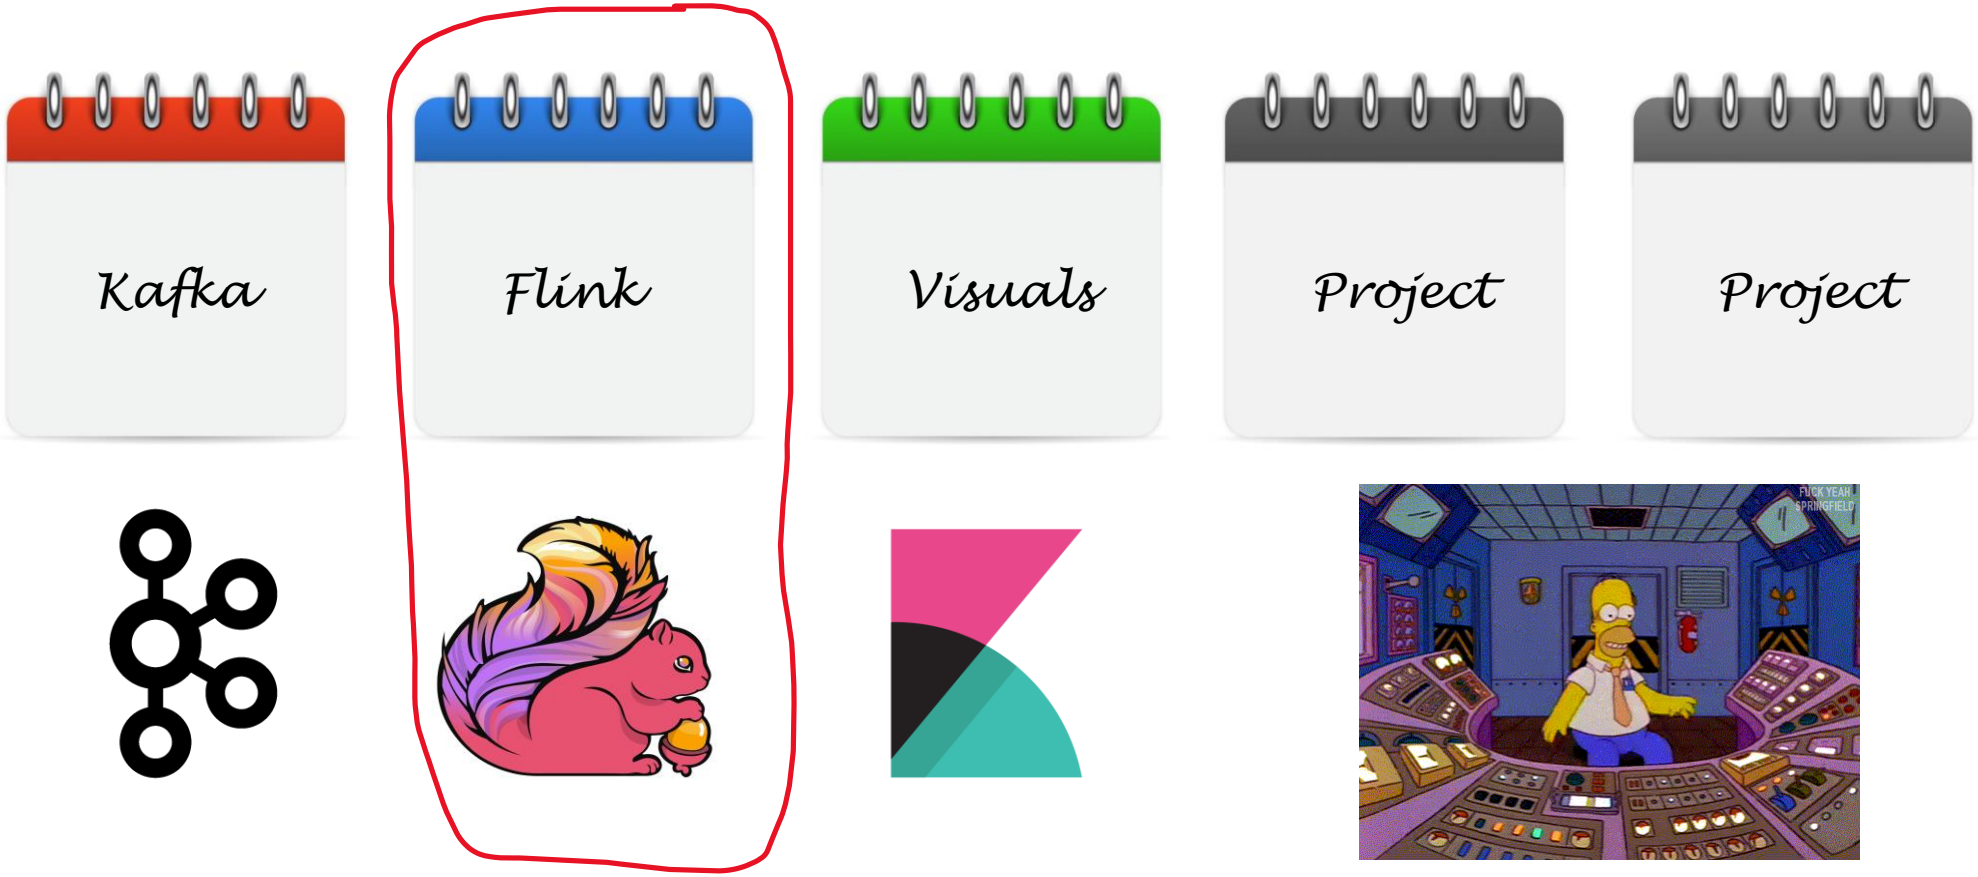

Agenda for today

Agenda for today

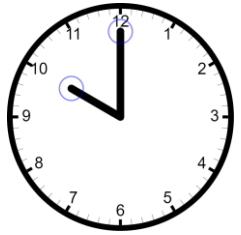

Data Streaming Concepts and Tools

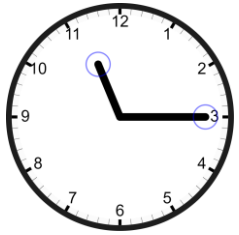

Agenda for today

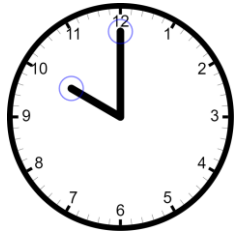

Data Streaming Concepts and Tools

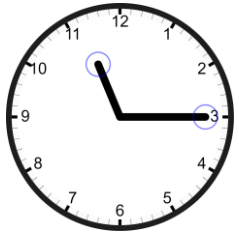

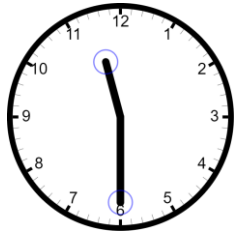

Apache Flink setup, practice

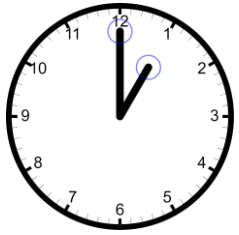

Agenda for today

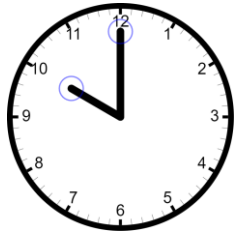

Data Streaming Concepts and Tools

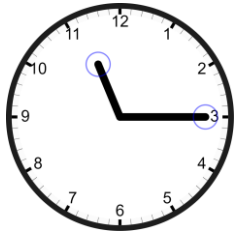

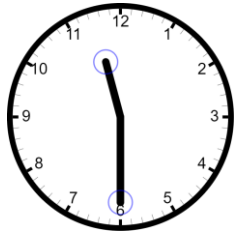

Apache Flink setup, practice

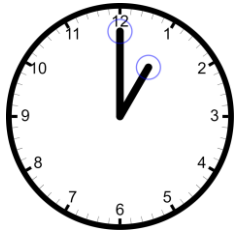

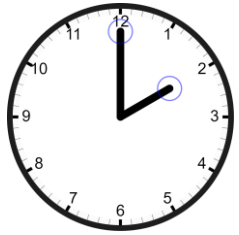

Connecting Apache Flink and Kafka

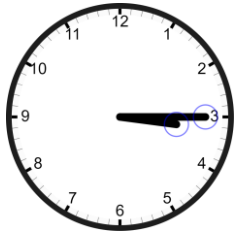

Use the right tool

Expectation

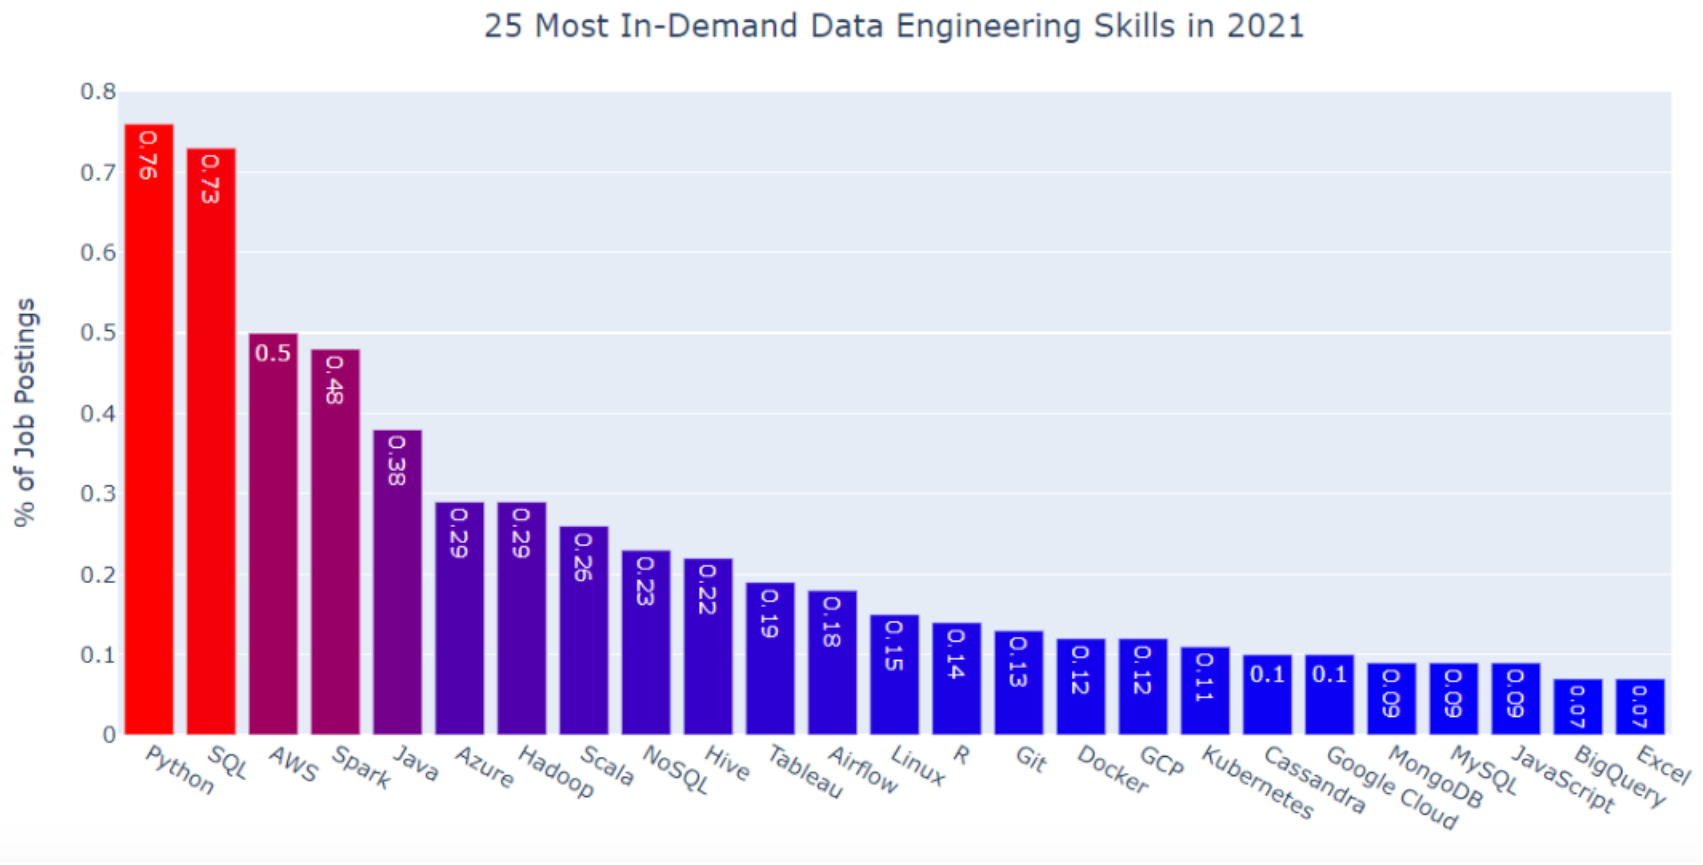

Source: https://www.kdnuggets.com/2021/05/most-demand-skills-data-engineers-2021.html

Classification

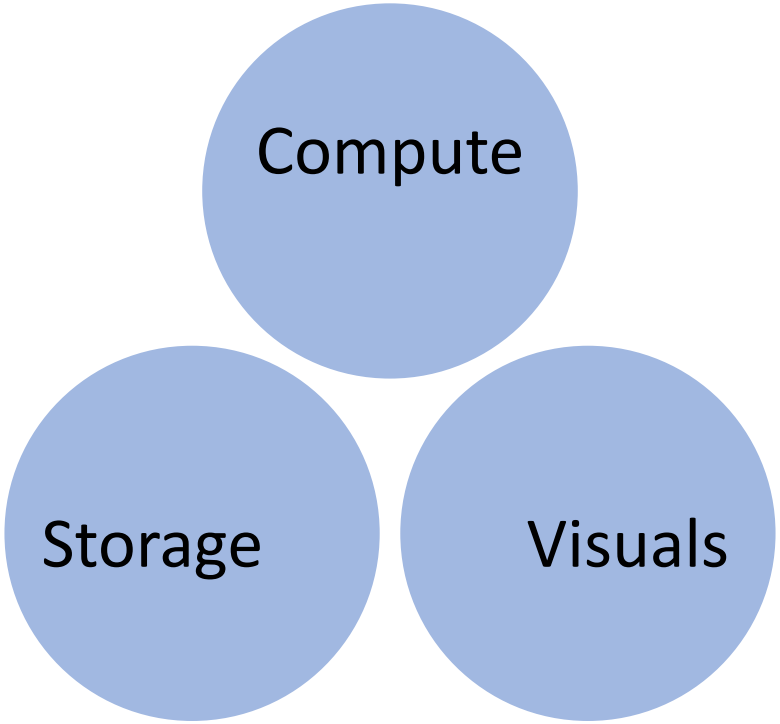

Classification

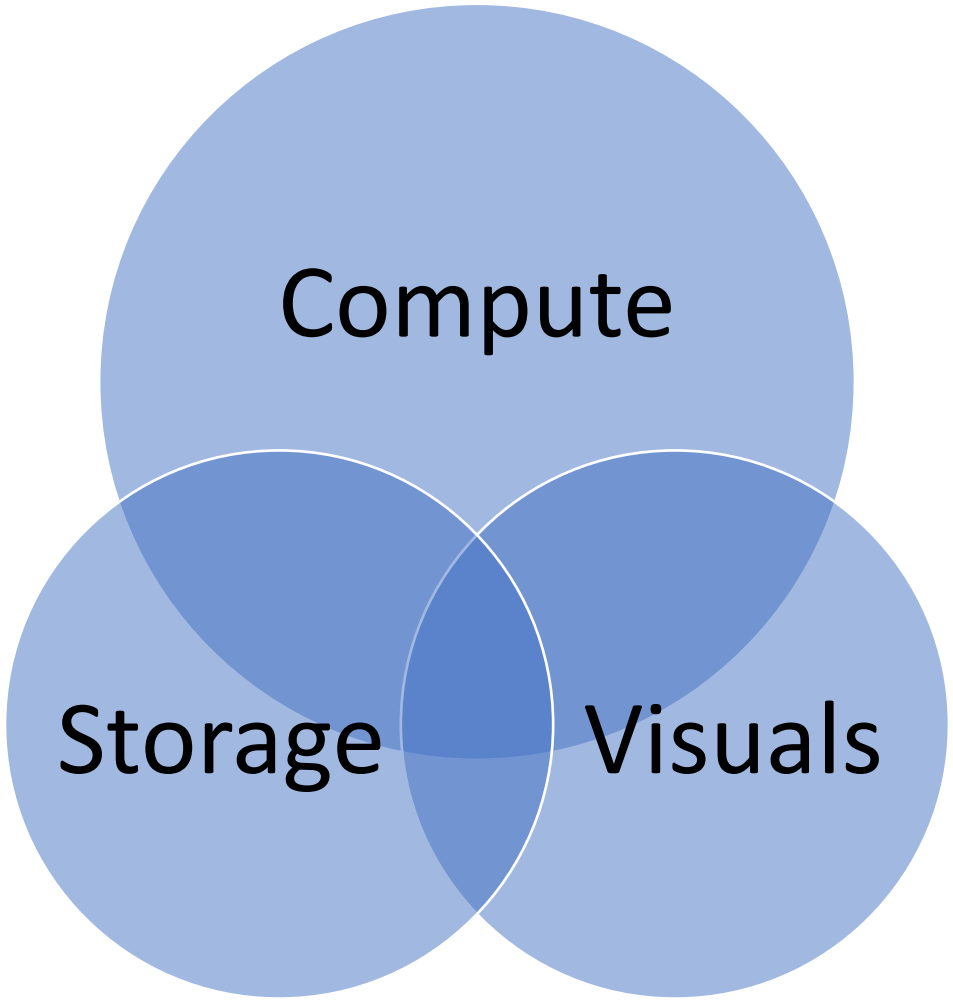

Classification

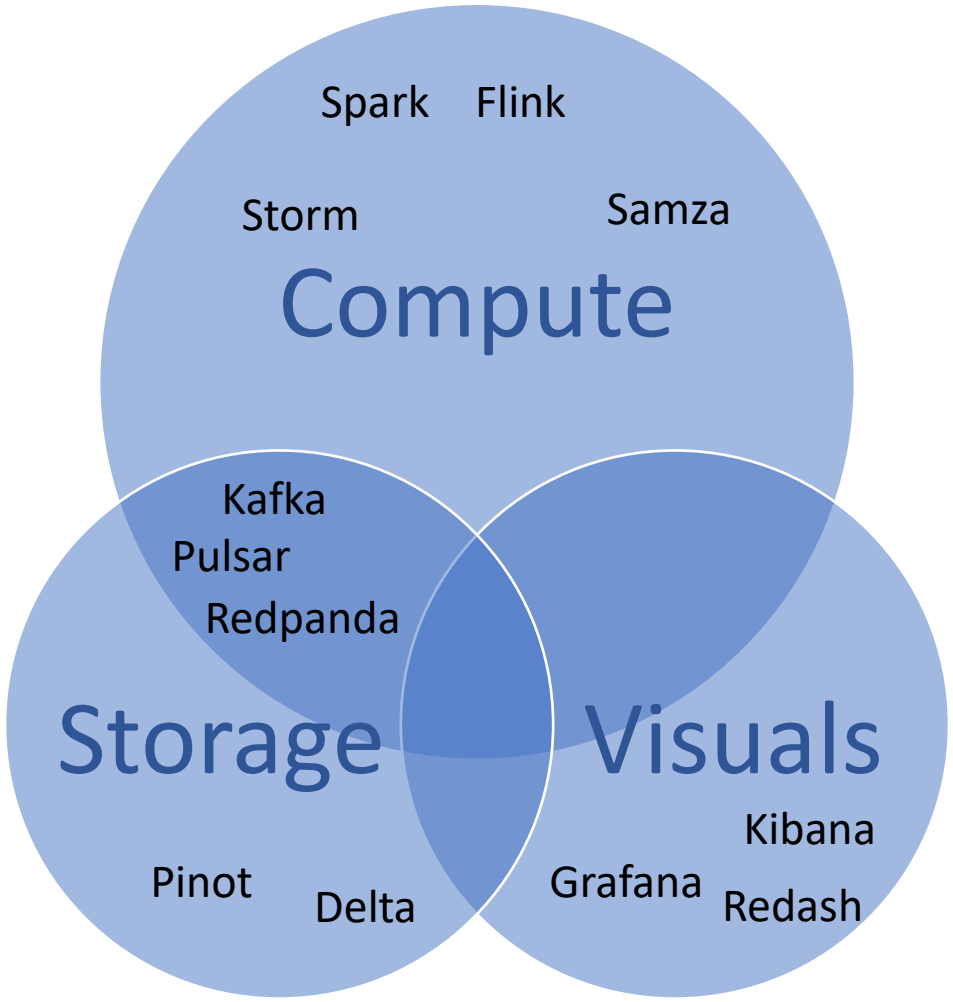

Note: arbitrary classification of open-source frameworks


BECAUSE...
IT'S IN THE CLOUD
HD
RY.COM
ator.net


aws


Microsoft
Azure


Google Cloud Platform

			
MSK	Event Hubs	Pub/Sub	Kafka

			
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink

aws	Microsoft Azure	Google Cloud Platform	
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink
Glue	Data Factory	Cloud Data Fusion	ETL

aws	Microsoft Azure	Google Cloud Platform	
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink
Glue	Data Factory	Cloud Data Fusion	ETL
EMR	HDInsight	Dataproc	Managed open-source (Hadoop)

aws	Microsoft Azure	Google Cloud Platform	
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink
Glue	Data Factory	Cloud Data Fusion	ETL
EMR	HDInsight	Dataproc	Managed open-source (Hadoop)
Redshift	Synapse	BigQuery	Data warehouse

aws	Microsoft Azure	Google Cloud Platform	
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink
Glue	Data Factory	Cloud Data Fusion	ETL
EMR	HDInsight	Dataproc	Managed open-source (Hadoop)
Redshift	Synapse	BigQuery	Data warehouse
S3	Blob storage	Cloud storage	Storage / Data lake

aws	Microsoft Azure	Google Cloud Platform	
MSK	Event Hubs	Pub/Sub	Kafka
Kinesis	Stream Analytics	Dataflow	Flink
Glue	Data Factory	Cloud Data Fusion	ETL
EMR	HDInsight	Dataproc	Managed open-source (Hadoop)
Redshift	Synapse	BigQuery	Data warehouse
S3	Blob storage	Cloud storage	Storage / Data lake
Quicksight	Power BI	Looker, Data Studio	Visualization

Cloud agnostic / SaaS vendors

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda

Kafka

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink
Fivetran | Talend | Hevo | Matillion | …	ETL

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink
Fivetran | Talend | Hevo | Matillion | …	ETL
Databricks | Cloudera	Managed open-source (Hadoop)

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink
Fivetran | Talend | Hevo | Matillion | …	ETL
Databricks | Cloudera	Managed open-source (Hadoop)
Snowflake | Vertica | Databricks	Data warehouse

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink
Fivetran | Talend | Hevo | Matillion | …	ETL
Databricks | Cloudera	Managed open-source (Hadoop)
Snowflake | Vertica | Databricks	Data warehouse
Databricks | Snowflake	Storage / Data lake

Note: arbitrary incomplete list

Cloud agnostic / SaaS vendors

Confluent | Conduktor | Redpanda	Kafka
Ververica | Aiven	Flink
Fivetran | Talend | Hevo | Matillion | …	ETL
Databricks | Cloudera	Managed open-source (Hadoop)
Snowflake | Vertica | Databricks	Data warehouse
Databricks | Snowflake	Storage / Data lake
Qlik | Tableau	Visualization

Note: arbitrary incomplete list

Community/trial versions to try out

- https://confluent.cloud/signup
  - 60 days
- https://databricks.com/try-databricks
  - Free (forever?)
- https://www.ververica.com/getting-started
  - Free, need to install on K8s cluster

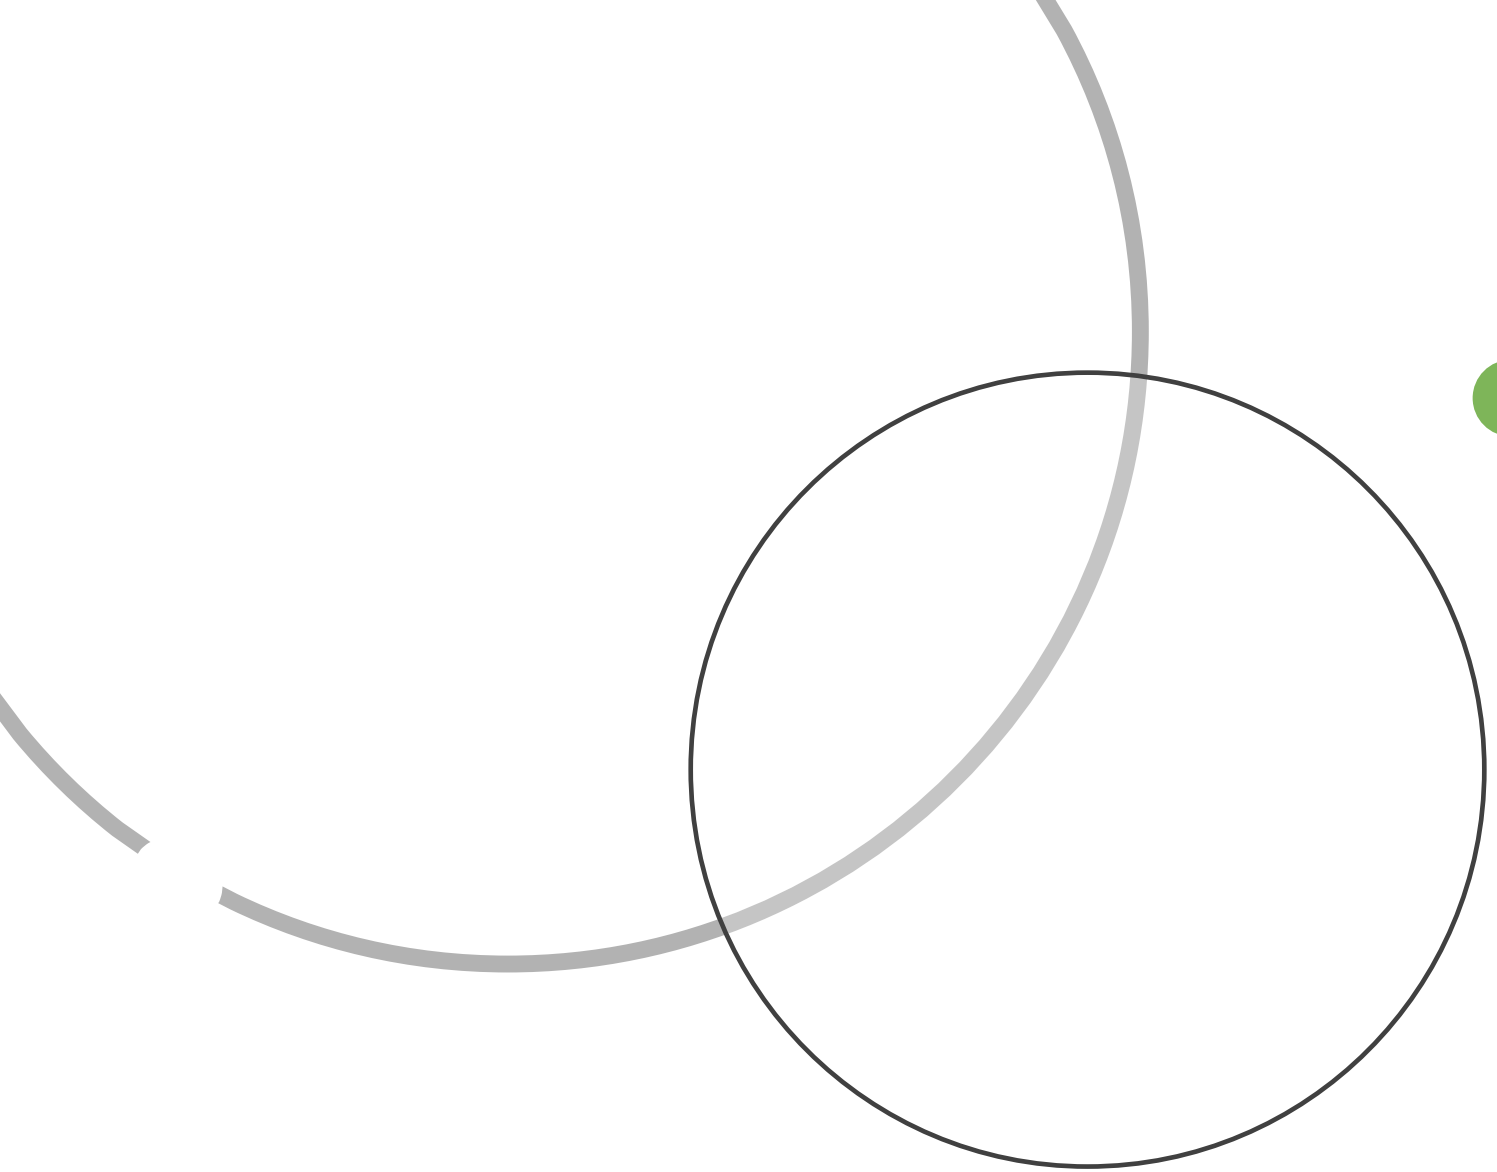

Flink

Flink

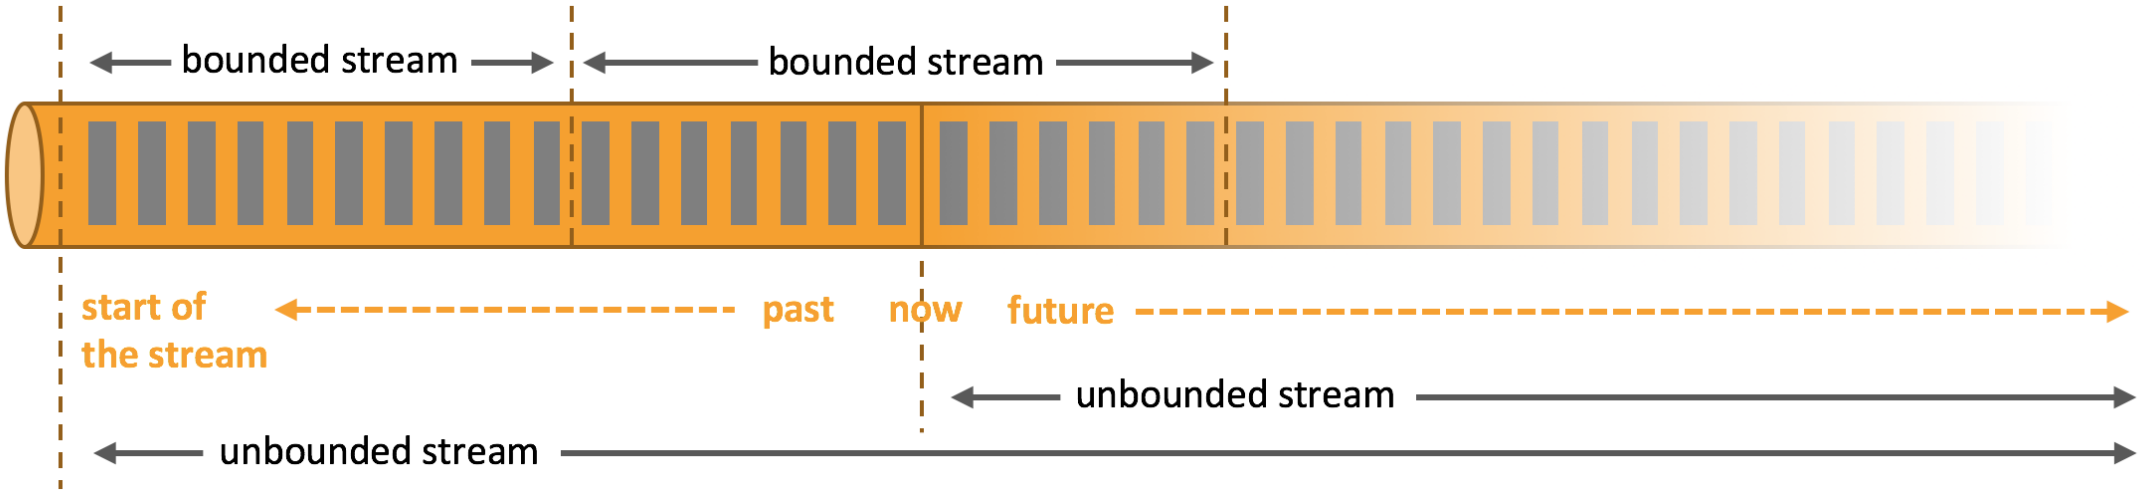
bounded stream
bounded stream
start of
the stream
past
now
future
unbounded stream
unbounded stream

Streaming Dataflow

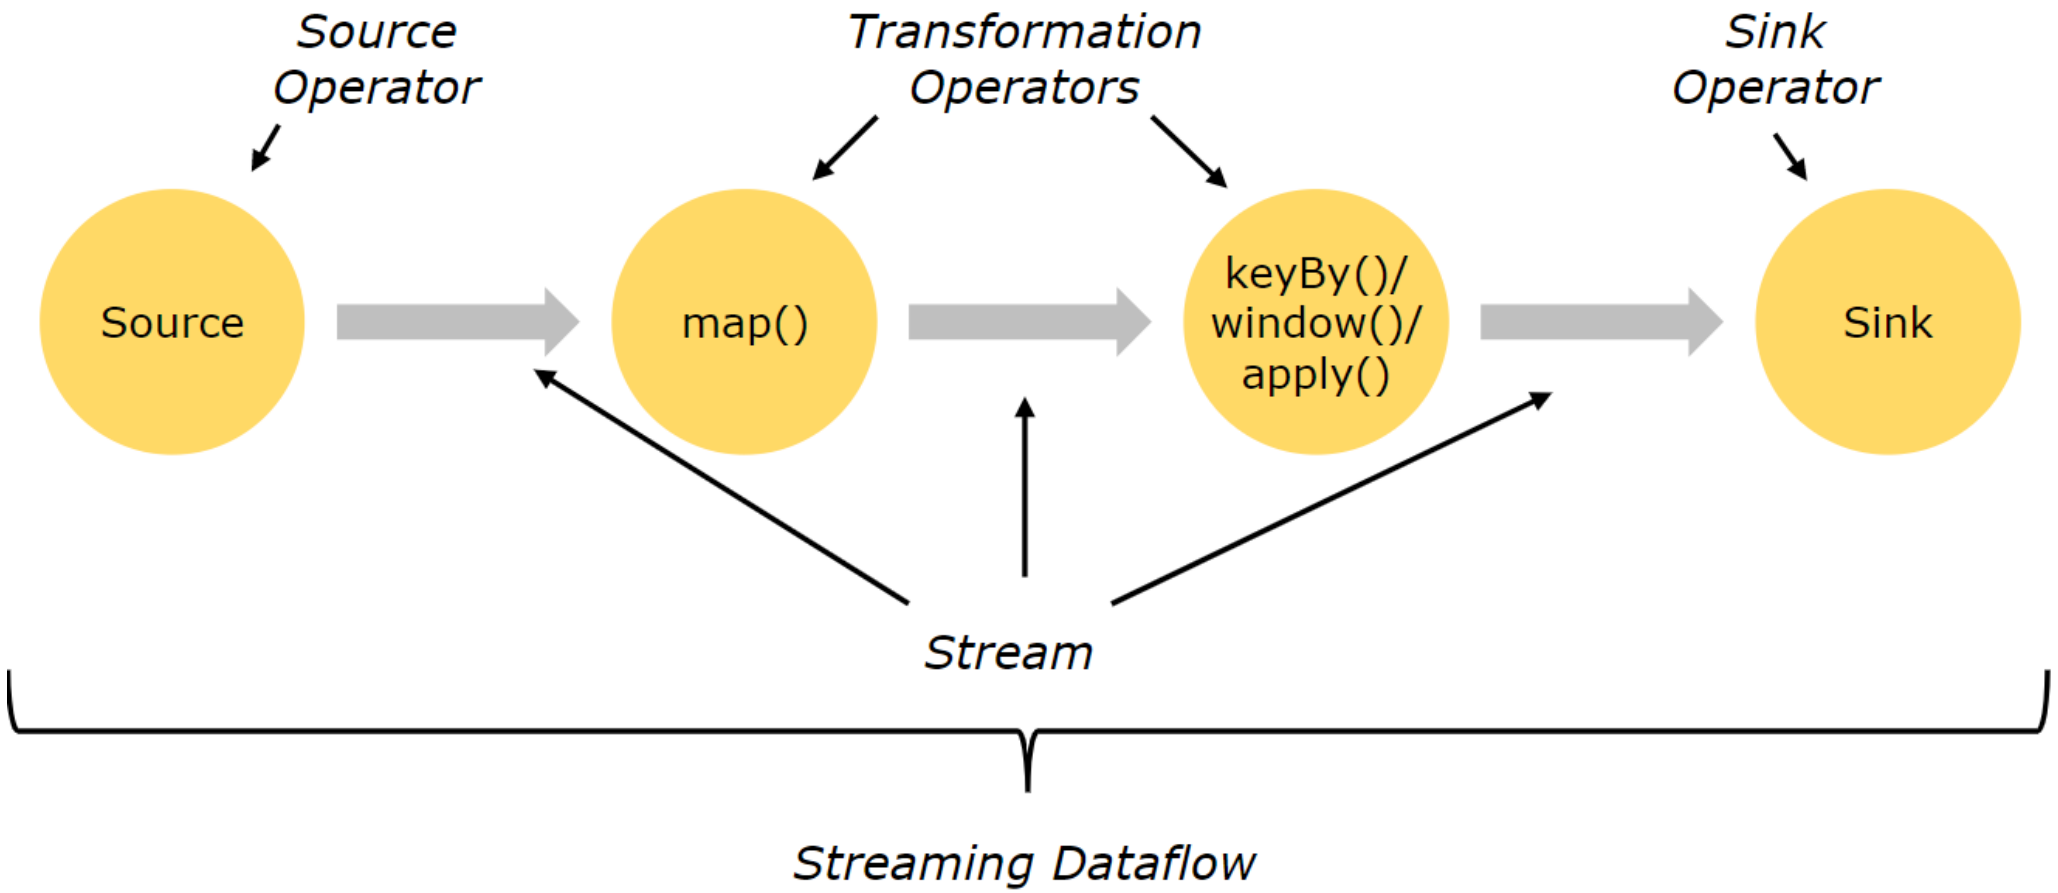

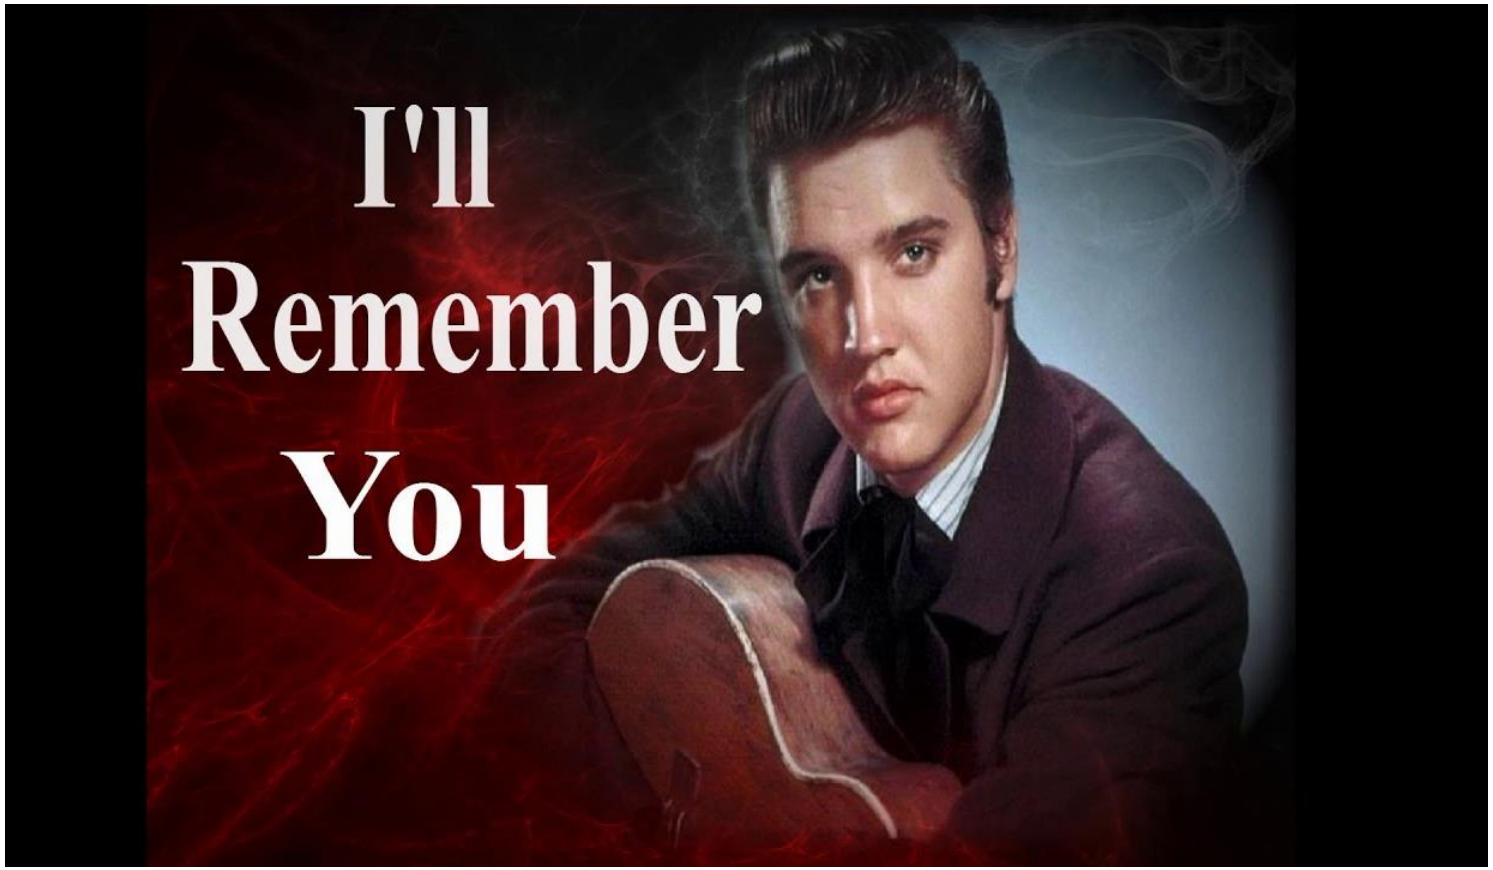

Stateful streaming

- One individual *event at a time* (for example an event parser)

Vs

- Remember information across multiple events (for example window operators). = **stateful**

Examples of stateful operations:

- When an application searches for certain event patterns, the state will store the sequence of events encountered so far.
- When aggregating events per minute/hour/day, the state holds the pending aggregates.
- When training a machine learning model over a stream of data points, the state holds the current version of the model parameters.
- When historic data needs to be managed, the state allows efficient access to events that occurred in the past.

What time is it?

What time is it?

Processing time: system time of the machine that is executing the respective operation.

Event time: the time that each individual event occurred on its producing device.

What time is it?

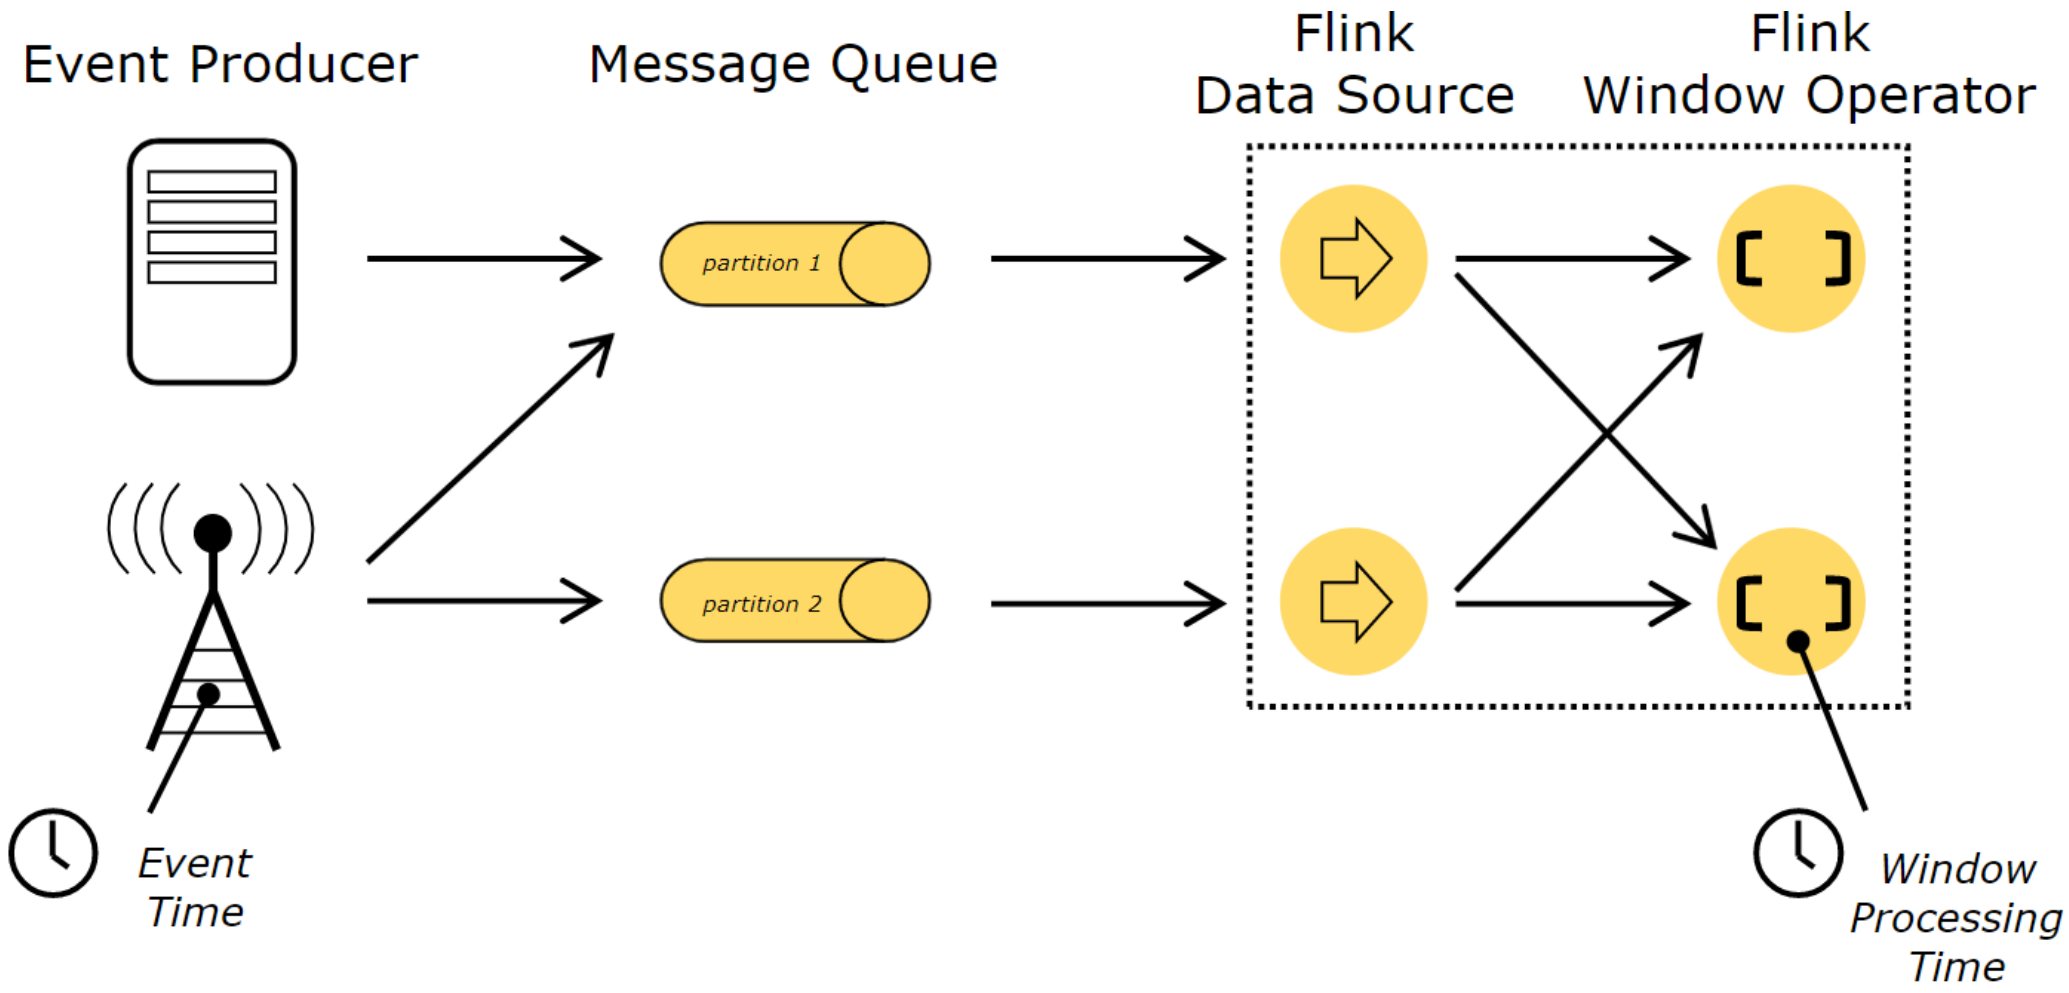

Windows

```
Flink SQL> SELECT * FROM Bid;
+------------------+-------+------+-------------+
|          bidtime | price | item | supplier_id |
+------------------+-------+------+-------------+
| 2020-04-15 08:05 | 4.00  | C    | supplier1   |
| 2020-04-15 08:07 | 2.00  | A    | supplier1   |
| 2020-04-15 08:09 | 5.00  | D    | supplier2   |
| 2020-04-15 08:11 | 3.00  | B    | supplier2   |
| 2020-04-15 08:13 | 1.00  | E    | supplier1   |
| 2020-04-15 08:17 | 6.00  | F    | supplier2   |
+------------------+-------+------+-------------+
```

```
Flink SQL> SELECT * FROM Bid;
+------------------+-------+------+-------------+
|          bidtime | price | item | supplier_id |
+------------------+-------+------+-------------+
| 2020-04-15 08:05 | 4.00  | C    | supplier1   |
| 2020-04-15 08:07 | 2.00  | A    | supplier1   |
| 2020-04-15 08:09 | 5.00  | D    | supplier2   |
| 2020-04-15 08:11 | 3.00  | B    | supplier2   |
| 2020-04-15 08:13 | 1.00  | E    | supplier1   |
| 2020-04-15 08:17 | 6.00  | F    | supplier2   |
+------------------+-------+------+-------------+
```

```
-- tumbling window aggregation
Flink SQL> SELECT window_start, window_end, SUM(price)
  FROM TABLE(
    TUMBLE(TABLE Bid, DESCRIPTOR(bidtime), INTERVAL '10' MINUTES))
  GROUP BY window_start, window_end;
+------------------+------------------+-------+
|     window_start |       window_end | price |
+------------------+------------------+-------+
| 2020-04-15 08:00 | 2020-04-15 08:10 | 11.00 |
| 2020-04-15 08:10 | 2020-04-15 08:20 | 10.00 |
+------------------+------------------+-------+
```

```
Flink SQL> SELECT * FROM Bid;
+------------------+-------+------+-------------+
|          bidtime | price | item | supplier_id |
+------------------+-------+------+-------------+
| 2020-04-15 08:05 | 4.00  | C    | supplier1   |
| 2020-04-15 08:07 | 2.00  | A    | supplier1   |
| 2020-04-15 08:09 | 5.00  | D    | supplier2   |
| 2020-04-15 08:11 | 3.00  | B    | supplier2   |
| 2020-04-15 08:13 | 1.00  | E    | supplier1   |
| 2020-04-15 08:17 | 6.00  | F    | supplier2   |
+------------------+-------+------+-------------+
```

```
-- cumulative window aggregation
Flink SQL> SELECT window_start, window_end, SUM(price)
  FROM TABLE(
    CUMULATE(TABLE Bid, DESCRIPTOR(bidtime), INTERVAL '2' MINUTES, INTERVAL '10' MINUTES))
  GROUP BY window_start, window_end;
+------------------+------------------+-------+
|     window_start |       window_end | price |
+------------------+------------------+-------+
| 2020-04-15 08:00 | 2020-04-15 08:06 | 4.00  |
| 2020-04-15 08:00 | 2020-04-15 08:08 | 6.00  |
| 2020-04-15 08:00 | 2020-04-15 08:10 | 11.00 |
| 2020-04-15 08:10 | 2020-04-15 08:12 | 3.00  |
| 2020-04-15 08:10 | 2020-04-15 08:14 | 4.00  |
| 2020-04-15 08:10 | 2020-04-15 08:16 | 4.00  |
| 2020-04-15 08:10 | 2020-04-15 08:18 | 10.00 |
| 2020-04-15 08:10 | 2020-04-15 08:20 | 10.00 |
+------------------+------------------+-------+
```


THE
PORTER
This isn’t going to
stop, is it?

Watermark

- A *Watermark(t)* declares that event time has reached time *t* in that stream, meaning that there should be no more elements from the stream with a timestamp *t'* <= *t* (i.e. events with timestamps older or equal to the watermark).


Watermark
example

Watermark vs window

Watermark	Window
Progress of event time	Logical partition of events
Based on time	Based on time or count
Goal: Determine when is a portion of data processed, Handle out-of-order streams	Goal: Aggregations

Flink architecture

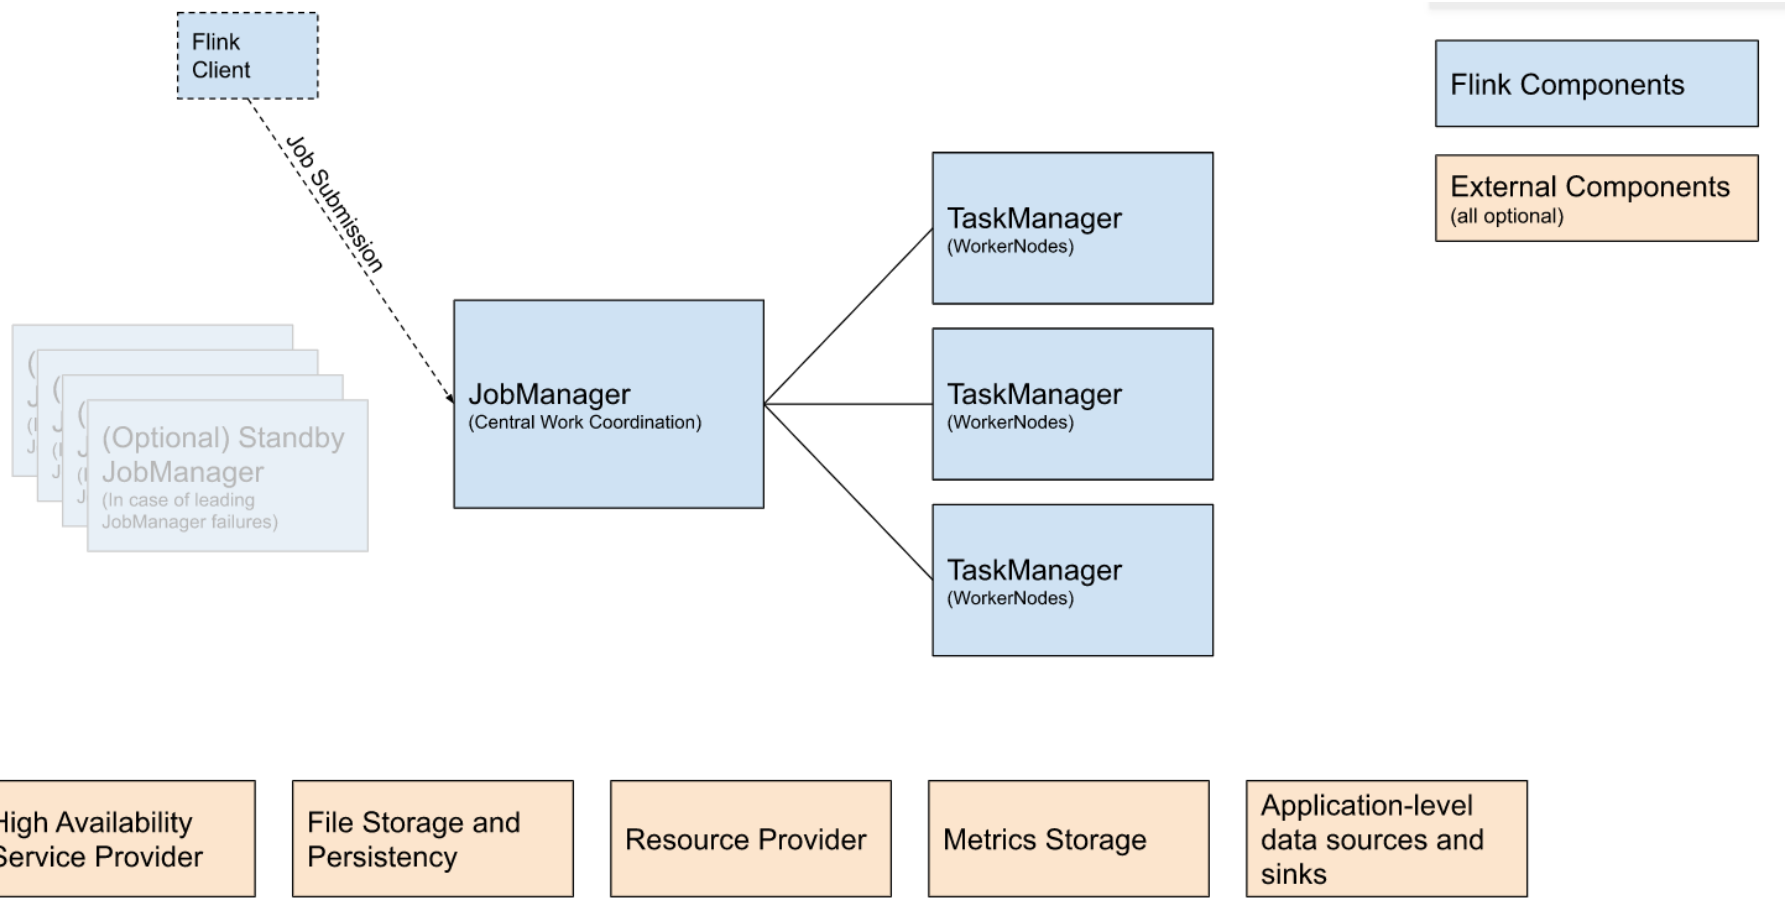

Flink architecture

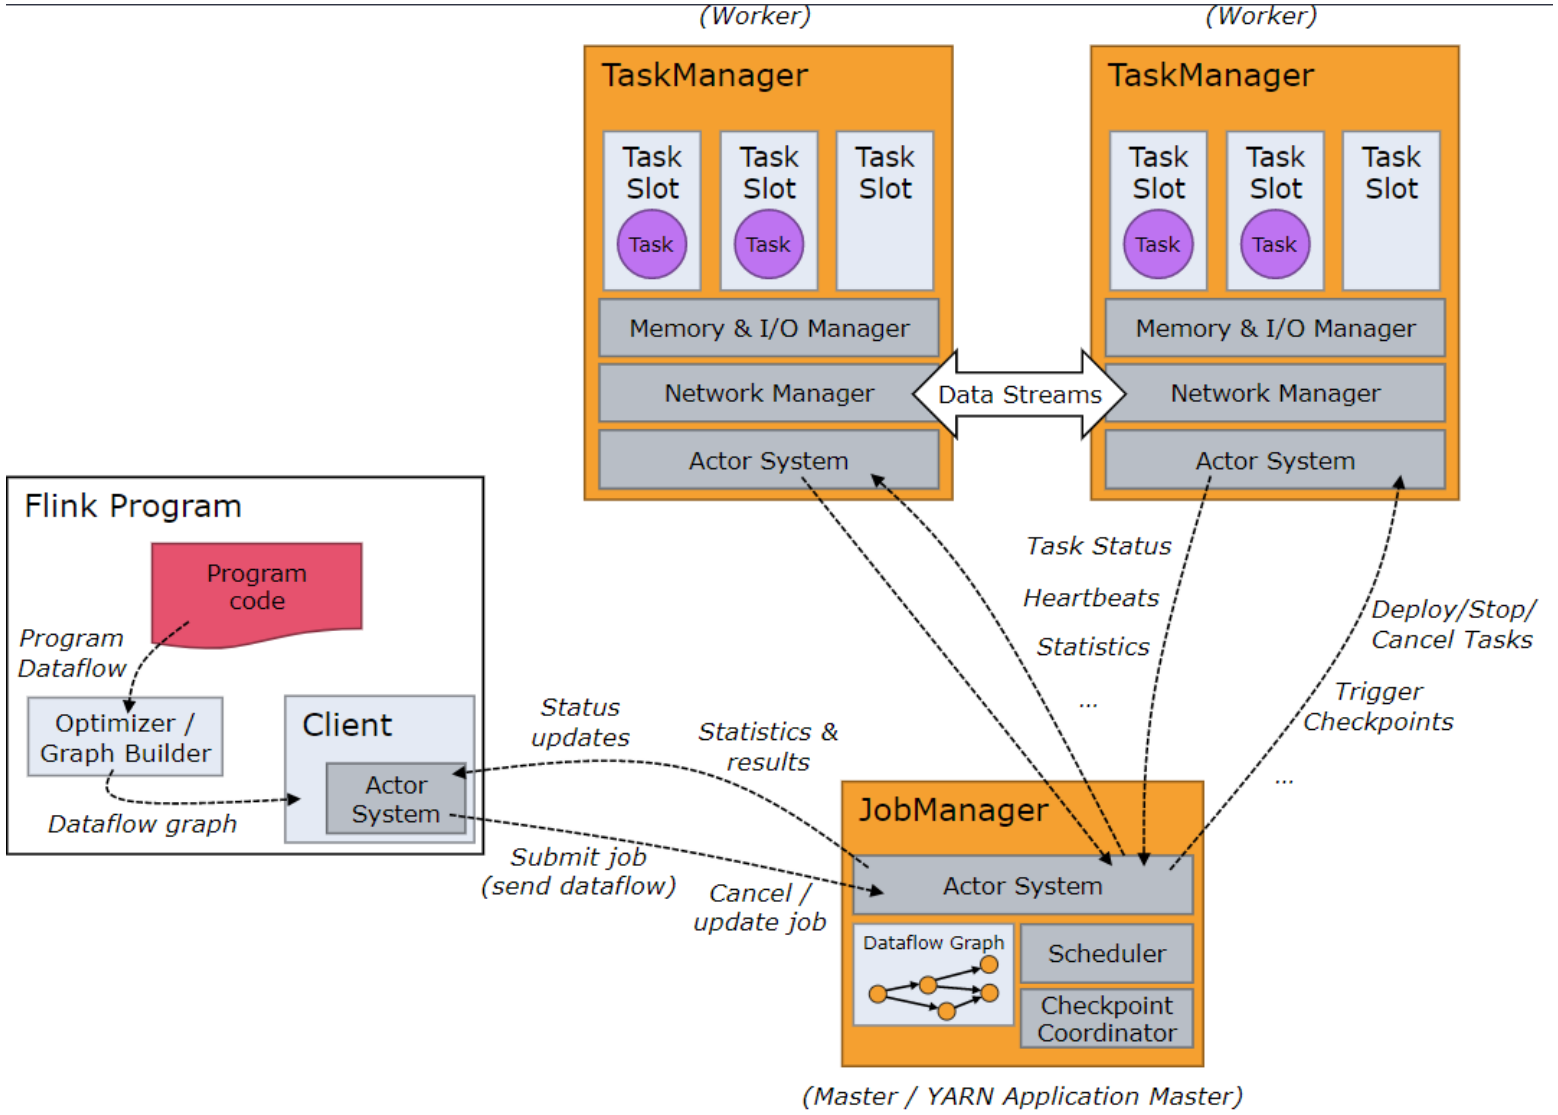

Flink Deployment

Application Mode

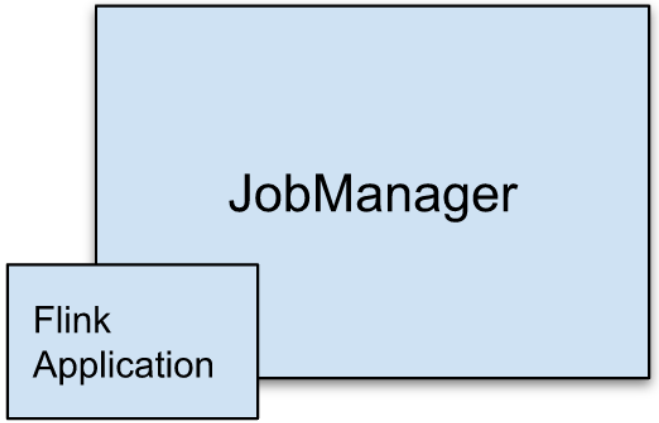

A dedicated JobManager is started for submitting the job. The JobManager will only execute this job, then exit.
The Flink Application runs on the JobManager.

Session Mode

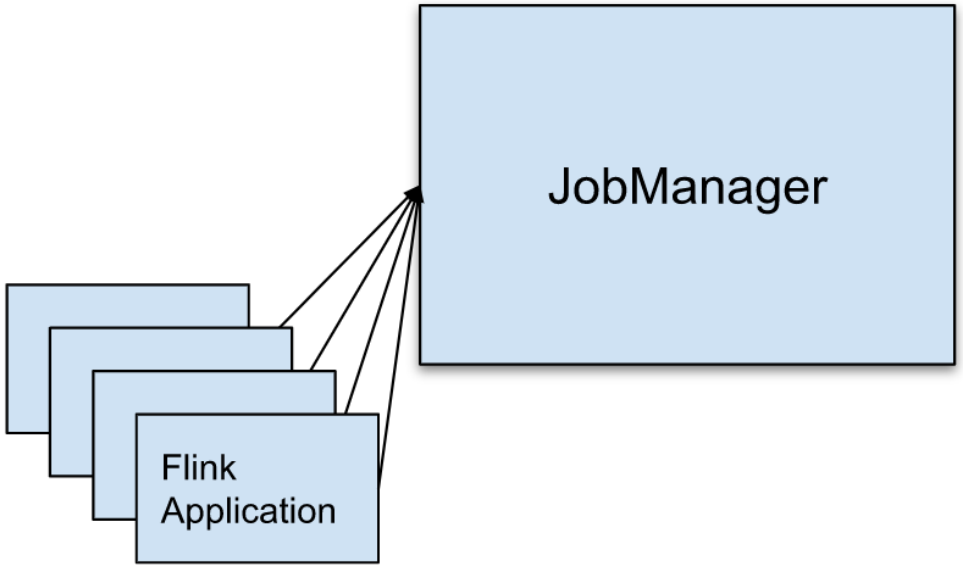

Multiple jobs share one JobManager.

Flink APIs

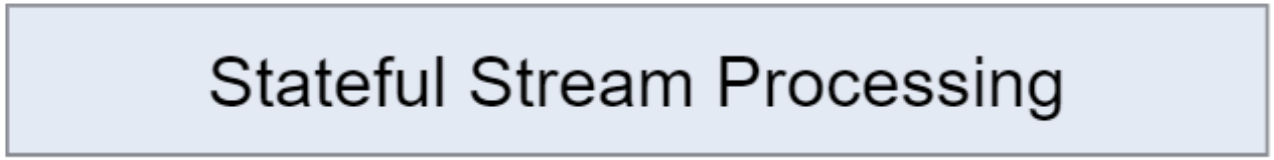

Low-level building block
(streams, state, [event] time)

Flink APIs

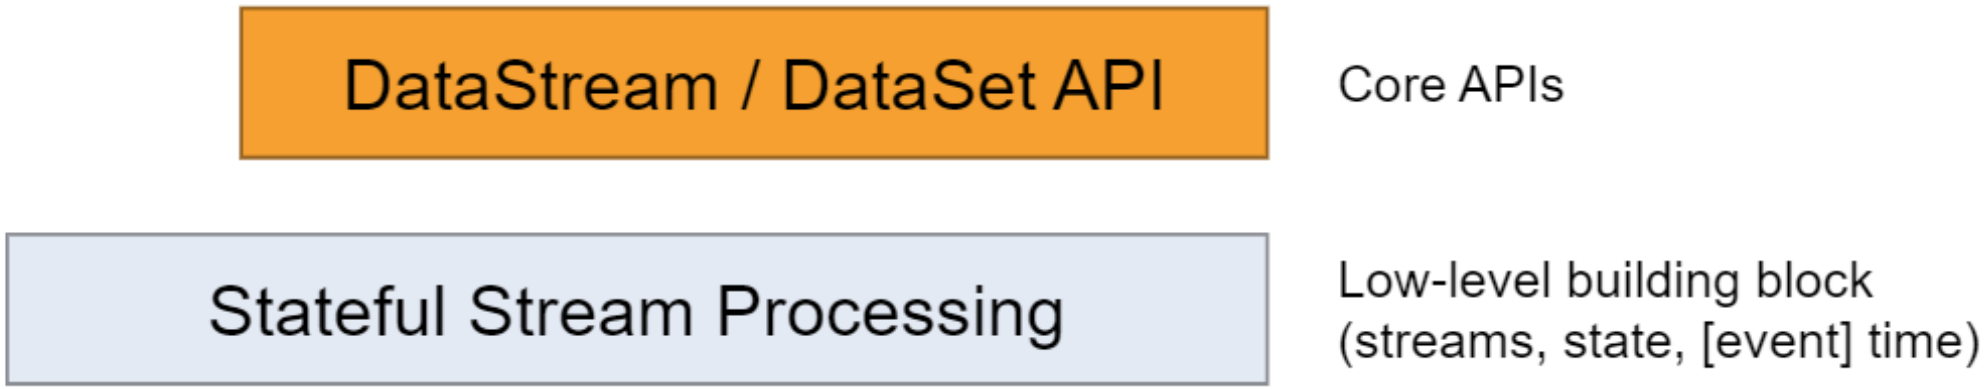
DataStream / DataSet API
Core APIs
Stateful Stream Processing
Low-level building block
(streams, state, [event] time)

Flink APIs

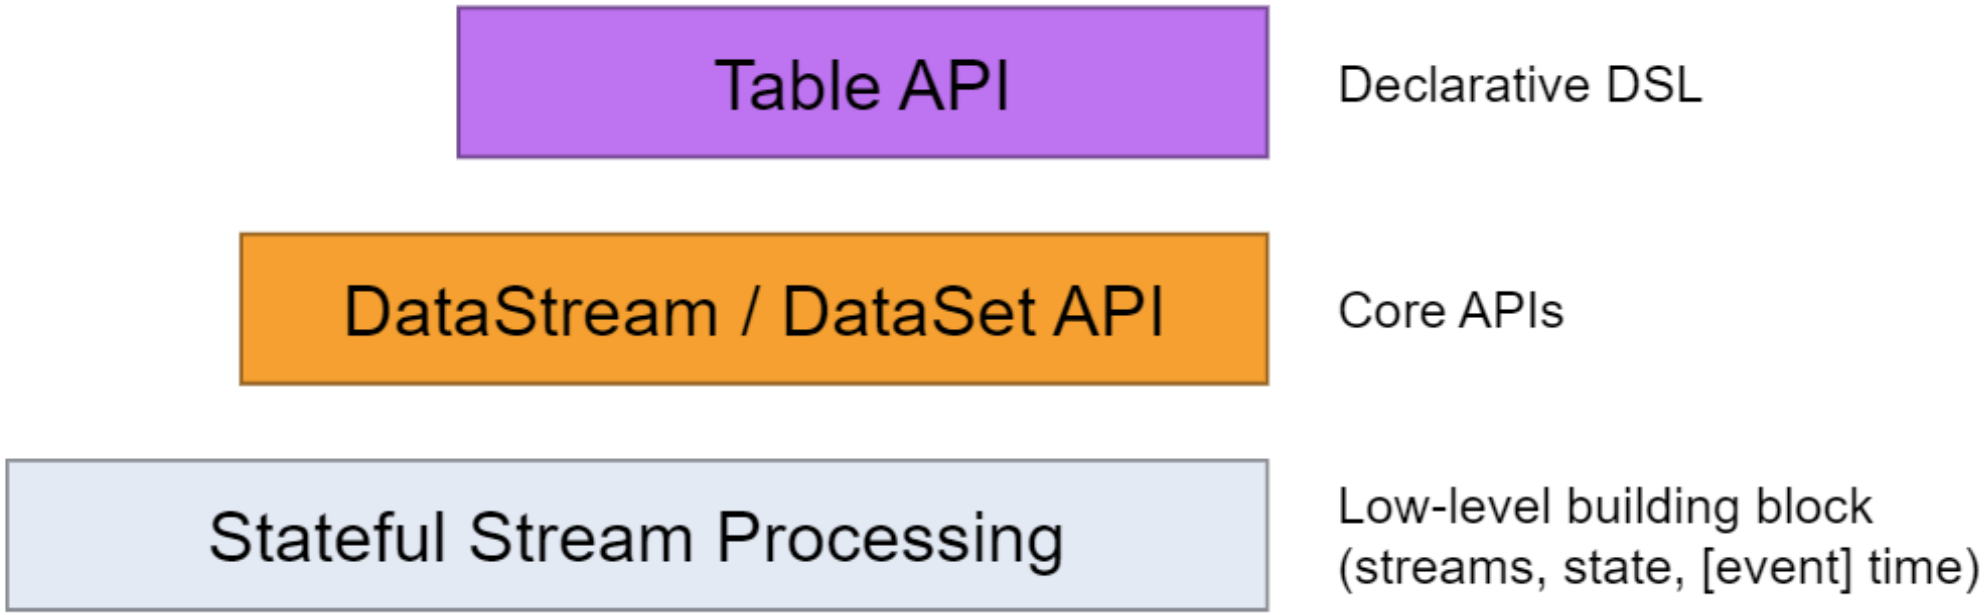
Table API
Declarative DSL
DataStream / DataSet API
Core APIs
Stateful Stream Processing
Low-level building block
(streams, state, [event] time)

Flink APIs

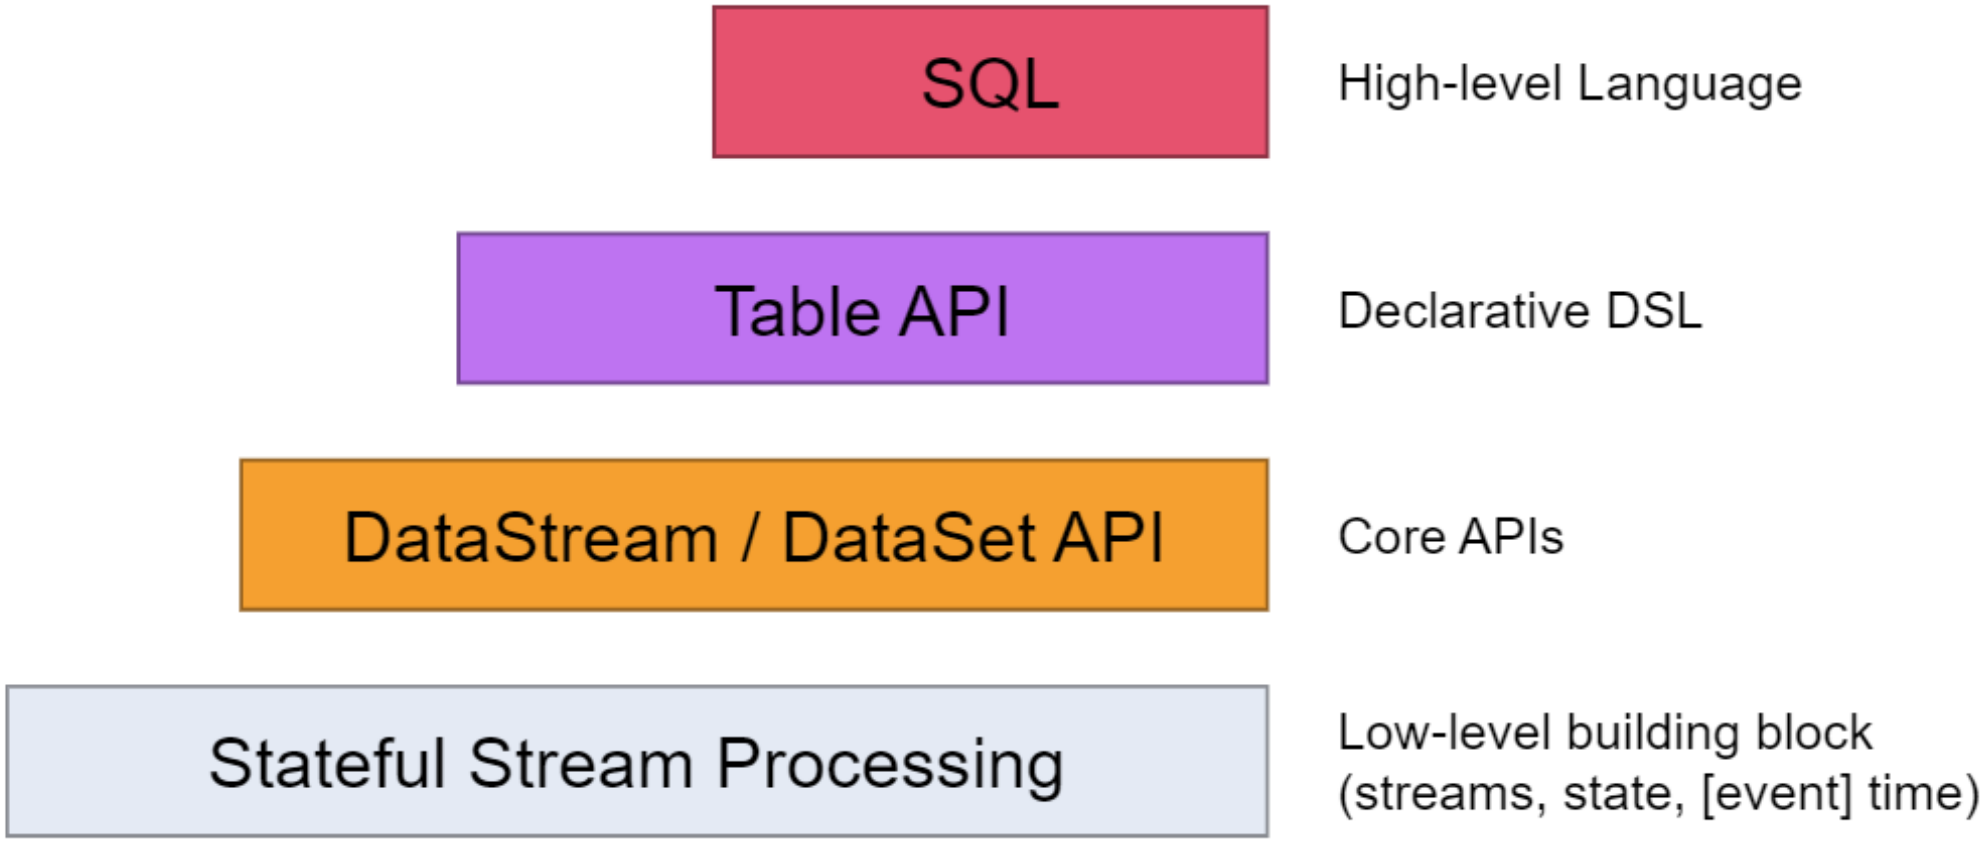
SQL
High-level Language
Table API
Declarative DSL
DataStream / DataSet API
Core APIs
Stateful Stream Processing
Low-level building block
(streams, state, [event] time)

Flink APIs

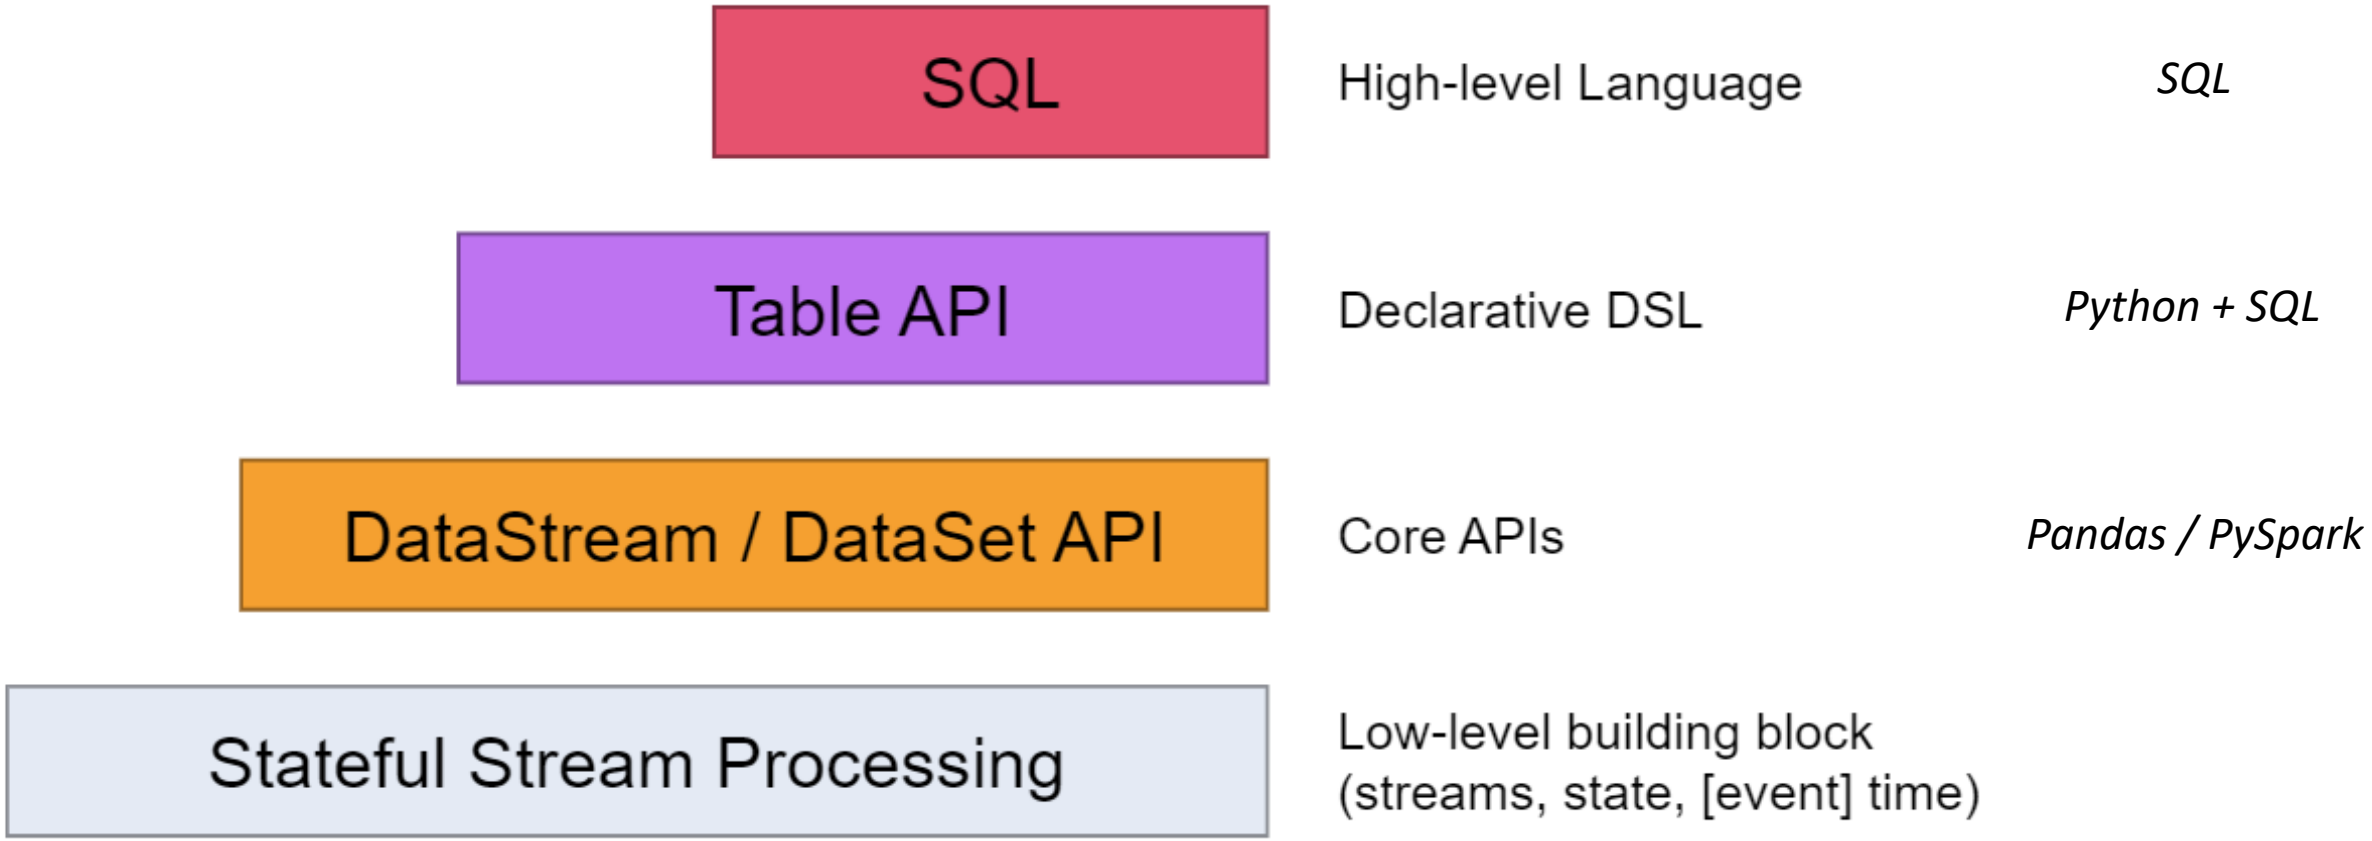
SQL
High-level Language
SQL
Table API
Declarative DSL
Python + SQL
DataStream / DataSet API
Core APIs
Pandas / PySpark
Stateful Stream Processing
Low-level building block
(streams, state, [event] time)


LETS ROLL
#silverdollarcity
SILVER DOLLAR CITY